



Emulex® Drivers for Linux User Manual

Copyright © 2003-2013 Emulex. All rights reserved worldwide. No part of this document may be reproduced by any means or translated to any electronic medium without the prior written consent of Emulex.

Information furnished by Emulex is believed to be accurate and reliable. However, no responsibility is assumed by Emulex for its use; or for any infringements of patents or other rights of third parties which may result from its use. No license is granted by implication or otherwise under any patent, copyright or related rights of Emulex.

Emulex, the Emulex logo, AutoPilot Installer, AutoPilot Manager, BlockGuard, Connectivity Continuum, Convergenomics, Emulex Connect, Emulex Secure, EZPilot, FibreSpy, HBAnyware, InSpeed, LightPulse, MultiPulse, OneCommand, OneConnect, One Network. One Company., SBOD, SLI, and VEngine are trademarks of Emulex. All other brand or product names referenced herein are trademarks or registered trademarks of their respective companies or organizations.

Emulex provides this manual "as is" without any warranty of any kind, either expressed or implied, including but not limited to the implied warranties of merchantability or fitness for a particular purpose. Emulex may make improvements and changes to the product described in this manual at any time and without any notice. Emulex assumes no responsibility for its use, nor for any infringements of patents or other rights of third parties that may result. Periodic changes are made to information contained herein; although these changes will be incorporated into new editions of this manual, Emulex disclaims any undertaking to give notice of such changes.

Emulex, 3333 Susan Street

Costa Mesa, CA 92626

OpenSolaris DHCHAP Notice

Contains portions of Covered Software subject to the Common Development and Distribution License (CDDL) Version 1.0. Such portions of Covered Software in Source Code form may be obtained from the website www.opensolaris.org, or by contacting online support from the website www.emulex.com. Derived from the RSA Data Security, Inc. MD5 Message-Digest Algorithm. Copyright (C) 1991-2, RSA Data Security, Inc. Created 1991. All rights reserved.

Supported Driver Versions

The following table lists the Emulex-supported Fibre Channel (FC), Fibre Channel over Ethernet (FCoE), Ethernet, and internet Small Computer System Interface (iSCSI) drivers that are applicable in this manual.

A check mark “✓” indicates the type of driver distribution that is supported.

Driver Version	Driver Distribution			Supported Operating Systems
	Inbox	DUP/DUD	Out-of-Box	
FC and FCoE Drivers				
8.2.2.15			✓	RHEL 5.6-5.9, SLES 10 SP3-SP4, OL 5.6-5.9, CentOS 5.6-5.9
8.3.7.18			✓	RHEL 6.1-6.4, SLES 11 SP1-SP2, OL 6.1-6.4, CentOS 6.1-6.4, Oracle Linux UEK R1, R2 with OL 5, Oracle Linux UEK R1, R2 with OL 6
Ethernet Drivers				
4.6.142.5			✓	RHEL 5.5-5.9, RHEL 6.1-6.4, SLES 10 SP3-SP4, SLES 11 SP1-SP2, OL 5.5-5.8, OL 6.0-6.4, CentOS 5.5-5.8, CentOS 6.0-6.4, Oracle Linux UEK R1, R2 with OL 5, Oracle Linux UEK R1, R2 with OL 6, Debian 6.0.7, Ubuntu 11.10, Ubuntu 12.10
4.6.148.0			✓	RHEL 5.5-5.9, RHEL 6.1-6.4, SLES 10 SP3-SP4, SLES 11 SP1-SP2
iSCSI Drivers				
4.6.142.0			✓	RHEL 5.5-5.9, RHEL 6.1-6.4, SLES 10 SP3-SP4, SLES 11 SP1-SP2, OL 5.5-5.8, OL 6.0-6.4, CentOS 5.5-5.8, CentOS 6.0-6.4, Oracle Linux UEK R1, R2 with OL 5, Oracle Linux UEK R1, R2 with OL 6
4.6.148.0			✓	RHEL 5.5-5.9, RHEL 6.1-6.4, SLES 10 SP3-SP4, SLES 11 SP1-SP2

Table of Contents

Supported Driver Versions.....	3
List of Tables	8
1. Introduction	9
Overview	9
Known Issues.....	10
Abbreviations	10
2. Installing and Uninstalling	14
General Installation Requirements	14
Binary RPM FC and FCoE Driver Kit.....	14
Installing the Binary RPM FC and FCoE Driver Kit.....	15
Uninstalling the Binary RPM FC and FCoE Driver Kit.....	15
Ethernet Driver Kit.....	16
Installing the Ethernet Driver Kit.....	16
Uninstalling the Ethernet Driver Kit	16
iSCSI Driver Kit.....	17
Installing the iSCSI Driver Kit	17
Uninstalling the iSCSI Driver Kit	18
Bootting from a Non-Zero LUN Attached to an Emulex LPFC Adapter.....	18
OneCommand Manager Application.....	18
3. Configuration	19
FC and FCoE Driver Configuration	19
FC and FCoE Driver Parameters	19
Static FC and FCoE Driver Parameters.....	19
Dynamic FC and FCoE Driver Parameters	22
Configuring FC and FCoE Driver Parameters.....	24
Temporary Configuration with modprobe.....	24
Persistent Configuration with modprobe.conf or lpfc.conf	24
Configure Parameters with a Read/Write to sysfs.....	25
Creating a New Ramdisk Image	26
Dynamically Recognizing LUNs and Targets (Using scan)	26
Persistent Naming	27
Using udev to Discover Logical to Physical Mappings for sd Devices.....	27
Configuring the System to Boot Using Persistent Names	27
Using udev with st Devices	28
Persistent Naming References	30

Ethernet Driver Configuration	31
Ethernet Driver Configuration Parameters.....	31
VLAN Support on UMC	31
SR-IOV Configuration	31
Introduction.....	31
Setting Up SR-IOV	32
Configuring VFs.....	36
Viewing VF Properties.....	38
iSCSI Driver Configuration	39
iSCSI Driver Parameter Descriptions	39
Configuring the iSCSI Driver Parameters	40
Configuring the ETO and LDTO Parameters.....	40
Configuring the large_io Parameter	40
Configuring the im_policy Parameter.....	41
iSCSI Error Handling	42
Error Handling using ETO and LDTO Parameters	42
Error Handling Under MultiPath I/O (MPIO) and Cluster Configurations.....	42
Virtual Port (VPort) Configuration	43
VPort Configuration Prerequisites.....	43
Creating, Deleting, and Displaying VPorts	43
Creating VPorts Using the mkvport.sh Script	43
Deleting VPorts Using the rmvport.sh Script	44
Displaying VPorts Using the lsvport.sh Script.....	44
VPort sysfs	44
VPort sysfs Tree	44
VPort sysfs Entries	46
Monitoring VPorts with fc_vport (FC and FCoE Drivers)	48
For FC 8.2.0.x Drivers	48
For FC 8.3.5.x Drivers	48
VPort Configuration Limits	49
DHCHAP Authentication and Configuration	50
Enabling Authentication	50
fcauthd Daemon.....	50
fcauthd Daemon Parameters.....	50
lpfc_enable_auth Module Parameter	51
Authentication Configuration Parameters	51
Setting Remote and Local Passwords	51
LPFC Driver Performance Tuning	51
Overview	51
lpfc_fcp_io_channel.....	51

lpfc_fcp_io_sched	52
lpfc_fcp_imax	52
lpfc_vector_map.sh	53
Network Performance Tuning	53
Memory Bandwidth Considerations.....	53
Enabling Optimal Bandwidth Options	54
Populate DIMM Slots	54
Disabling Memory Mirroring	54
Using a Fast Clock Speed for the Front Side Bus (FSB)	54
Network Memory Limits.....	54
TCP Segmentation Offload (TSO)	55
Flow Control	55
RX Frame Coalescing	56
Maximum Transmission Unit (MTU)	57
Interrupt Coalescing.....	57
Setting the Interrupt Delay Duration to a Range of Values (AIC)	57
Setting the Interrupt Delay Duration to a Constant Value.....	57
Receive Side Scaling (RSS).....	58
Analyzing Performance Issues.....	58
4. Troubleshooting	60
Situations and their Resolutions	60
FC and FCoE Driver Situations and their Resolutions.....	60
Ethernet Driver Situations and their Resolutions	64
iSCSI Driver Situations and their Resolutions	65
Log Messages	66
FC and FCoE Driver Log Messages	66
Retrieving FC and FCoE Driver Log Messages	66
LPFC Error Log Messages and their Descriptions	66
Ethernet Driver Log Messages.....	143
Retrieving Ethernet Driver Log Messages.....	143
Ethernet Driver Log Messages and their Descriptions	143
iSCSI Driver Log Messages.....	146
Retrieving iSCSI Driver Error Log Messages	146
iSCSI Driver Error Log Messages and their Descriptions	146
Appendix A. Open-iSCSI Support	148
Discovering and Adding iSCSI Targets	148
RHEL 5.6 and RHEL 5.7.....	148
Discovering and Adding iSCSI Targets.....	148
Adding iSCSI Targets Directly (without Discovering)	152

SLES 10 SP3, SLES 11 SP1, RHEL 5.8, and Later Versions	153
Discovering and Adding iSCSI Targets.....	153
iscsiadm Commands for Configuring the Target	153
Appendix B. Configuring iSCSI Through DHCP	155
IP Address Reservation.....	155
DHCP Option 43 (Vendor-Specific Information).....	155
DHCP Option 43 Parameter Descriptions	155
DHCP Option 43 Format.....	156
Default Initiator Name and Data Digest Settings Example	156
Default TCP and Mutual CHAP Settings Example.....	157
Appendix C. OneConnect 10 GbE UCNA Port Speed Specifications	158
Port Speed Negotiation on Non-Mezzanine Cards	158
Port Speed on UCNA Mezzanine Cards	158
Appendix D. Updating Ethernet Firmware	159
Appendix E. Ethtool -S Option Statistics	160

List of Tables

Table 3-1	Static FC and FCoE Driver Parameters	19
Table 3-2	Dynamic FC and FCoE Driver Parameters.....	22
Table 3-3	Ethernet Driver Configuration Parameters	31
Table 3-4	Total VF Counts that are Supported for Various Adapter Configurations	35
Table 3-5	iSCSI Driver Configuration Parameters.....	39
Table 3-6	Interrupt Moderation Policy (im_policy) Settings	41
Table 3-7	VPort sysfs Entries	46
Table 4-1	FC and FCoE Driver Situations and their Resolutions	60
Table 4-2	Ethernet Driver Situations and their Resolutions.....	64
Table 4-3	iSCSI Driver Situations and their Resolutions	65
Table 4-4	LPFC Error Log Messages and their Descriptions	66
Table 4-5	Ethernet Driver Log Messages of Type	143
Table 4-6	iSCSI Driver Error Log Messages	146
Table B-1	DHCP Option 43 (Vendor-Specific Information)	155
Table C-1	OneConnect 10 GbE UCNA Negotiated Speed Specifications	158
Table E-1	Ethtool -S Option Statistics	160
Table E-2	Transmit/Receive Queue Statistics.....	161

1. Introduction

Overview

This *Emulex Drivers for Linux User Manual* provides installing, updating, uninstalling, configuring, and troubleshooting procedures for several types of Emulex-supported drivers for Linux.

This manual is applicable to several versions of Linux drivers, operating systems, firmware, and adapters.

- For a list of supported Emulex drivers for Linux and their associated compatible operations systems, see “Supported Driver Versions” on page 3.
- For supported firmware versions and their latest release, see the “Downloads” page on the Emulex website for the specific adapter.
- For a list of Emulex adapters that are compatible with the standalone driver kits, see the “Downloads” page on the Emulex website for the specific driver.

The Emulex drivers for Linux include:

- Support for the LightPulse™ Fibre Channel (LPFC) family of adapters.
- Support for the OneConnect™ family of universal converged network adapters (UCNAs), including 10 GbE UCNAs and OCe11102-xT 10GBASE-T UCNAs.
- Supported protocols:
 - FC initiator mode and FCoE
 - Small Computer System Interface - Fibre Channel Protocol (SCSI-FCP)
 - FCP-2 (FC-Tape profile, including use of ADISC instead of PLOGI)
 - Ethernet (supported Linux kernel is SLES 11 SPx, supported kernel variants for x86 and x86_64)
 - iSCSI (supported Linux kernel is SLES 11 SPx, supported kernel variants for x86 and x86_64)

- Supported FC topologies: point-to-point, FC Arbitrated Loop (FC-AL), and fabric with auto-topology negotiation
- Supported FC speeds: 1, 2, 4, 8, and 16 Gb/s FC adapters with auto-rate negotiation

Note: The FC drivers do not support a 1 Gb/s rate on 8 or 16 Gb/s FC adapters.

- For FC 8.2.0.x drivers, supports Fibre Channel Security Protocol (FC-SP) Diffie-Hellman Challenge Handshake Authentication Protocol (DHCHAP).
- Tested up to 32 adapter ports
- Dynamic parameter setting using the Emulex OneCommand™ Manager application as part of a master kit:

Enabling graphic user interface (GUI) - based driver configuration, including in-band (FC) and out-of-band (TCP/IP) remote storage area network (SAN)

management capability, diagnostics (loopback and diagnostics dump), virtual port support, personality change and more.

For more information about the OneCommand Manager application, see the latest *OneCommand Manager Application User Manual*, which is available on the Emulex website.

- Support for the common HBA application programming interface (API).
- Batch firmware download capability
- Support for the sysfs (Linux virtual file system) interface. See “Configure Parameters with a Read/Write to sysfs” on page 25.
- Peripheral Component Interconnect (PCI) hot plug support
- Vital product data (VPD) support
- Storage Networking Industry Association - Conformance Testing Program (SNIA-CTP) compliant Storage Management Initiative specification (SMI-S 1.1) provider
- “Linux Tools” hyperlink provided in the Linux portion on the Emulex website: <http://www.emulex.com/files/downloads/linux/tools.html>
- Supports N_Port ID virtualization (NPIV) virtual ports.

NPIV is supported on Emulex Service Level Interface - 4 (SLI-4) OCe10102 series UCNA's. NPIV is also supported on SLI-3, 4, 8, and 16 Gb/s adapters. Emulex enterprise class (5-digit adapter model number) and midrange class (4-digit adapter model number) adapters support SLI-3. For SLI-4 and SLI-3 supported adapters, use the latest recommended firmware for NPIV support.

The FC 8.2.0.x and FC 8.3.5.x drivers support adapters running SLI-2, but NPIV support is unavailable in SLI-2 mode.

- Support for single root I/O virtualization (SR-IOV) for the Ethernet driver with kernel-based virtual machine (KVM) as the hypervisor. SR-IOV is supported on the OneConnect OCx11102-xT UCNA, or later models. For details, see “SR-IOV Configuration” on page 31.

Known Issues

Known issues are defined in the *Emulex Drivers for Linux Release Notes*, which are available on the driver's “Downloads” page on the Emulex website. Also, some known issues are detailed in chapter 4., “Troubleshooting,” on page 60.

Abbreviations

ADISC	Discover Address
AIC	adaptive interrupt coalescing
AMD	Advanced Micro Devices
API	application programming interface
ASIC	application-specific integrated circuit
BIOS	basic input/output system

CentOS	Community Enterprise Operating System
CLI	command line interface
DHCHAP	Diffie-Hellman Challenge Handshake Authentication Protocol
DIMM	dual in-line memory module
DMA	direct memory access
DUD	driver update disc
DUP	driver update package
ELS	extended link service
ETO	extended timeout
FC	Fibre Channel
FC-SP	Fibre Channel Security Protocol
FC-AL	Fibre Channel - Arbitrated Loop
FCoE	Fibre Channel over Ethernet
FCP	Fibre Channel Protocol
FSB	front side bus
GCC	GNU Compiler Collection
Gb/s	gigabits per second
GRO	Generic Receive Offload
GRUB	Grand Unified Bootloader
GUI	graphical user interface
HBA	host bus adapter
INTx	PCIe legacy interrupts, where “x” is variable
IOCB	input/output control block
IOPs	I/O operations per second
IOV	I/O virtualization
IP	Internet Protocol
IPL	initial program load
IQN	iSCSI Qualified Name
iSCSI	internet Small Computer System Interface
ISNS	internet storage name service
KB	kilobyte (1024 bytes)
KVM	kernel-based virtual machine
LACP	Link Aggregation Control Protocol
LDTO	link down timeout
LLC	logical link control
LOM	LAN on motherboard
LPFC	LightPulse Fibre Channel

LRO	large receive offload
LUN	logical unit number
MAC	media access control address
MBR	master boot record
MPIO	multipath I/O
MSI	message signaled interrupts
MSI-X	message signaled interrupts - extended
MSS	maximum segment size (parameter for TCP)
MTU	maximum transmission unit
NAA	network address authority
NCSI	Network Communication Services Interface
NIC	network interface card (or controller)
NPIV	N_Port ID virtualization
OL	Oracle Linux
PCI	Peripheral Component Interconnect
PCIe	Peripheral Component Interconnect Express
PF	PCIe physical function
PID	process ID
PLOGI	port login
POST	power-on self-test
QoS	Quality of Service
RHEL	Red Hat Enterprise Linux
RPI	remote port indicator
RPM	resource package manager
RSCN	registered state change notification
RSS	receive-side scaling
SAN	storage area network
SCSI	Small Computer System Interface
SLES	SUSE Linux Enterprise Server
SLI	Service Level Interface
SMI-S	Storage Management Initiative specification
SMP	symmetric multiprocessing
SNIA-CTP	Storage Networking Industry Association - Conformance Testing Program
SNAP	Sub Network Access Protocol
SNMP	simple network management protocol
SR-IOV	single-root I/O virtualization

TCP	Transmission Control Protocol
UEK	Unbreakable Enterprise Kernel
TSO	TCP Segmentation Offload
UCNA	universal converged network adapter
UDP	User Datagram Protocol
UMC	universal multichannel
VF	PCIe virtual function
VG	virtual guest tagging
VLAN	virtual local area network
VM	virtual machine
VPD	vital product data
WWPN	worldwide port name

2. Installing and Uninstalling

Emulex releases Linux binary RPMs that are digitally signed using the GNU Privacy Guard (GPG) standard. This will allow certification of the contents of the RPMs and verification that the contents have not been modified since they were created by Emulex. The RPMs have been digitally signed by Emulex with a GPG private key that is only held by Emulex. Instructions for creating the Emulex GPG public key file are located on the Emulex website at <http://www.emulex.com/downloads/linux-key.html>

General Installation Requirements

Prior to driver installation, follow these general requirements:

- Install a supported Emulex adapter in the system. Refer to the adapter's installation manual for specific hardware installation instructions.
- Use a supported operating system. The standalone driver kit supports the following distributions:
 - CentOS 5.5, 5.6, and 5.7 (Intel x86, Intel Itanium2, Intel EM64T, AMD64, and PowerPC 64-bit architectures)
 - CentOS 6.0 and 6.1 (Intel x86, Intel EM64T, AMD64, and PowerPC 64-bit architectures)
 - OL 5.5, 5.6, and 5.7 (Intel x86, Intel Itanium2, Intel EM64T, AMD64, and PowerPC 64-bit architectures)
 - OL 6.0 and 6.1 (Intel x86, Intel EM64T, AMD64, and PowerPC 64-bit architectures)
 - OL 5.6 UEK and OL 6.0 UEK (Intel x86 architectures)
 - RHEL 5.5, 5.6, 5.7, and 5.8 (Intel x86, Intel Itanium2, Intel EM64T, AMD64, and PowerPC 64-bit architectures)
 - RHEL 6.x (Intel x86, Intel EM64T, AMD64, and PowerPC 64-bit architectures)
 - SLES 10 SP3 and SP4 (Intel x86, Intel Itanium2, Intel EM64T, AMD64, and PowerPC 64-bit architectures)
 - SLES 11 SPx (Intel x86, Intel Itanium2, Intel EM64T, AMD64, and PowerPC 64-bit architectures)

Binary RPM FC and FCoE Driver Kit

The binary RPM FC and FCoE driver kit contains the following:

- A zipped tar file that includes the driver binary RPMs for a specific driver version and Linux distribution.

Note: Use only officially-released Linux distribution kernels. The binary RPM packages only support officially-released Linux distribution kernels, and do not support pre-release distribution kernels.

- An installation script, `elx_lpfsc_install.sh`, that installs by default the FC and FCoE driver binary RPM that corresponds to the target system's architecture and kernel memory variant.
- A README file that provides a description of the kit structure, its contents, and distribution support scope.
- The driver changelog file.

Installing the Binary RPM FC and FCoE Driver Kit

Note: You must uninstall any FC driver kits that are not part of this distribution. For example, you must uninstall any previous FC and FCoE driver kits that were installed from the Emulex website before installing this driver kit. This installation fails if a previous version of the FC and FCoE driver kit is detected. For more information, see “Uninstalling the Binary RPM FC and FCoE Driver Kit” on page 15.

To install the binary RPM FC and FCoE driver:

1. Download the appropriate driver kit from the Emulex website.
2. Log in as “root” to a terminal, and unpack the tarball:

```
tar xzf elx-lpfc-dd-<Linux distribution version>-<driver version>.tar.gz
```
3. Change to the directory that is extracted:

```
cd elx-lpfc-dd-<Linux distribution version>-<driver version>/
```
4. Run the `elx_lpfsc_install.sh` script without options to install the driver kit:

```
./elx_lpfsc_install.sh
```

Once the `elx_lpfsc_install.sh` script has finished running successfully, the Emulex FC and FCoE driver is loaded, and devices that are connected to the system are accessible.

5. Reboot the system to enable the newly added driver options in the ramdisk. You can also reboot the system later if you want.

Uninstalling the Binary RPM FC and FCoE Driver Kit

Note: You must run the uninstall script that shipped with the version of the driver kit you want to remove.

To uninstall the binary RPM FC and FCoE driver:

1. Log in as “root”.
2. If possible, exit all applications that use FC-attached drives, then unmount the drives. If you cannot exit all applications that use FC-attached drives, the uninstall script works properly, but you must reboot after the uninstallation is complete.
3. Run the `elx_lpfsc_install.sh` script with the “--uninstall” option:

```
./elx_lpfsc_install.sh --uninstall
```

Ethernet Driver Kit

The Ethernet driver kit includes the driver that supports the NIC protocol. The Ethernet driver kit contains the following:

- A zipped tar file that includes the driver binary RPMs for a specific driver version, and for all of the supported Linux distribution kernels.
Note: Use only officially-released Linux distribution kernels. The binary RPM packages only support officially-released Linux distribution kernels, and do not support pre-release distribution kernels.
- An installation script, `elx_net_install.sh`, which installs (by default) the Ethernet driver binary RPM that corresponds to the target system's architecture and kernel memory variant.
- A README file that provides a description of the kit structure, its contents, and distribution support scope.

Installing the Ethernet Driver Kit

Note: Remove any previously installed Ethernet driver kits that were installed from the Emulex website (that is, those that were not part of a distribution's kernel), before proceeding. See "Uninstalling the Ethernet Driver Kit" on page 16 for more information.

To install the Ethernet driver:

1. Download the appropriate driver kit from the Emulex website.
2. Log in as "root" to a terminal, and unpack the tarball:

```
tar xzf elx-be2net-dd-<driver version>.tar.gz
```
3. Change to the directory that is extracted:

```
cd elx-be2net-dd-<driver version>/
```
4. Run the `elx_net_install` script without options to install the driver kit:

```
./elx_net_install.sh
```

Once the `elx_net_install.sh` script has finished running successfully, the Emulex Ethernet driver is loaded, and devices that are connected to the system are accessible.

5. Reboot the system to enable the newly added driver options in the ramdisk. You can also reboot the system later if you want.

Uninstalling the Ethernet Driver Kit

Note: You must run the uninstall script that shipped with the version of the driver kit you want to remove.

To uninstall the Ethernet driver:

1. Log in as "root".

2. If possible, exit all applications that use Ethernet-attached drives, then unmount the drives. If you cannot exit all applications that use Ethernet-attached drives, the uninstall works properly, but you must reboot after the uninstallation is complete.
3. Run the `elx_net_install.sh` script with the “--uninstall” option:

```
./elx_net_install.sh --uninstall
```

iSCSI Driver Kit

The iSCSI driver kit includes the driver that supports the iSCSI protocol. The iSCSI driver kit contains the following:

- A zipped tar file that includes the binary RPMs for a specific driver version, and for all of the supported Linux distribution kernels.

Note: Use only officially-released Linux distribution kernels. The binary RPM packages only support officially-released Linux distribution kernels, and do not support pre-release distribution kernels.

- An installation script, `elx_iscs`
- `i_install.sh`, that installs by default the iSCSI driver binary RPM that corresponds to the target system’s architecture and kernel memory variant.
- A README file that provides a description of the kit structure, its contents, and distribution support.

Installing the iSCSI Driver Kit

Note: Remove any previously installed iSCSI driver kits and/or Application Helper Modules that were installed from the Emulex website (that is, those that were not part of a distribution’s kernel) before proceeding. See “Uninstalling the iSCSI Driver Kit” on page 18 for more information.

To install the iSCSI driver:

1. Download the appropriate driver kit from the Emulex website.
2. Log in as “root” to a terminal, and unpack the tarball:

```
tar xzf elx-be2iscsi-dd-<driver version>.tar.gz
```
3. Change to the directory that is extracted:

```
cd elx-be2iscsi-dd-<driver version>/
```
4. Run the `elx_iscsi_install.sh` script with no options to install the driver kit:

```
./elx_iscsi_install.sh
```
5. Once the `elx_iscsi_install` script has finished running successfully:
 - For an iSCSI boot case, you must reboot the system now to load the driver.
 - For all other iSCSI cases, the Emulex iSCSI driver is loaded, and devices that are connected to the system are accessible. Reboot the system now to enable the newly added driver options in the ramdisk. You can also reboot the system later.

Uninstalling the iSCSI Driver Kit

Note: You must run the uninstall script that shipped with the version of the driver kit you want to remove.

To uninstall the iSCSI driver:

1. Log in as “root”.
2. If possible, exit all applications that use iSCSI-attached drives, then unmount the drives. If you cannot exit all applications that use iSCSI-attached drives, the uninstall works properly, but you must reboot after the uninstallation is complete.
3. Run the `elx_iscsi_install.sh` script with the “--uninstall” option:

```
./elx_iscsi_install.sh --uninstall
```

Booting from a Non-Zero LUN Attached to an Emulex LPFC Adapter

To configure SLES 10 SPx or SLES 11 SPx to boot from an FC-attached disk device other than `/dev/sda`, see the *Emulex Boot for the Fibre Channel Protocol User Manual*.

OneCommand Manager Application

The OneCommand Manager application is a powerful and centralized adapter management suite. It provides discovery, reporting, and management of local and remote adapters from a single console anywhere in the SAN and across platforms. Both a GUI and CLI are provided. This remote configuration capability can be provided by either FC access via host systems on the same FC SAN or by TCP/IP access from IP addresses of remote machines. For instructions on installing and using the OneCommand Manager application, see the latest *OneCommand Manager Application User Manual*, which is available on the Emulex website.

3. Configuration

FC and FCoE Driver Configuration

FC and FCoE Driver Parameters

The FC and FCoE driver parameters determine some aspects of the driver's behavior. There are two main types, static and dynamic. Changes to the static parameters require a driver reload for the change to take effect. Changes to the dynamic parameters take effect immediately. See the following section and "Dynamic FC and FCoE Driver Parameters" on page 22, respectively.

Static FC and FCoE Driver Parameters

Changes to static parameters require a driver reload for the change to take effect. Table 3-1 lists the static FC and FCoE driver parameters.

Table 3-1 Static FC and FCoE Driver Parameters

Parameter	Description	sysfs Visible
lpfc_ack0	When enabled, ACK0 is used for Class 2. The enabled value is 1. The disabled value is 0 (default).	Yes
lpfc_dev_loss_initiator	When enabled, engage the devloss timeout for initiators. The enabled value is 1. The disabled value is 0 (default). Note: This parameter is applicable to FC 8.2.x.x drivers only.	Yes
lpfc_discovery_threads	Specifies the maximum number of ELS commands that can be outstanding for a discovery. Note: The lpfc_discovery_threads parameter defaults to a value of 64 for private loop topologies regardless of the configured value. If there are multiple ports configured on the host the value of 64 is only used for those ports that are connected in a private loop topology. The configured value is used for all other ports. The minimum value is 1. The maximum value is 64. The default value is 32.	No
lpfc_enable_da_id	When enabled, the FC and FCoE driver issues a DA_ID CT command to the fabric when VPorts log out of the fabric. The enabled value is 1. The disabled value is 0 (default).	No
lpfc_enable_hba_heartbeat	When enabled, the heartbeat logic in the FC and FCoE driver is able to detect whether the adapter is functional. If the heartbeat logic detects the adapter is not functional, the driver will shut down the adapter. The enabled value is 1 (default). The disabled value is 0.	Yes

Table 3-1 Static FC and FCoE Driver Parameters (Continued)

Parameter	Description	sysfs Visible
lpfc_enable_hba_reset	When enabled, the FC and FCoE drivers can pass resets to the adapter. This is typically used for debugging purposes. The enabled value is 1 (default). The disabled value is 0.	Yes
lpfc_enable_npiv	When enabled, the FC and FCoE driver can use NPIV to create VPorts (if supported by the fabric). The enabled value is 1 (default). The disabled value is 0.	Yes
lpfc_fcp_class	Specifies either FC Class 2 or 3 for FCP data transmission. For Class 2, the value is 2. For Class 3, the value is 3 (default).	Yes
lpfc_fcp_eq_count	Note: This parameter is deprecated in 8.3.7.x drivers. For OneConnect adapters and LPe16000 adapters, specifies the number of fast-path FCP event queues, if available. The minimum value is 1. The maximum value is 8. The default value is 4. Note: For LPe12000, LPe11000, and LP11000 adapters, this parameter is not applicable and has no effect.	Yes
lpfc_fcp_wq_count	Note: This parameter is deprecated in 8.3.7.x drivers. For OneConnect adapters and LPe16000 adapters, specifies the number of fast-path FCP work queues, if available. The minimum value is 1. The maximum value is 32. The default value is 4. Note: For LPe12000, LPe11000, and LP11000 adapters, this parameter is not applicable and has no effect.	Yes
lpfc_hba_queue_depth	Specifies the maximum number of FCP commands that can queue to an Emulex adapter. The minimum value is 32. The maximum value is 8192 (default).	Yes
lpfc_lun_queue_depth	Specifies the default maximum number of commands sent to a single logical unit (disk drive). The minimum value is 1. The maximum value is 128. The default value is 30.	Yes
lpfc_max_luns	Specifies the highest available LUN ID that is valid, per target. For example, a value of 19 means that LUN IDs from 0 to 19 are valid for the target. The SCSI layer scans each target until it reaches this specified LUN ID. The minimum value is 0. The maximum value is 65535. The default value is 255.	Yes
lpfc_max_scscmpl_time	Uses command completion time to control queue depth. The units are in milliseconds. The minimum value is 0 (default). The maximum value is 6000.	Yes
lpfc_multi_ring_rctl	When lpfc_multi_ring_support is enabled, identifies the routing control (R_CTL) for the additional ring configuration. The minimum value is 1. The maximum value is 255. The default value is 4.	Yes

Table 3-1 Static FC and FCoE Driver Parameters (Continued)

Parameter	Description	sysfs Visible
lpfc_multi_ring_support	Determines the number of primary SLI rings over which to spread IOCB entries. The minimum value is 1 (default). The maximum value is 2.	Yes
lpfc_multi_ring_type	When lpfc_multi_ring_support is enabled, identifies the TYPE of the additional ring configuration. The minimum value is 1. The maximum value is 255. The default value is 5 (LLC/SNAP).	Yes
lpfc_restrict_login	When enabled, restricts VPorts login to remote initiators. The enabled value is 1 (default). The disabled value is 0.	No
lpfc_scan_down	When enabled, selects the “scan down” method (scanning the AL_PA from high to low) to assign a SCSI ID. The enabled value is 1 (default). The disabled value is 0.	Yes
lpfc_sg_seg_cnt	Controls the scatter/gather maximum segment count passed to the FC and FCoE driver. This variable is applicable per SCSI command. For LPe12000, LPe11000, and LP11000 adapters, the minimum value is 64 (default), and the maximum value is 4096. For OneConnect and LPe16000 adapters, the minimum value is 64 (default), and the maximum value is 510.	Yes (sg_table size)
lpfc_sli_mode	For LPe12000, LPe11000, and LP11000 adapters, this parameter allows you to force the SLI mode requested by the adapter driver. The possible values are: <ul style="list-style-type: none"> 0 = Auto-select (default) 2 = SLI-2 3 = SLI-3 Note: For OneConnect and LPe16000 adapters, this parameter is not applicable and has no effect.	No
lpfc_use_msi	When enabled, determines whether the driver uses MSI or MSI-X. <ul style="list-style-type: none"> 0 = MSI disabled; INTx mode is used (default for FC 8.2.x.x drivers). 1 = MSI; allows a maximum of 32 interrupts. 2 = MSI-X; allows a maximum of 2048 interrupts (default for FC 8.3.x.x drivers). 	Yes

Dynamic FC and FCoE Driver Parameters

Changes to the dynamic parameters take effect immediately. All LPFC dynamic parameters are read/write using sysfs. Table 3-2 lists the dynamic FC and FCoE driver parameters.

Table 3-2 Dynamic FC and FCoE Driver Parameters

Parameter	Description
lpfc_cr_count	For LPe12000, LPe11000, and LP11000 adapters, this parameter determines the value for I/O coalescing for lpfc_cr_count outstanding commands. The minimum value is 1 (default). The maximum value is 255. Note: For OneConnect and LPe16000 adapters, this parameter is not applicable and has no effect.
lpfc_cr_delay	For LPe12000, LPe11000, and LP11000 adapters, this parameter determines the value for I/O coalescing for lpfc_cr_delay (milliseconds) outstanding commands. The minimum value is 0 (default). The maximum value is 63. Note: For OneConnect and LPe16000 adapters, this parameter is not applicable and has no effect.
lpfc_devloss_tmo	Specifies the number of seconds to hold an I/O error when a device disappears. The minimum value is 0. The maximum value is 255. The default value is 30.
lpfc_enable_auth	Specifies whether DHCHAP support is enabled. When set to 1, DHCHAP is enabled. When set to 0, DHCHAP is disabled. Note: This property requires a link reset to activate. Note: This parameter is applicable to FC 8.2.x.x drivers only.
lpfc_fcp_io_channel	For LPe16000 adapters using 8.3.7.x drivers, defines the number of I/O channels supported by the driver. For more information, see “LPFC Driver Performance Tuning” on page 51. The default value is 4 I/O channels. For LPe12000, LPe11000, and LP11000 adapters, and for LPe16000 adapters using 8.2.x.x drivers, this parameter is not applicable and has no effect.
lpfc_fcp_io_sched	For LPe16000 adapters using 8.3.7.x drivers, determines which algorithm to use when scheduling an FCoE I/O to an I/O channel. For more information, see “LPFC Driver Performance Tuning” on page 51. The default value is 0, configuration by round-robin scheduling. A value of 1 sets configuration to CPU scheduling. For LPe12000, LPe11000, and LP11000 adapters, and for LPe16000 adapters using 8.2.x.x drivers, this parameter is not applicable and has no effect.
lpfc_fcp_imax	For OneConnect and LPe16000 adapters, specifies the maximum number of fast-path FCP interrupts per second. The minimum value is 636. The maximum value is 651042. The default value is 10000. Note: For LPe12000, LPe11000, and LP11000 adapters, this parameter is not applicable and has no effect.

Table 3-2 Dynamic FC and FCoE Driver Parameters (Continued)

Parameter	Description
lpfc_fdmi_on	Specifies the type of FDMI support. The enabled values are 1 or 2 depending on the type needed. The disabled value is 0 (default).
lpfc_link_speed	Specifies the FC link speed. The possible values are: <ul style="list-style-type: none"> 0 = Auto-select (default) 1 = 1 Gb/s 2 = 2 Gb/s 4 = 4 Gb/s 8 = 8 Gb/s 16 = 16 Gb/s Note: This parameter does not affect FCoE 10 Gb/s adapters.
lpfc_log_verbose	Specifies the log verbosity level of the messages posted by the driver. Extra activity logging (bit mask). The minimum value is 0x0 (default). The maximum value is 0xFFFF.
lpfc_nodev_tmo (deprecated)	Note: This is a deprecated parameter and lpfc_devloss_tmo should be used instead. This parameter will not work if you altered lpfc_devloss_tmo. Specifies the number of seconds to hold an I/O error when a device disappears. The minimum value is 1. The maximum value is 255. The default value is 30.
lpfc_pci_max_read	Specifies the maximum DMA read byte count. The possible values are 512, 1024, 2048 (default), and 4096.
lpfc_poll	Sets the FCP ring polling mode control. The possible values are: <ul style="list-style-type: none"> 0 = no polling (default) 1 = poll with interrupts enabled 3 = poll and disable FCP ring interrupts
lpfc_poll_tmo	Specifies the number of milliseconds that the driver waits between polling FCP ring interrupts. The minimum value is 1. The maximum value is 255. The default value is 10.
lpfc_topology	For FC adapters, this parameter sets the link topology. The possible values are: <ul style="list-style-type: none"> 0x0 = loop first; if loop fails, then point-to-point (default) 0x2 = point-to-point only 0x4 = loop only 0x6 = point-to-point first; if point-to-point fails, then loop Note: For FCoE adapters, this parameter is not applicable and has no effect.
lpfc_use_adisc	When enabled, an ADISC is sent instead of a PLOGI for device discovery or RSCN. The enabled value is 1. The disabled value is 0.

Configuring FC and FCoE Driver Parameters

You can configure the FC and FCoE driver parameters by using:

- the modprobe Linux program for temporary configuration
- the modprobe.conf file (FC 8.2.x.x drivers) or the lpfc.conf file (FC 8.3.x.x drivers) for persistent configuration
- the sysfs interface (to view and modify parameters after loading the FC and FCoE driver)
- the OneCommand Manager application (See the *OneCommand Manager Application User Manual* for more information.)

Note: FC and FCoE driver parameter changes made using modprobe.conf or the OneCommand Manager application persist if the FC and FCoE driver is uninstalled. To return to the default settings, you must reset them in the modprobe.conf file and reload the driver.

Temporary Configuration with modprobe

When you manually load the FC and FCoE driver as a module using the modprobe command, and you change one or more driver parameter values in the command line, the configuration is temporary. These changes are considered temporary because they are valid for the current session only or until the FC and FCoE driver is unloaded.

Modprobe uses the modprobe.conf file, but parameters passed to it using the command line override the parameters in the modprobe.conf file. Values can be expressed in hexadecimal or decimal notation.

If you want to temporarily set `lun_queue_depth` to 20 (default is 30) for all HBAs in your system, load the FC and FCoE driver with the following command:

```
modprobe lpfc lpfc_lun_queue_depth=20
```

Persistent Configuration with modprobe.conf or lpfc.conf

For FC 8.2.0.x drivers, to make the FC and FCoE driver parameters persist across module loads and reboots, modify the `/etc/modprobe.conf` file. The FC and FCoE driver parameters are specified in `/etc/modprobe.conf` via the “options” command. For example, the following command sets the verbose flag:

```
options lpfc lpfc_log_verbose=0xffffffff
```

For FC 8.3.5.x drivers, to make the FC and FCoE driver parameters persist across module loads and reboots, follow these two steps:

1. In the `/etc/modprobe.d` directory, create a file with the driver name “`lpfc.conf`”.
2. In `/etc/modprobe.d/lpfc.conf`, use the “options” command to add the appropriate FC and FCoE driver parameters and their desired values. For example, adding the following command to the `lpfc.conf` file sets the verbose flag:

```
options lpfc lpfc_log_verbose=0x3ffff
```

If driver parameters are modified in the `modprobe.conf` file or added to the `lpfc.conf` file, the FC and FCoE driver must be reloaded for the parameters to take effect. Also, a

new ramdisk image is required if you want the changes to take effect in the next boot. See “Creating a New Ramdisk Image” on page 26.

If the same parameter is specified on the modprobe command line and also in the modprobe.conf or lpfc.conf file, then the value specified in the modprobe command line takes precedence.

Configure Parameters with a Read/Write to sysfs

Sysfs is a virtual file system that exposes the structure of the system. It also includes interfaces to driver parameters through which the FC and FCoE driver parameters can be viewed and modified. Since these interfaces are available only after driver load, only dynamic FC and FCoE driver parameters can be changed. However, both static and dynamic FC and FCoE driver parameters can be read through sysfs.

Note: Sysfs changes exist only during driver load and are lost when the FC and FCoE driver is unloaded or the system is rebooted.

Viewing Parameters with sysfs

The sysfs file system is mounted and available as /sys. You must first identify the scsi_host that represents the adapter for which you want to modify the FC and FCoE driver parameters. All scsi_hosts bound to the FC and FCoE driver can be viewed with the following command:

```
ls -d /sys/bus/pci/drivers/lpfc/*/host*
```

Assuming you are interested in adapter scsi_host 7, you can list the FC and FCoE driver parameters for this particular adapter as:

```
ls -l /sys/class/scsi_host/host7/lpfc*
```

An example output follows:

```
-r--r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/lpfc_ack0
-r--r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/lpfc_fcp_class
-rw-r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/
lpfc_fdmi_on
-r--r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/lpfc_link_speed
-rw-r--r-- 1 root root 4096 Feb 28 15:34 /sys/class/scsi_host/host7/lpfc_log_verbose
-r--r--r-- 1 root root 4096 Feb 28 17:03
/sys/class/scsi_host/host7/lpfc_lun_queue_depth
-rw-r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/lpfc_max_luns
-rw-r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/lpfc_nodev_tmo
-rw-r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/lpfc_scan_down
-r--r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/lpfc_topology
-rw-r--r-- 1 root root 4096 Feb 28 17:03 /sys/class/scsi_host/host7/lpfc_use_adisc
```

Temporary Configuring Parameters with sysfs

In the previous example, notice that the FC and FCoE driver parameters are available as files. Reading a file displays the current value of a driver parameter. If the permissions allow it, you can write a value to the file and it will take effect immediately.

Reading the lpfc_log_verbose file may show that its value is “0”:

```
cat /sys/class/scsi_host/host7/lpfc_log_verbose
0
```

To modify the `lpfc_log_verbose` value to `0xffffffff`:

```
echo 0xffffffff > /sys/class/scsi_host/host7/lpfc_log_verbose
```

Reading the `lpfc_log_verbose` file now shows a value of `0xffff`:

```
cat /sys/class/scsi_host/host7/lpfc_log_verbose
0xffffffff
```

Creating a New Ramdisk Image

The `lpfc-install` script creates a ramdisk image containing the FC and FCoE driver for the currently running kernel.

Note: You must create a new ramdisk image whenever the LPFC options in `/etc/modprobe.conf` are changed and you want the change to take effect on the next reboot.

To create a new initial ramdisk image for inbox FC and FCoE drivers and installed binary RPM FC driver kits:

- For SLES 10 PPC64 architecture distributions, type

```
mkinitrd -k vmlinux -i initrd
```
- For SLES 10 non-PPC64 architecture distributions and SLES11 SPx distributions, type

```
mkinitrd -k vmlinuz -i initrd
```
- For RHEL 5.x distributions, type

```
mkinitrd -f /boot/initrd-<kernel-version>.img  
<kernel-version>
```
- For RHEL 6.x distributions, type

```
dracut -f /boot/initramfs-<kernel-version>.img  
<kernel-version>
```

Dynamically Recognizing LUNs and Targets (Using scan)

The FC and FCoE driver enables you to dynamically recognize LUNs and targets without unloading or reloading the LPFC module and without resetting the adapter.

To rescan an adapter's targets with `sysfs`, given the adapter's host number (in this example, 3), type

```
echo "- - -" > /sys/class/scsi_host/host3/scan
```

To limit the rescan to a particular target, given the adapter's host number (in this example, 3) and the target number (in this example, 2), type

```
echo "- 2 -" > /sys/class/scsi_host/host3/scan
```

You can also use the Emulex `lun_scan` script in the `/usr/sbin/lpfc` directory.

Persistent Naming

The generic device manager for the Linux kernel is “udev”, which primarily manages device nodes in the /dev directory.

Using udev to Discover Logical to Physical Mappings for sd Devices

In Linux, the driver for SCSI disk drives is “sd”. A disk device name has an sd prefix. Persistent names for sd devices are provided in the /dev/disk/by-id directory. To find the persistent udev name for the disk, which is currently “sdc”, type

```
cd /dev/disk/by-id
ls -l | grep sdc
```

The sample output is:

```
lrwxrwxrwx 1 root root 9 2006-08-01 19:08 scsi-32000000c5005d6e6 -> ../../sdc
```

In the previous example, the disk has no partitions. If the disk had two partitions, the output would look like the following:

```
lrwxrwxrwx 1 root root 9 2006-08-01 19:08 scsi-32000000c5005d6e6 -> ../../sdc
lrwxrwxrwx 1 root root 10 2006-08-01 19:08 scsi-32000000c5005d6e6-part1 -> ../../sdc1
lrwxrwxrwx 1 root root 10 2006-08-01 19:08 scsi-32000000c5005d6e6-part2 -> ../../sdc2
```

Configuring the System to Boot Using Persistent Names

For SLES 10 SPx and SLES 11 SPx

Note: SLES 10 SPx and SLES 11 SPx are configured by default with udev to provide persistent names for hard disks, including FC-attached disks.

To use a persistent name for a boot device with SLES 10 SPx and SLES 11 SPx:

1. In /boot/grub/menu.lst, find the kernel line for the default boot. For example:

```
kernel /boot/vmlinuz root=/dev/sda2 vga=0x314
```
2. Find the persistent name for the root partition (following “root=” on the kernel line) by using the instructions in “Using udev to Discover Logical to Physical Mappings for sd Devices” on page 27.
3. In the same file, /boot/grub/menu.lst, replace the text after “root=” with the partition’s persistent name. For example:

```
kernel /boot/vmlinuz root=/dev/disk/by-id/scsi-32000000c5005d6e6-part2 vga=0x314
```
4. Change any mounts listed in /etc/fstab that refer to this root partition by either its /dev/sd name or a file system label to use the persistent name as well.

For RHEL 5.x and RHEL 6.x

To use a persistent name for a boot device with RHEL 5.x and RHEL 6.x:

1. In /boot/grub/grub.conf, find the kernel line for the default boot. For example:

```
kernel /boot/vmlinuz -<kernel version> ro root=/dev/sda2
```

2. Find the persistent name for the root partition (following “root=” on the kernel line) by using the instructions in “Using udev to Discover Logical to Physical Mappings for sd Devices” on page 27.
3. In the same file, /boot/grub/menu.lst, replace the text after “root=” with the partition's persistent name. For example:

```
kernel /boot/vmlinuz -<kernel version> ro
root=/dev/disk/by-id/scsi-32000000c5005d6e6-part2
```
4. Change any mounts listed in /etc/fstab which refer to this root partition by either its /dev/sd name or a file system label to use the persistent name as well.

Using udev with st Devices

In Linux, the driver for SCSI tape drives is “st”. A tape device name has an “st” prefix. The udev rules for tape devices are the same as for disk devices. There must be a unique ID that persists across initiator reboots and persists regardless of discovery order.

You must consider whether the tape device is an FC tape device or an FC-SCSI tape device (in which there are multiple SCSI tape devices that reside behind an FC controller). If it is an FC tape device, then the WWPN is unique and can be used to create the persistent name. In this case, the `scsi_id` command should return this as the unique identifier with a single digit prefix. If the FC controller has multiple SCSI tape devices behind it, the WWPN is not unique, and the persistent name must use multiple information elements to build the unique ID. The following are examples of each scenario.

FC Tape Device Examples

The following examples use the `scsi_id` command to retrieve and generate a unique SCSI identifier:

```
scsi_id [options]
```

For these examples, the following [options] are used:

- `-g` Treats the device as white listed. It is needed on the command line or in the `scsi_id.config` file for the `scsi_id` command to generate any output. In the examples, the `-g` option is needed on the command line because the vendor and model for this tape device were not in the `/etc/scsi_id.config` file.
- `-s` Generates an id for the sysfs-device. Note that “-s” is an invalid option for `scsi_id` version 147.

Note: Since the [options] can vary depending on the version of the `scsi_id` command, see the `scsi_id` man page on your system for the correct and complete list of the [options].

The following example is an FC tape device using the SCSI generic driver (sg) rather than the SCSI tape driver. The value returned has a leading prefix of 3, which is the NAA type. The remaining digits represent the FC controller's WWPN.

```
scsi_id -g -s /sys/class/scsi_generic/sg0
```

```
350060b000029b592
```

The following example is an FC tape device using the SCSI tape driver. The value returned is the same as the previous example.

```
scsi_id -g -s /sys/class/scsi_tape/nst0
350060b000029b592
```

The following example uses a different FC tape vendor. Notice that the value returned is similar to the previous examples, with respect to the leading digit and the WWPN.

```
/sbin/scsi_id -g -s sys/class/scsi_tape/nst0
35005076300015101
```

FC-SCSI Tape Device Example

The following is an example of a FC controller with multiple SCSI tape devices behind it (FC-SCSI tape device). When the Emulex driver is loaded, the SCSI mid-level discovers the SCSI tape devices as follows:

```
scsi scan: INQUIRY to host 14 channel 0 id 0 lun 0
scsi: unknown device type 12
Vendor: ADIC      Model: SNC 4000      Rev: 42d4
Type:   RAID      ANSI SCSI revision: 03
Attached scsi generic sg5 at scsi14, channel 0, id 0, lun 0, type 12
scsi scan: INQUIRY to host 14 channel 0 id 0 lun 1
Vendor: ADIC      Model: Scalar 24     Rev: 227A
Type:   Medium Changer ANSI SCSI revision: 02
Attached scsi generic sg6 at scsi14, channel 0, id 0, lun 1, type 8
scsi scan: INQUIRY to host 14 channel 0 id 0 lun 2
Vendor: IBM       Model: ULTRIUM-TD2   Rev: 38D0
Type:   Sequential-Access ANSI SCSI revision: 03
Attached scsi tape st0 at scsi14, channel 0, id 0, lun 2
st0: try direct i/o: yes (alignment 512 B), max page reachable by HBA 4503599627370495
Attached scsi generic sg7 at scsi14, channel 0, id 0, lun 2, type 1
scsi scan: INQUIRY to host 14 channel 0 id 0 lun 3
Vendor: IBM       Model: ULTRIUM-TD2   Rev: 38D0
Type:   Sequential-Access ANSI SCSI revision: 03
Attached scsi tape st1 at scsi14, channel 0, id 0, lun 3
st1: try direct i/o: yes (alignment 512 B), max page reachable by HBA 4503599627370495
Attached scsi generic sg8 at scsi14, channel 0, id 0, lun 3, type 1
```

This log output shows a controller at LUN 0, the medium changer at LUN 1, and two SCSI tape devices at LUNs 2 and 3.

The following example is the result of a `scsi_id` call:

```
scsi_id -g -s /sys/class/scsi_tape/nst0
1IBM      ULTRIUM-TD2      1110133831
scsi_id -g -s /sys/class/scsi_tape/nst1
1IBM      ULTRIUM-TD2      1110133994
```

Notice that the unique ID comprises three values with space delimiters. A udev rule must have a unique ID for the device, so that all three parts of this returned string are required. To do this, use the following command:

```
scsi_id -u -g -s /sys/class/scsi_tape/nst0
1IBM_____ULTRIUM-TD2_____1110133831
scsi_id -u -g -s /sys/class/scsi_tape/nst1
1IBM_____ULTRIUM-TD2_____1110133994
```

Creating the udev Persistent Name for SCSI Tape Device

After you know the SCSI ID call needed to extract a unique ID, use the same process to create a udev persistent name for a SCSI tape device as on a SCSI disk device.

The rule for the FC tape device is

```
BUS="scsi", SYSFS{vendor}="HP", SYSFS{model}="ULTRIUM 3-SCSI",
PROGRAM="/sbin/scsi_id -p 0x83 -u -g -s
/sys/class/scsi_tape/nst%n", RESULT="350060b000029b592",
SYMLINK="fc_lun_st%n"
```

The rule for the FC-SCSI tape device is

```
BUS="scsi", SYSFS{vendor}="IBM", SYSFS{model}="ULTRIUM-TD2",
PROGRAM="/sbin/scsi_id -p 0x83 -u -g -s
/sys/class/scsi_tape/nst%n", RESULT="1IBM_____ULTRIUM-TD2_____111013
3831", SYMLINK="fc_lun_st%n"
BUS="scsi", RESULT="1IBM_____ULTRIUM-TD2_____1110133994",
SYMLINK="fc_lun_st%n"
```

Create a new file named `/etc/udev/rules.d/45-local.rules` and put the appropriate rule in it. Then run “udevtrigger” to reload the udev rules, and the output of the rules will follow:

```
udevtrigger
ls -al /dev/fc*
lrwxrwxrwx 1 root root 3 Apr  7 15:03 fc_lun_st0 -> st0
lrwxrwxrwx 1 root root 3 Apr  7 15:03 fc_lun_st1 -> st1
```

Persistent Naming References

See the following references for more information on persistent naming:

- <http://www.reactivated.net/udevrules.php> by Daniel Drake (dsd)
- http://kernel.org/pub/linux/utils/kernel/hotplug/udev_vs_devfs by Greg Kroah-Hartman

Ethernet Driver Configuration

Ethernet Driver Configuration Parameters

The Ethernet driver supports the user-configurable parameters described in Table 3-3.

Table 3-3 Ethernet Driver Configuration Parameters

Parameter	Description
num_vfs	In systems supporting SR-IOV, when IOV is enabled, this parameter indicates the number of VFs to be enabled per PF. For configuring SR-IOV, see “SR-IOV Configuration” on page 31. The default value is 0 (SR-IOV is not enabled). The possible values are 0 to 32.
rss_on_mc	Enables receive-side scaling (RSS) on multichannel functions that have the capability. The default value is 0 (disabled). The enabled value is 1.
rx_frag_size	The size of fragments used to DMA received data. The possible values are 2048 (default), 4096 and 8192.
num_rings	The number of rings on which packets can be transmitted and received can be varied using this parameter. The default value is 8. The maximum supported value is 12.

You can configure Linux to automatically load the driver with any of these options after each reboot. To do so, add a line to `/etc/modprobe.conf` with the required options. For example, to load the driver with the fragment size of 4096 and create eight VFs per PF in an SR-IOV-capable system, add the following line:

```
options be2net rx_frag_size=4096 num_vfs=8
```

VLAN Support on UMC

In universal multichannel (UMC) mode, an LPVID for each channel configured in the BIOS is transparently added in the transmit path and removed in the receive path by the UCNA. When a VLAN is configured in the host using `vconfig` on any of the functions, the host VLAN ID overrides the corresponding LPVID channel configured in the BIOS. Up to 15 VLAN IDs can be configured in the host for each PF using `vconfig`. The VLAN IDs configured in the host should be different from the channel VLAN IDs configured in the BIOS. For additional information on UMC, refer to the *Emulex Universal Multichannel Reference Guide*.

Note: You cannot run LACP when UMC is enabled.

SR-IOV Configuration

Introduction

This section contains requirements and instructions to use SR-IOV with the following host operating systems:

- Windows 2012 - 64-bit Hyper-V
- Red Hat Enterprise Linux 6.3 - 64-bit KVM
- Red Hat Enterprise Linux 6.4- 64-bit KVM
- SuSE Linux Enterprise Server 11 SP1 - 64-bit Xen
- SuSE Linux Enterprise Server 11 SP2 - 64-bit Xen
- VMware ESXi 5.1 Update 1 (IBM Customized Image)
- SLES and RHEL supported with Xen

These environments support capabilities of OneConnect to enable multiple PCIe virtual functions (VFs) for a PCIe physical function (PF). Each of these VFs can be assigned to virtual machines (VMs). A VF enables the guest operating system direct access to OneConnect, such that guest performance is not limited by the overheads of the hypervisor.

With SR-IOV, VMs directly drive I/Os on the NIC. Therefore, SR-IOV has the following advantages over traditional virtualized I/O:

- Reduced CPU utilization
- Reduced latency
- Increased scalability
- Improved device performance for virtual guests

Known issues include:

- Any attempt to disable SR-IOV by the driver in hypervisor when VFs are assigned to VMs leads to undefined behavior. This known issue needs to be corrected in the kernel. As a result of this issue, once an Ethernet driver is loaded with a non-zero value for `num_vfs`, the driver can never be unloaded. The only way to reload the driver with a different `num_vfs` value is to reboot the system.
- The Kdump (kernel dump) feature is not supported when SR-IOV is enabled.

Setting Up SR-IOV

Prerequisites

To set up SR-IOV on your system, you need the following:

- A server/blade with an SR-IOV-capable motherboard BIOS.

Note: Currently, SR-IOV is supported only in non-multichannel setups.

Note: Configuration mechanisms for parameters such as MAC address, VLAN and QoS for VF are supported in RHEL6, RHEL 6.1, and SLES 11.3 distributions.

- OneConnect OCx11102-xT UCNA, or later UCNA versions.
- If a KVM hypervisor is installed, it must contain the `qemu-kvm` packages.

Procedure

Depending on your system, perform one or more of the following tasks to set up your BIOS. For more information, see the manufacturer's instructions for your system.

- a. Enable SR-IOV in the system BIOS. This supports SR-IOV functionality.
- b. Enable Intel Virtualization Technology support for Direct I/O VT-d.

You can use the PXESelect utility or the UEFI (HII) utility to set up SR-IOV.

To enable SR-IOV in the OneConnect firmware using the UEFIBoot utility, see the *Boot for NIC, iSCSI, and FCoE Protocols User Manual*.

To enable SR-IOV in the OneConnect firmware using the PXESelect menu, refer to the instructions in "Enabling SR-IOV Using the PXESelect Menu" below.

Enabling SR-IOV Using the PXESelect Menu

1. Enter the PXE Select utility menu by pressing <Ctrl> + <P> when the PXE BIOS banner appears during system boot.

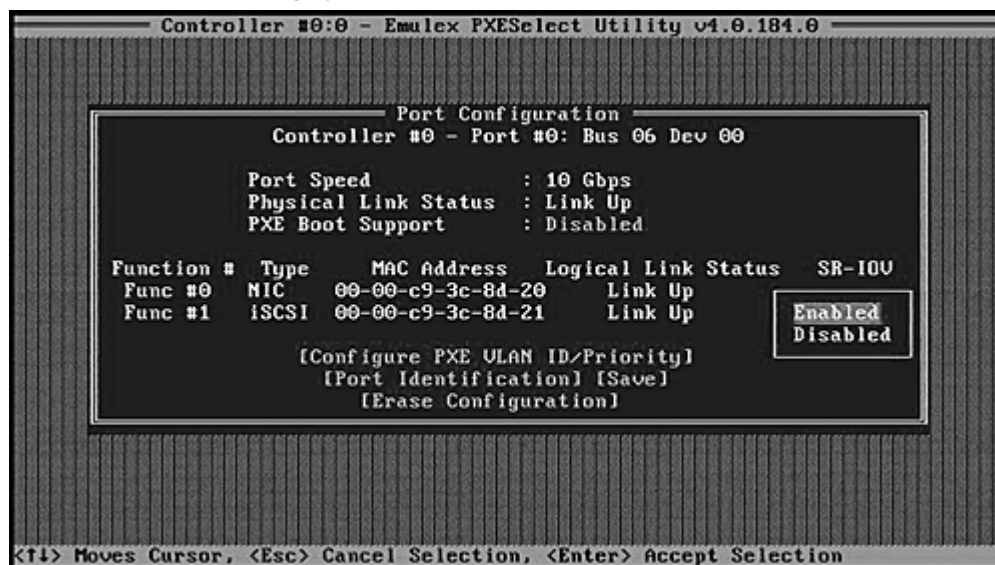


Figure 3-1 PXESelect Utility, SR-IOV Enabled

2. Install the required Linux operating system that serves as a hypervisor.
3. Update the /boot/grub/menu.lst file to include the following command line load parameter for the Linux kernel:

```
intel_iommu=on
```

4. Install the appropriate version of the Emulex Ethernet driver that supports SR-IOV for the operating system version that you are using.
5. Reboot the server for the new changes to become operational.
6. Use the "lspci -vvv" output command to check if SR-IOV is properly enabled:

```
lspci -vvv
```

This command returns an SR-IOV capability section for each OneConnect Ethernet PF. The Initial VFs and Total VFs should be non-zero. Make a note of the lspci

command output in the hypervisor. This output is needed in step 11 on page 36. Specifically make note of the pci-id of the VFs that have been created.

Example of the SR-IOV capabilities section output of the PF with SR-IOV enabled:

```
Capabilities: [180] Single Root I/O Virtualization (SR-IOV)
IOVCap: Migration-, Interrupt Message Number: 000
IOVCtl: Enable- Migration- Interrupt- MSE- ARIHierarchy-
IOVSta: Migration-
Initial VFs: 32, Total VFs: 32, Number of VFs: 0, Function Dependency
Link: 00
VF offset: 0, stride: 1, Device ID: 0710
Supported Page Size: 00000557, System Page Size: 00000001
Region 0: Memory at 0000000000000000 (64-bit, non-prefetchable)
VF Migration: offset: 00000000, BIR: 0
```

Example of the SR-IOV capabilities section output of the PF with SR-IOV disabled:

```
Capabilities: [180] Single Root I/O Virtualization (SR-IOV)
IOVCap: Migration-, Interrupt Message Number: 000
IOVCtl: Enable- Migration- Interrupt- MSE- ARIHierarchy-
IOVSta: Migration-
Initial VFs: 0, Total VFs: 0, Number of VFs: 0, Function Dependency
Link: 00
VF offset: 0, stride: 1, Device ID: 0710
Supported Page Size: 00000557, System Page Size: 00000001
Region 0: Memory at 0000000000000000 (64-bit, non-prefetchable)
VF Migration: offset: 00000000, BIR: 0
```

7. Create a virtual machine using the Virtual Machine Manager utility and install the RHEL 6.x operating system on the VM.
8. Shut down the VM.
9. Reload the newly installed Ethernet driver in the host (hypervisor) with the number of VFs/PF as module parameter:

```
rmmod be2net
modprobe be2net num_vfs=X
```

Where “X” is the number of VFs per PF. The possible values are 0 to 32 per physical port. The default value is 0 (SR-IOV is not enabled).

The total number of VFs can be distributed among available ports as required, but each port has a maximum of 32 VFs. Table 3-4 on page 35 lists the total number of VF counts that are supported for various adapter configurations.

Notes:

- On 4-port adapters, VFs can be configured only for ports 0 and 1, not for ports 2 or 3.
- VFs are supported for network functions only; they are not supported for storage functions.

Table 3-4 Total VF Counts that are Supported for Various Adapter Configurations

Adapter Configuration	Number of Ports	PF Count	VF Count per ASIC	Comments
2-port 10Gb NIC, stand-alone card	2	2	60	Per port VF count is 30. P0/P1: 30/30
2-port 10G NIC, stand-alone, NCSI-enabled IPL	2	2	59	VF count per port can be P0/P1: 30/29 or P0/P1: 29/30.
2-port 10Gb, HP/IBM LOM/Mezzanine	2	2	60	Per port VF count is 30. P0/P1: 30/30
2-port 10Gb, HP/IBM LOM/Mezzanine, NCSI-enabled IPL	2	2	59	VF count per port can be P0/P1: 30/29 or P0/P1: 29/30.
2-port 10Gb + 2-port 1Gb (4-port), IBM LOM/Mezzanine	4	4	56	Per port VF count is 28. VFs are not supported on 1G ports.
2-port 10Gb + 2-port 1Gb (4-port), IBM LOM/Mezzanine	4	4	55	Per port VF count can be P0/P1: 28/27 or P0/P1: 27/28. VFs are not supported on 1G ports.
2-port 10G HP, Flex-10 enabled (All IPLs)	2	4/6/8	48	Per port VF count is 24. VFs will be distributed across the PFs in multiples of 8.

10. Detach VFs from the host before adding them to the guest.

Example 1 (using only one VF per physical port):

RHEL6.x KVM: - View Emulex PCI devices with the lspci command:

```
# lspci | grep Emulex      16:00.0 Ethernet controller: Emulex
Corporation OneConnect 10Gb NIC (be3) (rev 03) <-- LOM Port 0 (Function 0)
16:00.1 Ethernet controller: Emulex Corporation OneConnect 10Gb NIC (be3)
(rev 03) <-- LOM Port 1 (Function 1)      16:04.0 Ethernet controller: Emulex
Corporation OneConnect 10Gb NIC (be3) (rev 03) <-- VF LOM Port 0 (Function
0)      16:08.0 Ethernet controller: Emulex Corporation OneConnect 10Gb NIC
(be3) (rev 03) <-- VF LOM Port 1 (Function 1)
```

Detach desired VFs:

```
# virsh nodedev-dettach pci_0000_16_04_0      Device pci_0000_16_04_0
detached <-- VF LOM Port 0 (Function 0)      # virsh nodedev-dettach
pci_0000_16_08_0      Device pci_0000_16_08_0 detached <-- VF LOM Port 1
(Function 1)
```

Example 2 (using only one VF per physical port):

SLES11.x Xen: - View Emulex PCI devices with the lspci command:

```
# lspci | grep Emulex      16:00.0 Ethernet controller: Emulex
Corporation OneConnect 10Gb NIC (be3) (rev 03) <-- LOM Port 0 (Function 0)
16:00.1 Ethernet controller: Emulex Corporation OneConnect 10Gb NIC (be3)
(rev 03) <-- LOM Port 1 (Function 1)      16:04.0 Ethernet controller: Emulex
Corporation OneConnect 10Gb NIC (be3) (rev 03) <-- VF LOM Port 0 (Function
```

```

0)      16:08.0 Ethernet controller: Emulex Corporation OneConnect 10Gb NIC
(be3) (rev 03) <-- VF LOM Port 1 (Function 1)      # lspci -nn | grep Emulex
16:00.0 Ethernet controller: [0200]: Emulex Corporation OneConnect 10Gb NIC
(be3) [19a2:0710] (rev 03) <-- LOM Port 0 (Function 0)      16:00.1 Ethernet
controller: [0200]: Emulex Corporation OneConnect 10Gb NIC (be3)
[19a2:0710] (rev 03) <-- LOM Port 1 (Function 1)      16:04.0 Ethernet
controller: [0200]: Emulex Corporation OneConnect 10Gb NIC (be3)
[19a2:0710] (rev 03) <-- VF LOM Port 0 (Function 0)      16:08.0 Ethernet
controller: [0200]: Emulex Corporation OneConnect 10Gb NIC (be3)
[19a2:0710] (rev 03) <-- VF LOM Port 1 (Function 1)

```

- List Emulex PCI device by device code

```

# virsh nodedev-list | grep 19a2      pci_19a2_710 <-- LOM Port 0
(Function 0)      pci_19a2_710_0 <-- LOM Port 1 (Function 1)
pci_19a2_710_1 <-- VF LOM Port 0 (Function 0)      pci_19a2_710_2 <-- VF LOM
Port 1 (Function 1)      # virsh nodedev-dettach pci_19a2_710_1      Device
pci_19a2_710_1 detached <-- VF LOM Port 0 (Function 0)      # virsh
nodedev-dettach pci_19a2_710_2      Device pci_19a2_710_2 detached <-- VF LOM
Port 1 (Function 1)

```

11. Use the Virtual Machine Manager GUI to attach the VF (step 6 on page 33) to the guest operating system by using the add physical device option.

Note: To reconfigure a system that is already set up, remove the attached VF from the guest operating system by selecting the VF and using the remove option. See the documentation for the host operating system for information on using the Virtual Machine Manager to attach and remove virtual interfaces.

12. Start the RHEL 6.x guest operating system. Once the guest operating system is booted, use the `lspci` command to confirm the visibility of the OneConnect NIC function. The output shows a OneConnect NIC function, for example:

```

03:00.0 Ethernet controller: Emulex Corp. Emulex OneConnect 10Gb
NIC (be3)

```

13. The OneConnect Ethernet driver automatically loads with the out-of-box driver, and creates the network interfaces. Use the `ifconfig` command to verify that the interface is created.
14. After configuring the network interfaces with proper IP addresses, you can send and receive network traffic from the VM. See the documentation for the host and guest operating systems for information on network configuration.

Configuring VFs

In operating system distributions with newer IP commands that support VF configuration options, the host administrator can perform the following by using the “`ip link set`” command:

- Change the default MAC address.
- Configure VLAN.
- Configure the transmission rate.
- Set the QoS parameter on VFs.

The “`ip link set`” command is part of RHEL 6.x distributions.

MAC Address Configuration

The OneConnect UCNA is shipped with factory-configured MAC addresses for the network interfaces corresponding to the PFs. The driver generates random MAC addresses for the network interfaces corresponding to the VFs based on the factory-configured MAC address. Other MAC addresses can be assigned for the interfaces corresponding to the VF using IP utility commands in the hypervisor.

To configure the MAC address for the virtual function, run the following command in the hypervisor:

```
# ip link set eth<X> vf <VFN> mac <MAC-ADDR>
```

where:

- eth<X> is the interface corresponding to the physical function.
- <VFN> is the VF number (0-based) corresponding to the interface for which you are configuring the MAC address.
- <MAC-ADDR> is the MAC address you are configuring.

For example, to configure the MAC address for eth0 for VF number 0, run the following command in the hypervisor:

```
# ip link set eth0 vf 0 mac 00:16:88:AA:BB:AA
```

Additionally, if the VM is already running and the VF driver is loaded, run the following command in the guest operating system:

```
# ifconfig eth0 hw ether 00:16:88:AA:BB:AA
```

where “eth0” is the interface corresponding to the VF in the hypervisor.

VLAN Configuration

This section includes examples for transparent tagging and virtual guest tagging (VGT)

When transparent tagging is configured for a VF, the NIC transparently tags all traffic from the VF with the configured VLAN ID. The VM is not aware of the VLAN tag.

To assign a transparent VLAN ID to the VF, run the following IP command in the hypervisor:

```
# ip link set eth<X> vf <VFN> vlan <VLAN>
```

where:

- eth<X> is the interface corresponding to the physical function.
- <VFN> is the VF number corresponding to the interface for which you are configuring the VLAN.
- <VLAN> is the VLAN ID you are configuring.

Example:

```
# ip link set eth0 vf 0 vlan 5
```

Note: With transparent VLAN tagging, only one VLAN ID per VF can be configured.

VGT allows the VMs to configure their own VLAN interfaces for the VFs attached to them. To enable VGT, run the following IP command in the hypervisor before starting the VM to which the VF is attached. For example:

```
# ip link set eth0 vf 0 vlan 4095
```

Note: With VGT, multiple VLANs can be created on one VF.

Transmission Rate Configuration

Configure the transmission rate limit (TX-RATE) on a VF interface from the hypervisor using the following IP command syntax:

```
# ip link set eth<X> vf <VFN> rate <TX-RATE>
```

where:

- eth<X> is the interface corresponding to the physical function.
- <VFN> is the VF number corresponding to the interface that you are configuring the TX-rate.
- <TX-RATE> is the transmission rate limit, in Mb/s.

For example, to set a TX-RATE of 5000 Mb/s for the VF 0, run

```
# ip link set eth0 vf 0 rate 5000
```

Viewing VF Properties

To view the properties configured to VFs attached to a PF, use the following IP command syntax:

```
# ip link show eth<X>
```

where eth<X> is the interface corresponding to the physical function.

For example, to view the properties of a PF at “eth0” (which has VFs 0, 1 associated with it), run the following command in the hypervisor:

```
# ip link show eth0
```

Expected example output:

```
eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP qlen 1000
link/ether 00:00:c9:bb:16:ee brd ff:ff:ff:ff:ff:ff
vf 0 MAC 00:00:c9:9d:90:80, tx rate 10000 (Mbps)
vf 1 MAC 00:00:c9:9d:90:81, tx rate 10000 (Mbps)
```

iSCSI Driver Configuration

iSCSI Driver Parameter Descriptions

Table 3-5 lists the user-configurable iSCSI driver parameters. It includes a description of the parameters, default values, and their configuration value range.

Note: If the value provided for a parameter exceeds the parameter's defined minimum or maximum value, the driver logs an error in the event log, and continues to load using the parameter's default value.

Table 3-5 iSCSI Driver Configuration Parameters

Parameter	Description
ETO	<p>The Extended Timeout (ETO) parameter specifies the amount of time, in seconds, that the initiator driver waits for the target to become available once it has lost the connection to the target during an I/O operation.</p> <p>The minimum value is 0. The maximum value is 3600. The default value is 30.</p> <p>Note: If the ETO value is set between 0 - 19, the driver assumes a value of 20 internally. You will not see any modifications in the registry.</p> <p>For more information, see “Configuring the ETO and LDTO Parameters” on page 40 and “Error Handling using ETO and LDTO Parameters” on page 42.</p>
large_io	<p>The Large I/O (large_io) parameter specifies the maximum transfer size in a single I/O request in KBs (1KB = 1024 bytes). By default, the OneConnect iSCSI driver supports up to 128 KB of data in a single I/O request. When large_io is set to 512, then up to 512 KB of data can be supported in a single I/O request.</p> <p>The minimum value is 128 (default). The maximum, and only other, value is 512.</p> <p>For more information see “Configuring the large_io Parameter” on page 40.</p>
LDTO	<p>The Link Down Timeout (LDTO) parameter determines the amount of time, in seconds, that the initiator driver waits for the controller's physical link to become available before reporting that the LUNs are unavailable to the operating system.</p> <p>The minimum value is 0. The maximum value is 3600. The default value is 20.</p> <p>For more information, see “Configuring the ETO and LDTO Parameters” on page 40 and “Error Handling using ETO and LDTO Parameters” on page 42.</p>
im_policy	<p>The Interrupt Moderation policy (im_policy) parameter specifies the rate of interrupts for OneConnect UCNAs. The possible values are:</p> <ul style="list-style-type: none"> 0 - disables the interrupt algorithm 1 - highest interrupt rate 2 - moderate interrupt rate (default) 3 - an interrupt rate between moderate and lowest 4 - lowest interrupt rate <p>For more information, see “Configuring the im_policy Parameter” on page 41.</p>

Configuring the iSCSI Driver Parameters

Configuring the ETO and LDTO Parameters

ETO and LDTO values are configurable during insmod time. The ETO value specified during insmod is the default ETO value that is applied to all targets.

The following example configures the LDTO and ETO during insmod of the driver. These settings must be used every time the iSCSI driver is loaded by insmod.

```
insmod be2iscsi.ko ldto=25 eto=20
```

To pass module parameters to the iSCSI driver when it is being used as a boot controller, edit the file or the /etc/modprobe.conf file. For example, add the following line:

```
options be2iscsi ldto=25 eto=20
```

Then type

```
# mkinitrd
```

The change takes effect the next time the system is booted.

For information on these parameters in relation to SCSI error handling, see “Error Handling using ETO and LDTO Parameters” on page 42.

Configuring the large_io Parameter

The large I/O (large_io) parameter specifies the maximum transfer size in a single SCSI command of 128 or 512 KB.

By default, large_io=128, which means the iSCSI driver supports up to 128 KB and 32 scatter gather entries in a single SCSI command. If applications issue I/O requests that are larger than 128 KB or need more than 32 scatter gather entries, the request is split into multiple requests by the driver.

When large_io=512, the iSCSI driver can support up to 512 KB of data and a total of 128 scatter gather entries in a single SCSI command. In this case, the iSCSI driver supports a larger maximum transfer size, but also consumes a larger amount of physical memory. Intermediate sizes between 64 and 512 KB are accepted, but the memory used by the driver is the same as for a value of 512. If applications issue I/O requests that are larger than 512 KB or need more than 128 scatter gather entries, the request is split into multiple requests by the driver.

Manually Setting the large_io Parameter

To manually set the large_io parameter, type

```
insmod be2iscsi.ko large_io=512
```

Permanently Setting the large_io Parameter

To permanently set large_io on every iSCSI boot, edit the /etc/modprobe.conf file. In the following example, note that the lines beginning with “alias” are added automatically when the driver is installed. The lines beginning with “options” need to be added manually for non-default parameters, such as large_io.


```
alias eth0 tg3
alias scsi_hostadapter ata_piix
alias scsi_hostadapter1 usb-storage
alias scsi_hostadapter2 be2iscsi
options be2iscsi eto=0 ldto=0 large_io=512
alias eth1 be2net
alias eth2 be2net
```

After adding the options line, save the file, and then rebuild initrd.

Configuring the im_policy Parameter

The interrupt moderation policy setting (im_policy parameter) controls the rate of interrupts for the OneConnect UCNA. Table 3-6 defines the available im_policy settings.

Table 3-6 Interrupt Moderation Policy (im_policy) Settings

im_policy Value	Setting Type	Description
0	Disabled	Interrupt moderation algorithm is turned off in the driver.
1	Aggressive	Highest interrupt rate among available settings.
2	Moderate	Default value; The interrupt rate varies between 3500 to 10000 interrupts per second.
3	Conservative	Lower interrupt rate than Moderate.
4	Very conservative	The minimum interrupt rate among available settings.

By default, the driver implements an interrupt moderation scheme that is based on the I/O load and the interrupt rate. Changing the interrupt moderation policy setting should be based on the initiator system configuration, the number of connected iSCSI targets, the I/O load, and the throughput and latency offered by these iSCSI targets.

The default Moderate setting (im_policy = 2) varies the interrupt rate between 3500 to 10000 interrupts per second. While the Moderate setting may work for most configurations, there are instances when the setting may need to be altered.

On systems capable of sustaining a higher interrupt rate and when a few number of targets are connected (up to 8), the Aggressive setting (im_policy = 1) results in lower latency and higher values of I/O operations per second (IOPs). However, this higher interrupt rate could also result in system stalls and freezes, especially during higher values of queue depth and smaller sized I/O requests.

In a configuration that involves a large number of iSCSI targets (more than 32 or 64) and higher values of queue depth, the Conservative (im_policy = 3) or Very Conservative (im_policy = 4) setting may be more desirable. Though these settings increase the latency of an I/O request, the lowered interrupt rate may allow the system to be functional under a high load.

For example, to manually set the im_policy parameter for the highest interrupt rate (Aggressive), type

```
insmod be2iscsi.ko im_policy=1
```

To permanently set `large_io` on every iSCSI boot, edit the `/etc/modprobe.conf` file, type

```
options be2iscsi eto=0 ldto=0 large_io=512 im_policy=1
```

After adding the options line, save the file, and then rebuild `initrd`.

iSCSI Error Handling

Error Handling using ETO and LDTO Parameters

The goal of iSCSI error handling is to be tolerant of link-level and target-level failures up to configured timeout values so that I/O errors are not seen by the application or operating system.

The error handling is triggered under the following conditions:

- Loss of connection to the target due to target and/or network disconnection at the target.

If the driver has I/O requests pending with the target and the target becomes unavailable (due to the target going down or failing over, or network issues at the target), the driver queues up the I/O requests internally for a configured period of time. The threshold value of this period is the ETO value.

- Loss of immediate link to the initiator (such as cable disconnect or port failure)

The UCNA firmware detects and notifies the driver of a loss of the link. When this happens, the driver queues up the I/O requests internally to a configured period of time so that the operating system does not see I/O errors. The threshold value of this period is the LDTO value.

When the configured ETO or LDTO value is reached, and the initiator is still unable to connect to the target, the driver fails all I/O requests. At this point, I/O errors are seen by the application and operating system.

Note: Following a link up, switch ports can take a long time to initialize and go to the forwarding state. Because of this, additional time should be added to the ETO and LDTO settings to eliminate I/O disruption or target unavailability. If the switch port is connected to a single host, then the PortFast mode can be enabled on the switch port to eliminate delays when it transitions to the forwarding state.

Error Handling Under MultiPath I/O (MPIO) and Cluster Configurations

In an MPIO or cluster configuration, fault tolerant software is present on the system that makes the iSCSI driver error handling redundant. These configurations also require that I/O errors be reported as soon as they are detected so that the software can fail over to an alternate path or an alternative node as quickly as possible.

When the iSCSI driver runs under these configurations, the error handling implemented in the driver must be turned off by setting the default values of LDTO and ETO to zero. The changes take effect during the next driver load.

Virtual Port (VPort) Configuration

VPort Configuration Prerequisites

Before configuring VPorts, note the following:

- Ensure you are using the latest recommended firmware for VPort functionality. Check the Emulex website for the latest firmware.
- Loop devices and NPIV are not supported on the same port simultaneously. If you are running a loop topology and you create a VPort, the VPort's link state is offline.
- You can create VPorts only on 4, 8, 10, and 16 Gb/s adapters. You cannot create VPorts on 1 and 2 Gb/s adapters.
- VPorts do not persist across system reboots.

Creating, Deleting, and Displaying VPorts

VPorts are created through sysfs entries that are presented in the physical port's sysfs directory. The `vport_create` and `vport_delete` sysfs entries are discussed in "VPort sysfs Entries" on page 46, but there are also three scripts for creating, deleting and displaying VPorts. The scripts reside in the `/usr/sbin/lpfc` directory and are part of the OneCommand Manager application kit.

When NPIV is enabled and VPorts are configured, it may take longer for the adapter to finish discovery in some cases because each VPort must perform discovery independently. As more VPorts are configured, the amount of time that the driver and adapter take to finish discovery of remote ports on the SAN increases. To compensate for this extended amount of time taken in discovery, set the `lpfc_devloss_tmo` parameter to 60 seconds when NPIV is enabled.

Creating VPorts Using the `mkvport.sh` Script

You can use the `mkvport` script to create VPorts. To see the usage information, run the script with no parameters specified. The `mkvport.sh` script uses the following syntax:

```
./mkvport.sh <Physical Port's Host number> <Port Name> <Node Name>
```

You must supply the physical port's host number, WWPN, and WWNN when using the `mkvport.sh` script. For example, to create a VPort with port name of 10000000c94ac63a and a node name of 20010000c94ac63a on the physical port with scsi_host name "host7", type

```
./mkvport.sh host7 10000000c94ac63a 20010000c94ac63a
```

This script fails if the VPort is not created.

Note: It is possible for a VPort to be created successfully but to be in a failed state. For example, loop devices and NPIV are not supported on the same port simultaneously. If you are running a loop topology and you create a VPort, the VPort's link state will be offline.

Deleting VPorts Using the rmvport.sh Script

Note: You must un-map, un-mount, and flush I/Os to VPort-connected devices before deleting the VPort.

You can use the rmvport script to delete VPorts. To see the usage information, run the script with no parameters specified. The rmvport.sh script uses the following syntax:

```
./rmvport.sh <VPort's Host number>
```

-or-

```
./rmvport.sh <Port Name> <Node Name>
```

To delete the VPort with a port name of 10000000c94ac63a and a node name of 20010000c94ac63a, type

```
./rmvport.sh 10000000c94ac63a 20010000c94ac63a
```

This script may take up to 30 seconds to finish. The script fails if the VPort is not deleted.

Displaying VPorts Using the lsvport.sh Script

You can use the lsvport script to display the VPorts and physical ports that are present on the system. Run the script with no parameters to display port information. For example:

```
./lsvport.sh
lpfc0: host6 10000000c93a5b5e:20000000c93a5b5e LP10000 NPIV Not Supported
lpfc1: host7 10000000c93a5b5d:20000000c93a5b5d LP10000 NPIV Not Supported
lpfc2: host8 10000000c93cc8dd:20000000c93cc8dd LPe12000 NPIV Physical
lpfc3: host9 10000000c93cc8dc:20000000c93cc8dc LPe12000 NPIV Physical
lpfc4: host10 10000000c94ac63a:20010000c94ac63a NPIV Virtual (VPI 1)
```

In reference to the previous example:

- For LPFC0 and LPFC1, “NPIV Not Supported” indicates that this adapter/firmware combination does not support the creation of VPorts.
- For LPFC2 and LPFC3, “NPIV Physical” refers to a physical port of this adapter.
- For LPFC4, “NPIV Virtual” refers to a VPort of this adapter.

VPort sysfs

VPort sysfs Tree

For FC 8.2.0.x Drivers

When a VPort is created, two new directories are created in the class tree:

```
/sys/class/scsi_host/hostY/
/sys/class/fc_host/hostY/
```

Creating a new VPort also creates a new sysfs directory in the bus and devices tree:

```
ls /sys/bus/pci/drivers/lpfc/0000:07:00.0/host8/
```

```
fc_host:host8 host10 power scsi_host:host8 uevent
ls /sys/bus/pci/drivers/lpfc/0000:07:00.0/host8/host10
fc_host:host10 power scsi_host:host10 uevent
```

In this example, host 8 is the physical port, and host 10 is a VPort that was created on host 8.

For FC 8.3.5.x Drivers

When a VPort is created, three new directories are created in the class tree:

```
/sys/class/scsi_host/hostY/
/sys/class/fc_host/hostY/
/sys/class/fc_vports/vport-X:0-Z/-
```

Creating a new VPort also creates a new sysfs directory in the bus and devices tree:

```
/sys/bus/pci/drivers/lpfc/0000:A:B:C/hostX/vport-X:0-Z/hostY
/sys/devices/pci0000:A/0000:A:B:C/hostX/vport-X:0-Z/hostY
```

In both directories there is a hostY directory that contains the remote ports that this new host can access:

```
/sys/bus/pci/drivers/lpfc/0000:A:B:C/hostX/vport-X:0-Z/hostY
/sys/bus/pci/drivers/lpfc/0000:A:B:C/hostX/vport-X:0-Z/hostY/rport-Y:0-0
/sys/bus/pci/drivers/lpfc/0000:A:B:C/hostX/vport-X:0-Z/hostY/rport-Y:0-1
/sys/bus/pci/drivers/lpfc/0000:A:B:C/hostX/vport-X:0-Z/hostY/rport-Y:0-2
```

In this example:

- “Y” indicates the new host value for the VPort that was created.
- “X” indicates the host value for the parent fc_host that this VPort was created from.
- “Z” indicates the instance of VPort created from the parent fc_host. A, B, and C indicate the PCI hierarchy for each physical LPFC port.

In other words, “hostY” is the new host created for the new VPort. “vport-X:0-Z” uniquely identifies the VPort and indicates the parent host structure (X) that this VPort was created by.

For example, if a VPort is created from host5, a new scsi_host, a new fc_host, a new fc_vport, and a new entry under the bus tree are created as well.

```
ls /sys/class/scsi_host/
host0 host1 host4 host5 host6
ls /sys/class/fc_host/
host4 host5 host6
ls /sys/class/fc_vports/
vport-5:0-0
```

VPort sysfs Entries

Note: VPort sysfs entries in Table 3-7 are only present if the driver was loaded with `lpfc_enable_npiv` enabled.

Table 3-7 VPort sysfs Entries

VPort sysfs Entries	Type	Range /Input	Location and Description
<code>lpfc_peer_port_login</code>	Read/Write	0=Off(default) 1=On	<p><code>/sys/class/scsi_host/hostX/lpfc_peer_port_login</code></p> <p>This entry sets the port's behavior when discovering targets in the SAN. The default behavior (value=0) will log in only to N_Ports that are physically located on a different port. The port will still attempt to log in to targets on all other ports (including the other port in a dual-port adapter).</p> <p>If this parameter is turned on (value=1), then the port attempts to log in to all N_Ports, even if they are physically located on the same port.</p> <p>Note: This parameter was created to reduce the amount of hardware resources (for example, RPIs) that the driver requires. In a configuration where there are many VPorts on one physical port, this capability greatly reduces the number of RPIs that the driver uses.</p>
<code>lpfc_restrict_login</code>	Read/Write	0=Off 1=On (default)	<p><code>/sys/class/scsi_host/hostX/lpfc_restrict_login</code> (VPorts only)</p> <p>This entry sets the VPort's behavior when discovering targets in the SAN. The default behavior (value=1) prevents the VPort from logging into other initiator ports in the SAN. It also rejects logins from other ports in the SAN, because it assumes that all ports that send a PLOGI are initiators.</p> <p>If this sysfs entry is turned off, the driver attempts to log in to every port that it can access in the SAN, and accepts logins from all ports.</p> <p>Note: This parameter was created to reduce the amount of hardware resources (for example, RPIs) that the driver requires. In a SAN where there are other initiators, this capability greatly reduces the number of RPIs that the driver uses.</p>
<code>max_npiv_vports</code>	Read-only	integers	<p><code>/sys/class/fc_host/hostX/max_npiv_vports</code></p> <p>This entry displays the maximum number of VPorts that are supported by the <code>fc_host</code>'s underlying hardware.</p> <p>This sysfs entry exists only if the <code>vport_create</code> and <code>vport_delete</code> sysfs entries exist. If an <code>fc_host</code> does not support NPIV, this sysfs entry may not exist.</p> <p>Use this sysfs entry with the <code>npiv_vports_inuse</code> entry to determine whether the maximum number of VPorts have been created on this <code>fc_host</code>.</p>

Table 3-7 VPort sysfs Entries (Continued)

VPort sysfs Entries	Type	Range /Input	Location and Description
node_name	Read-only	16-byte hexadecimal value	<p>For FC 8.2.0.x drivers: /sys/class/fc_host/hostX/node_name</p> <p>For FC 8.3.5.x drivers /sys/class/fc_host/hostX/node_name/sys/class/fc_vports/vport-X:0-Z/node_name</p> <p>These entries display the physical or VPort's node name. You assign this value when the VPort is created, and it is transmitted to the fabric upon fabric login.</p>
npiv_vports_inuse	Read-only	integers	<p>/sys/class/fc_host/hostX/npiv_vports_inuse</p> <p>This entry displays the number of VPorts that were created on this fc_host.</p> <p>This sysfs entry exists only if the vport_create and vport_delete sysfs entries exist. If an fc_host does not support NPIV, this sysfs entry may not exist.</p> <p>Use this sysfs entry with max_npiv_vports to determine whether the maximum number of VPorts have been created on this fc_host.</p>
port_name	Read-only	16-byte hexadecimal value	<p>/sys/class/fc_host/hostX/port_name/sys/class/fc_vports/vport-X:0-Z/port_name</p> <p>This entry displays the physical or VPort's port name. You assign this value when the VPort is created, and it is transmitted to the fabric upon fabric login.</p>
vport_create	Write-only	WWPN; WWNN	<p>/sys/class/fc_host/hostX/vport_create</p> <p>This entry creates a VPort on the physical port that hostX is located on. The new VPort will have a WWPN and WWNN present on the fabric based on the WWPN and WWNN that are entered with this sysfs entry.</p> <p>This entry returns a "0" if the VPort creation was successful. A non-zero value indicates that the VPort was not created.</p> <p>If an fc_host does not support NPIV, then this sysfs entry may not exist.</p> <p>Note: It is possible for the VPort creation to succeed but for the VPort to be in a failed or inoperative state. Use the new sysfs tree created by the new VPort to check the state of the new VPort.</p>

Table 3-7 VPort sysfs Entries (Continued)

VPort sysfs Entries	Type	Range /Input	Location and Description
vport_delete	Write-only	WWPN; WWNN	<p>/sys/class/fc_host/hostX/vport_delete</p> <p>This entry deletes a VPort on the physical port that hostX is located on. The VPort matching the WWPN and WWNN is immediately deleted.</p> <p>This entry returns a “0” if the VPort deletion was successful. A non-zero value indicates that the VPort was not deleted.</p> <p>If an fc_host does not support NPIV, then this sysfs entry may not exist.</p> <p>Note: This entry deletes the VPort even if there are mounted file systems being accessed through this VPort, or if there are open files on it.</p>

Monitoring VPorts with fc_vport (FC and FCoE Drivers)

For FC 8.2.0.x Drivers

In the FC 8.2.0.x driver, the fc_vport directory does not exist (yet) so a link from the physical port to the VPort is present in the fc_host’s device directory.

```
ls /sys/class/fc_host/host5/device/
fc_host:host5 power scsi_host:host5
host6 uevent
```

In this example, host6 is a VPort of physical port host5.

To find the VPorts that have been created by a physical port, you can list the fc_host’s device directory for the physical port. This gives you a link to the fc_host and scsi_host directory as usual, and it also displays a list of VPorts (in the form of hostx) that were created on this physical port.

For FC 8.3.5.x Drivers

In the FC and FCoE 8.3.5.x driver, the transport creates an fc_vports directory that you can use to monitor VPorts. This directory is populated entirely of VPorts and has links from each to the fc_host associated with that VPort.

```
ls /sys/class/fc_vports/
vport-5:0-0
ls -ld /sys/bus/pci/drivers/lpfc/*/host*/*/host*
/sys/bus/pci/drivers/lpfc/0000:03:06.1/host5/vport-5:0-0/host6
ls /sys/devices/pci*/*/host5/vport-5*/host6
power rport-6:0-0 rport-6:0-1 rport-6:0-2 uevent
ls /sys/devices/pci*/*/host5/vport-5*/host6/rport-*
/sys/devices/pci00:03:00:03:06.1/host5/vport-5:0-0/host6/rport-6:0-0:
power uevent

/sys/devices/pci00:03:00:03:06.1/host5/vport-5:0-0/host6/rport-6:0-1:
power uevent
```



```
/sys/devices/pci00:03/00:03:06.1/host5/vport-5:0-0/host6/rport-6:0-2:  
power target6:0:0 uevent
```

In this example:

- There is a new entry in the `fc_vports` directory for the VPort (`vport-5:0-0`). The `vport-5:0-0` entry indicates that the VPort was created from `host5` and it is the first (0) VPort to be created on that `fc_host`.
- The new host for the VPort is `host6`, and it will appear in the usual directories.
- There is also a new directory in the bus tree. This new directory indicates that `host6` was created under `vport-5:0-0` (which was created from `host5`).

VPort Configuration Limits

VPort configuration limits are designated as enforced or unenforced. Enforced limits are limits that the driver enforces and prevents the user from exceeding. Unenforced limits are limits that the driver cannot enforce, but configurations that exceed them are unsupported.

The following VPort configuration limits have been tested with and are supported by the Emulex driver. Configurations that exceed one or more of these limits are unsupported.

- Before the VPort is deleted or the driver is unloaded, I/O devices accessed through a VPort must be stopped and file systems must be unmounted.
- For enterprise-class adapters, the maximum number of VPorts configurable on a physical port is 64. The hardware allows more than 64 VPorts to be created, but the driver has only been qualified at 64. For mid-range adapters, the maximum number of VPorts configurable on a physical port is 16.
- The maximum number of LUNs supported on each driver port is 256.
- The maximum number of targets supported for each driver port is 255.
- The maximum number of driver ports in one zone is 64. This limit is based on the system's ability to recover from link events within the time constraints of the default timers.

The NPIV use-cases that involve virtual server environment include associating a VPort with a virtual machine, and placing the virtual machine in its own zone. This results in one VPort per zone. In the case of load balanced environments, this can increase typically to two VPorts per virtual machine, to a practical limit of something far less than 50.

In the NPIV cases not related to virtual server environments, zoning is typically initiator-zoning, again resulting in one VPort, or a low number of VPorts in the case of load-balancing, within a given zone. If there are too many VPorts within a single zone, expected behavior includes devices being lost after link events.

- The minimum lifetime of a VPort is 60 seconds. There is an unenforced limit of 60 seconds between the creation of a VPort and the deletion of the same VPort. VPorts are designed to live for a long time in the system, and the creation of VPorts is asynchronous. This means that a VPort might not be finished with FC or SCSI discovery when the command to create a VPort is finished.

DHCHAP Authentication and Configuration

Note: This section is applicable to FC 8.2.0.x drivers only.

To activate FC-SP/ Authentication between the adapter host port and fabric F_Port using DHCHAP, modify the DHCHAP-associated driver properties in the driver configuration file.

The LPFC driver for Linux version 8.2.0.x supports MD5 and SHA-1 hash functions and supports the following DH groups: Null, 1024, 1280, 1536, and 2048.

Enabling Authentication

Enabling authentication is a two-step process. To enable authentication:

1. Start the `fcauthd` daemon.
2. Set the `lpfc_enable_auth` module parameter to 1 (enabled).

fcauthd Daemon

The LPFC driver requires the `fcauthd` daemon to perform authentication tasks for it. To enable authentication, you must have this daemon running. If you want to load the LPFC driver with authentication enabled, the `fcauthd` daemon should be running before the driver is loaded. The LPFC driver can start with authentication enabled if the daemon is not running, but all ports are placed into an error state.

When the daemon is started, the LPFC driver should discover the daemon and reset the adapter to enable the LPFC driver to perform authentication. To test if this daemon is running, start the daemon, or stop the daemon, you must use the `/etc/init.d/fcauthd` script.

The script syntax is `/etc/init.d/fcauthd <parameter>`.

fcauthd Daemon Parameters

The `fcauthd` daemon supports the following parameters:

- `start` - To start the `fcauthd` daemon, pass the `start` command to the `fcauthd` script. This command loads the daemon into memory, opens a netlink connection for the driver, and reads the authentication configuration database into memory for use by the LPFC driver.
- `stop` - To stop the `fcauthd` daemon, pass the `stop` command to the `fcauthd` script. This command takes down the netlink connection between the `fcauthd` daemon and the LPFC driver, and stops the `fcauthd` daemon.
- `reload` - The `reload` command reloads the authentication configuration database into memory. This is done whenever the database is changed by another application (such as the OneCommand Manager application) or by you. If the database is changed, the new configuration information is not used until the `fcauthd` daemon reloads the database.
- `status` - This command displays the current status of the `fcauthd` daemon. The status should be either running or stopped.

- restart - The restart command stops the fcauthd daemon and then restarts it.
- condrestart - The conditional restart command checks the status of the fcauthd daemon. If it is running, it issues a stop and then a start command. If the fcauthd daemon is not running, nothing happens.

lpfc_enable_auth Module Parameter

Use the lpfc_enable_auth module parameter to enable or disable authentication support. This module parameter can be set when the LPFC driver is loaded to enable or disable authentication on all Emulex adapters in the system, or it can be set dynamically after the LPFC driver is loaded to enable or disable authentication for each port (physical and virtual). The default setting for the lpfc-enable-auth module parameter is disabled. See “Dynamic FC and FCoE Driver Parameters” on page 22.

Authentication Configuration Parameters

You can configure each port’s authentication parameters using the OneCommand Manager application. See the latest *OneCommand Manager Application User Manual*.

Setting Remote and Local Passwords

You can configure each port’s password using the OneCommand Manager application. See the latest *OneCommand Manager Application User Manual*.

LPFC Driver Performance Tuning

Overview

Version 8.7.x.x of the lpfc driver includes the following configurable parameters that can enhance performance:

- lpfc_fcp_io_channel
- lpfc_fcp_io_sched
- lpfc_fcp_imax

These features are available through module parameters that are defined in the lpfc driver as well as sysfs entries defined by the Linux kernel.

In addition, you can use the lpfc_vector_map.sh script to map a specific I/O channel to a specific CPU.

This section provides more information about how the tuning parameters and script can improve Emulex adapter performance.

lpfc_fcp_io_channel

The lpfc_fcp_io_channel module parameter can be configured at driver load time. It defines the number of I/O channels supported by the driver. The driver is capable of supporting parallel I/O paths, and each I/O path is capable of posting and completing FCP commands independent of the other.

OneConnect and LPe16000 adapters that are running in MSI-X interrupt mode can use more than one I/O path. Each I/O channel is composed of a unique MSI-X vector-EQ/CQ/WQ tuple. This parameter will override the value of the `lpfc_fcp_eq_channel` and the `lpfc_fcp_wq_count` parameters.

Note: LPe12000, LPe11000 and LP11000 adapters only support one I/O path, so this parameter has no effect on them.

By default, the driver is configured for four I/O channels per port. The driver will also limit the number of I/O channels to not exceed the number of online “logical” CPUs (as reported by `/proc/cpuinfo`).

`lpfc_fcp_io_sched`

The `lpfc_fcp_io_sched` module parameter can be configured at driver load time. It can also be set dynamically as an `sysfs` entry. The driver uses the parameter to determine which algorithm to use when scheduling an FCP I/O to an I/O channel.

Note: This parameter is only applicable for OneConnect and LPe16000 adapters.

When multiple I/O channels are in use, I/Os can be scheduled to an I/O channel in a round-robin fashion, or by determining which CPU is running when the I/O is submitted.

The default value (0) configures the driver for round-robin scheduling. A value of 1 configures scheduling by CPU.

`lpfc_fcp_imax`

The `lpfc_fcp_imax` can be configured at driver load time. It can also be set dynamically as an `sysfs` entry. This parameter defines the maximum number of interrupts per second that each adapter port will support.

Note: This parameter is only applicable for OneConnect and LPe16000 adapters.

Considerations

- The lower the value set, the more completions are coalesced by the adapter, which causes the driver to handle multiple completions under the context of one interrupt. The higher the value, the faster an interrupt is generated for a completed command. Therefore, a balanced or “tuned” system must be found.
- A lower value equals higher interrupt latency; a higher value equals lower interrupt latency.
- Faster completions consume more system resources / CPU cycles, as the overhead of one interrupt completes fewer commands. The value is divided by the number of I/O channels, and each I/O channel is separately configured for its own interrupt latency.

By default, the module parameter is configured for 50,000 interrupts per second per adapter port.

lpfc_vector_map.sh

The `lpfc_vector_map.sh` script uses kernel sysfs entry points to map a specific I/O channel (MSI-X vector-EQ/CQ/WQ tuple) to a specific CPU. The script should be run immediately after the driver is loaded.

This script resides in `/usr/sbin/lpfc`. The recommended way to run this script is by adding the following line to `/etc/modprobe.d/lpfc.conf` and thru `/etc/rc.d` for the initial boot:

```
install lpfc /sbin/modprobe --ignore-install lpfc;  
/usr/sbin/lpfc/lpfc_vector_map.sh
```

The script maps each interrupt vector allocated by the driver to a CPU, thereby spreading the interrupt load of the ports across multiple CPUs. Each vector, with its associated I/O channel, is sequentially mapped to a CPU in a round-robin fashion. The number of vectors assigned to each adapter port is defined by the `lpfc_fcp_io_channel` module parameter.

I/O channels, which correspond to vectors, are typically mapped to unique CPUs to enhance the ability of the driver to run multiple FCP commands in parallel. In addition, running this script forces I/O scheduling to be by CPU (`lpfc_fcp_io_sched = 1`) which increases performance when a specific I/O channel is mapped to a specific CPU.

The script has two modes of operation: Driver mode and HBA mode. By default, the script runs in Driver mode.

Driver Mode

Driver mode maps all vectors for all driver ports, starting with CPU0, sequentially assigning a new CPU for each vector belonging to the entire driver. If there are more interrupt vectors than CPUs, the vector assignment wraps back to CPU0 as needed.

HBA Mode

HBA mode maps all vectors for each specific adapter port, starting with CPU0, sequentially assigning a new CPU for each vector belonging to that adapter port.

Network Performance Tuning

The OneConnect UCNA is an x8, Generation 2 (Gen2) PCI-Express (PCIe) device. To obtain the best two-port throughput from the OneConnect UCNA, it should be installed in an x8 or x16 Gen2 PCIe slot.

Memory Bandwidth Considerations

The availability of higher memory bandwidth leads to better network performance. The following sections describe how memory bandwidth can be increased.

Enabling Optimal Bandwidth Options

Most computers offer multiple distinct memory channels, or memory interleaves, which may not be enabled by default. Check the manufacturer's documentation and BIOS parameters for details on enabling optimal memory bandwidth options.

Populate DIMM Slots

Typically, all the dual in-line memory module (DIMM) slots must be populated to make use of all the memory channels. As a general rule, using more DIMMs provides better performance by allowing a higher degree of memory-access interleaving to occur.

Disabling Memory Mirroring

Some servers may allow memory mirroring, where the total memory is divided in half and each location is stored twice. This allows fault recovery if one memory location detects an error, but it greatly reduces the perceived memory bandwidth of the system. Consider disabling memory mirroring if it is not needed.

Using a Fast Clock Speed for the Front Side Bus (FSB)

Nearly any desktop or low-end server has enough memory bandwidth for OneConnect UCNAs to support DMA at 20 Gb/s of data (10 Gb/s read, 10-Gb/s write). However, most of the memory demands come from the processor accessing the data for either packet copies in the non-offloaded networking stack or application accesses. All processor memory accesses use the FSB. The clock speed of this bus is critical for allowing efficient memory bandwidth. A system with a faster processor FSB clock speed performs better than a system with a slower FSB clock speed.

Network Memory Limits

The default values of tunable parameters in the Linux network stack are optimal for most network applications involving several TCP/UDP streams. The optimal size for the network queues and buffers depend on several factors such as protocol, number of streams (connections), request size, and application behavior. The following network configuration settings are a good combination to get the best uni-directional transmit and receive performance with six or more TCP connections/UDP streams:

```
echo 4096 87380 4194304 > /proc/sys/net/ipv4/tcp_rmem
echo 4096 16384 4194304 > /proc/sys/net/ipv4/tcp_wmem
echo 64000000 > /proc/sys/net/core/rmem_default
echo 64000000 > /proc/sys/net/core/rmem_max
echo 32000000 > /proc/sys/net/core/wmem_default
echo 32000000 > /proc/sys/net/core/wmem_max
```

These settings assume ideal conditions such as low latency, zero (or close-to-zero) packet loss in the network, enough free memory, and 10 Gb/s path-to-peer system.

These `tcp_rmem` and `tcp_wmem` values are also the default values in recent RHEL 5 and SLES 10 distributions. If your application requires best throughput with very small

number of connections (less than four), it may help to increase the `tcp_rmem` and `tcp_wmem` to much larger values:

```
echo 4096 87380 16777216 > /proc/sys/net/ipv4/tcp_rmem
echo 4096 65536 16777216 > /proc/sys/net/ipv4/tcp_wmem
```

TCP Segmentation Offload (TSO)

In low-loss networks, TSO considerably improves performance, and therefore must be enabled. TSO is enabled by default in the OneConnect network driver.

The `/proc/sys/net/ipv4/tcp_tso_win_divisor` process variable controls how aggressive the network stack can be in making TSO requests. For low-loss networks, Emulex recommends TSO divisor values in the range of 2 to 16. In most distributions, the default value of 3 seems to be the optimal choice for a no-loss network.

Smaller divisor values result in larger TSO chunks and better throughput, as well as better CPU utilization. However, if the receiver or the network is dropping frames (too many re-transmits on the transmit side as indicated by `netstat -st`), it may help to make TSO chunks smaller (by increasing the divisor value) or to turn TSO off. For example, to set the divisor level to a value of 8, run

```
echo 8 > /proc/sys/net/ipv4/tcp_tso_win_divisor
```

To turn TSO on or off, run one of the following `ethtool` commands:

```
ethtool -K eth<N> tso on
ethtool -K eth<N> tso off
```

where `eth<N>` is the name of the Ethernet device you are working on (for example, `eth0`).

Flow Control

The OneConnect UCNA supports IEEE 802.3x standard flow control, which uses control packets to temporarily pause the transmission of packets between two endpoints. These control messages are point-to-point, and are not forwarded by switches or routers. The UCNA can respond to flow control packets by temporarily pausing transmits. The UCNA can send flow control pause packets when the transmitter is overwhelming the system's receive bandwidth.

Flow control can greatly improve performance, as described in the following examples:

- The UCNA is installed in 4x PCIe slot or an underpowered server system.
If the PCIe bus does not provide 10 Gb/s of throughput due to chipset limitations or the bus width, the UCNA cannot maintain 10 Gb/s of incoming receive data. It starts dropping packets quickly. In this situation, it may be beneficial to enable receive flow control in the UCNA, and enable flow control in the attached switch for all devices. This helps to slow down the transmitters.
- The UCNA transmits to 1 Gb/s devices, especially when using a non-TCP protocol.
If the UCNA transmits to a 10 Gb/s switch with attached 1 Gb/s clients, the UCNA may overwhelm the switch. The switch is then forced to start dropping

packets because, although it may receive a 10 Gb/s stream, the client can only sink a 1 Gb/s stream. In this situation, it may be beneficial to enable transmit flow control in the UCNA, and enable flow control for the 10 Gb/s switch port.

You can configure the OneConnect UCNA to respond to flow control pause frames from the other side (switch or router) using the following ethtool commands:

```
ethtool -A eth<N> pause rx on
ethtool -A eth<N> pause rx off
```

where eth<N> is the name of the Ethernet device you are working on (for example, eth0).

You can configure the UCNA to send flow control pause frames using the following ethtool commands:

```
ethtool -A eth<N> pause tx on
ethtool -A eth<N> pause tx off
```

where eth<N> is the name of the Ethernet device you are working on (for example, eth0).

RX and TX flow control are enabled by default in the OneConnect UCNA. When priority flow control is enabled in the adapter, normal flow control cannot be enabled.

Refer to the switch/router documentation to determine how link level flow control can be configured on the switch/router to which the OneConnect UCNA port is connected.

Note: In multichannel configurations where multiple PCI functions are exposed for a single 10G Ethernet port, the flow control parameter for a port can be configured through any interface associated with the physical port, and the configured property will apply to all interfaces associated with the port.

RX Frame Coalescing

The OneConnect Ethernet driver coalesces regular-sized TCP segments to a large frame before passing it to the network stack, which may improve TCP receive performance. RX frame coalescing is implemented using the GRO mechanism (in Linux driver versions that support GRO) or the LRO mechanism (in older Linux driver versions).

RX frame coalescing is enabled by default. In some configurations where the end point for the TCP connection to which the packets belong is not in the current server (for example, the end point is a router), RX coalescing should not be enabled.

GRO can be disabled using the -K option with the ethtool command:

```
ethtool -K eth<N> gro off
```

LRO can be disabled using the -C option with the ethtool command:

```
ethtool -C eth<N> rx-frames 1
```

where eth<N> is the name of the Ethernet device you are working on (for example, eth0).

Maximum Transmission Unit (MTU)

The OneConnect Ethernet driver supports MTUs between 256 bytes and 9000 bytes. The default MTU is set to 1500. If other elements in the network path support a larger MTU, you can increase the MTU up to 9000 using the `ifconfig` command. To do this run:

```
ifconfig <ethN> mtu 9000
```

where `eth<N>` is the name of the Ethernet device you are working on (for example, `eth0`).

Interrupt Coalescing

The OneConnect Ethernet driver tries to reduce the number of interrupts by delaying the interrupts from the UCNA. This reduces CPU utilization during a high traffic rate. The interrupt delay duration can be set to change dynamically within a range of values, depending on the receive rate (known as Adaptive Interrupt Coalescing (AIC)), or can be set to a constant value.

Setting the Interrupt Delay Duration to a Range of Values (AIC)

For receive interrupts, AIC is enabled by default. When AIC is enabled, the default low limit is 0 microseconds and the default high limit is 96 microseconds. In low traffic, the interrupt delay is set to 0 for best latency. As the number of interrupts per second increases, the delay is increased to higher values proportional to the receive rate, up to the default high limit of 96 microseconds. You can change the low and high limits using `ethtool`. For example, to set a low limit of 8 and a high limit of 40, run

```
ethtool -C eth<N> rx-usecs-low 8  
ethtool -C eth<N> rx-usecs-high 40
```

where `eth<N>` is the name of the Ethernet device you are working on (for example, `eth0`).

For transmit interrupts, AIC is not supported.

To disable AIC and set the interrupt delay duration to a constant value, see the following section.

Setting the Interrupt Delay Duration to a Constant Value

The interrupt delay duration can be set to a constant value for both receive and transmit interrupts. The possible interrupt delay duration values are 0 to 96 microseconds, in 8 microsecond increments.

For receive interrupts, disable AIC (since it is enabled by default) and set the interrupt delay duration using `ethtool`. For example, to disable AIC and set the constant RX interrupt delay to 8 microseconds, run

```
ethtool -C eth<N> adaptive-rx off rx-usec 8
```

where `eth<N>` is the number of the Ethernet interface you are working on.

If your application requires low or predictive latency, Emulex recommends that you turn off AIC and set `rx-usecs` to 0.

For transmit interrupts, the default interrupt delay duration is 96 microseconds. You can change this value using `ethtool`. For example, to set the TX interrupt delay to 64 microseconds run

```
ethtool -C eth<N> tx-usec 64
```

where `eth<N>` is the number of the Ethernet interface you are working on.

Receive Side Scaling (RSS)

Distributing the incoming traffic across several receive rings with separate interrupt vectors helps to distribute the receive processing across several CPU cores. This could reduce the packet drop and improve the packet rate in certain applications. RSS is enabled in non-SR-IOV and non-multichannel configurations. In multichannel configurations, RSS can be enabled in the first function of each port using the module parameter `rss_on_mc`. See Table 3-3, Ethernet Driver Configuration Parameters, on page 31.

Analyzing Performance Issues

MSI-x interrupts are required for RSS to work. If your motherboard and operating system version supports MSI-X, the OneConnect Ethernet driver automatically uses MSI-X interrupts. If there are not enough MSI-X vectors available, the Ethernet driver uses INTx interrupts, which may decrease performance. The `proc` node `/proc/interrupts` shows the interrupts and their types.

The Linux performance utility “top” can monitor the CPU utilization while troubleshooting performance issues. A low idle CPU percentage in any CPU core is an indication of excessive processing load for that CPU. The `proc` node `/proc/interrupts` shows the distribution of the interrupts across the CPU cores. If you see too many interrupts per second directed to one CPU, check to see if the `irqbalance` program is running. The `irqbalance` program is normally started at system boot. In some cases, you can get better performance by disabling `irqbalance` and manually distributing interrupts. You can manually distribute the interrupt load across the available CPU cores by setting the CPU affinity for any interrupt vector by setting the mask in the `proc` node `/proc/irq/<int-vector>/smp_affinity`.

Use the `netstat` command to look for excessive TCP retransmits or packet drops in the network stack.

Use the `-S` option of `ethtool` to see all statistics counters maintained by the OneConnect Ethernet and driver. Excessive drop or error counters are an indication of a bad link or defective hardware. See Table E-1, `Ethtool -S` Option Statistics, on page 160, and Table E-2, Transmit/Receive Queue Statistics, on page 161.

Turning off auditing and SELinux can improve CPU utilization, and in some cases increase throughput. You can disable auditing by appending `audit=0` in the boot command line. You can turn off SELinux by specifying `selinux=0` in the boot command line. For example, the following command boots the Linux kernel with the SELinux and auditing options disabled:

```
kernel /boot/vmlinuz-2.6.18 ro root=/dev/md0 selinux=0 audit=0
```

You can get better CPU utilization, and in some cases better throughput, by disabling kernel debug options such as `CONFIG_DEBUG_SLAB`. This requires you to build the kernel image and modules.

4. Troubleshooting

Situations and their Resolutions

This section explains some of the situations in which your system may operate in an unexpected manner, and some possible resolutions.

FC and FCoE Driver Situations and their Resolutions

Table 4-1 lists the FC and FCoE driver situations and their resolutions.

Table 4-1 FC and FCoE Driver Situations and their Resolutions

Situation	Resolution
FC link fails to come up.	<p>If an FC link fails to come up, verify that an 8 or 16 Gb/s adapter is not attempting to connect to a 1 Gb/s device. Only 2, 4, and 8 Gb/s devices are supported on 8 Gb/s adapters. Only 2, 4, 8, and 16 Gb/s devices are supported on 16 Gb/s adapters.</p> <p>For LP21000 series adapters, ensure the adapter is not in maintenance mode and it is not running the manufacturing firmware.</p>
“Authentication is enabled but authentication service is not running.” Error Message	<p>If you see this message in /var/log/messages and the adapter is in an error state, the fcauthd daemon probably is not running. To determine whether fcauthd is running, run</p> <pre>/etc/init.d/fcauthd status.</pre> <p>To start fcauthd, run</p> <pre>/etc/init.d/fcauthd start.</pre>
If a SAN configuration has 256 targets mapped by the FC and FCoE driver, any additional added targets do not get a target ID mapping by the driver and cause target discovery to fail.	<p>Removing targets or re-initializing the link does not solve this issue.</p> <p>Unload and reload the driver to reset available target IDs. Ensure that the SAN configuration is correct prior to rebooting the driver. This clears the driver’s consistent binding table and frees target IDs for new target nodes.</p>
rmmod fails to unload FC and FCoE driver module due to “ERROR: Module lpfc is in use.”	<p>This message can appear when you attempt to remove the driver and there is a Logical Volume Group dependent on the driver.</p> <p>If you have configured boot from a SAN, you must reboot the system. Otherwise, use these steps to resolved this situation:</p> <ol style="list-style-type: none"> 1) Make the Logical Volume Group unavailable. Type <pre>lvchange -a n xxxxxxxx</pre> <p>The “xxxxxxx” parameter is the Volume Group Name.</p> 2) Stop the OneCommand Manager application. 3) Stop Device Mapper.
rmmod of LPFC driver hangs and module reference count is 0.	<p>Due to a small race condition in the kernel, it is possible for an rmmod command to hang. Issue the “rmmod -w” command. If this does not help, reboot the computer.</p>

Table 4-1 FC and FCoE Driver Situations and their Resolutions (Continued)

Situation	Resolution
rmmod fails to unload driver due to device or resource busy.	<p>This message occurs when you attempt to remove the driver without first stopping the OneCommand Manager application or the fcauthd daemon when the OneCommand Manager application is installed and running, or when FC disks connected to a LightPulse adapter are mounted. To resolve this situation:</p> <ol style="list-style-type: none"> 1) Stop the OneCommand Manager application before attempting to unload the driver. The script is located in the /usr/sbin/ocmanager directory. Type <code>./stop_ocmanager</code> 2) Unmount any disks connected to the adapter. 3) Unload the driver. Type <code>rmmod lpfc</code>
An lspci shows recent Emulex adapters as unknown.	<p>This situation occurs because of the delay of getting new product IDs into the Linux development cycle.</p> <p>There is no resolution at this time.</p>
Slow targets or extended link faults on the storage side may result in storage being marked offline by the mid-level and remaining offline (not recovered) when the link faults are corrected.	<p>This version of the driver should eliminate this issue. However, if you experience offline device issues, increase the SCSI command timeout to a value greater than or equal to 60 seconds. Emulex also provides a script which addresses this issue.</p> <p>To access the lun_change_state.sh script, go to http://www.emulex.com/files/downloads/linux/tools.html.</p>
Under certain conditions of an I/O load, some targets cannot complete an I/O issued by a Linux initiator within the default timeout of 30 seconds given by the SCSI mid-level.	<p>If the situation is not corrected, the initiator-to-target condition deteriorates into abort/recovery storms, leading to I/O failures in the block layer. These types of failures are preceded by a SCSI I/O error of hex 6000000.</p> <p>Emulex provides a script that addresses this issue.</p> <p>To access the set_target_timeout.sh script, go to http://www.emulex.com/files/downloads/linux/tools.html.</p>
The FC or FCoE driver fails to recognize an adapter and logs “unknown IOCB” messages in the system log during driver load.	<p>The adapter is running outdated firmware.</p> <p>Install the latest firmware on the adapter.</p>
Loading the FC and FCoE driver on SLES 10 SPx and SLES 11 SPx reports “unsupported module, tainting kernel” in system log.	<p>This message is logged by the kernel whenever a module that is not shipped with the kernel is loaded.</p> <p>This message can be ignored.</p>
The system panics when it is booted with a failed adapter installed.	Remove the failed adapter and reboot the system.

Table 4-1 FC and FCoE Driver Situations and their Resolutions (Continued)

Situation	Resolution
<p>Unloading the FC and FCoE driver on SLES 10 SPx or SLES 11 SPx may cause a message to be logged in the system log such as the following:</p> <pre>umount: /dev/disk/bypath/pci-0000:02: 04.0-scsi-0:0:1:0: not mounted</pre>	<p>These messages are normal output from the SLES 10 SPx and SLES 11 SPx hotplug scripts and can be safely ignored.</p>
Driver installation fails.	<p>The lpfc-install script fails to install the driver. The install script may fail for the following reasons:</p> <ul style="list-style-type: none"> • A previous version of the driver is installed. Run the lpfc-install --uninstall script and then try to install the driver. • The current driver is already installed. • Run a supported RHEL or SLES kernel.
<p>“No module lpfc found for kernel KERNELVERSION” RPM error message when upgrading the kernel.</p> <p>A recently upgraded kernel cannot find the ramdisk. After upgrading the kernel, the kernel cannot find the ramdisk, which halts or panics the system.</p>	<p>These three situations can be resolved by upgrading the kernel. There are two ways to install the driver into an upgraded kernel. The method you use depends on whether you are updating the driver.</p> <ul style="list-style-type: none"> • Upgrade the kernel using the same version of the driver. • Upgrade the kernel using a new version of the driver. <p>See chapter 2., “Installing and Uninstalling,” on page 14 for these procedures.</p>
The driver is not loaded after a system reboot after upgrading the kernel.	
Driver uninstallation fails.	<p>The lpfc-install --uninstall script fails with an error. Try the following solutions:</p> <ul style="list-style-type: none"> • Uninstall the OneCommand Manager application; see the <i>OneCommand Manager Application User Manual</i> for instructions. • Unmount all FC disk drives. • Unload the lpfcdriver and FC and FCoE driver. • Use rpm -e lpfcdriver and -e ocmanager and uninstall the new kits.
lpfc-install script exit code.	<p>The lpfc-install script contains exit codes that can be useful in diagnosing installation issues. See the lpfc-install script for a complete listing of codes and definitions.</p>

Table 4-1 FC and FCoE Driver Situations and their Resolutions (Continued)

Situation	Resolution
The Emulex driver for Linux does not load in ramdisk for a custom-built kernel.	<p>Custom built kernels are not supported by Emulex. However, the Emulex install script attempts to install the driver into a ramdisk that follows the naming scheme used by Red Hat or SLES kernels.</p> <ul style="list-style-type: none"> • The SLES naming scheme for an Intel Itanium IA64 ramdisk images is: <code>/boot/efi/efi/suse/initrd.</code> • The SLES naming scheme for ramdisk images on all other architectures is: <code>/boot/initrd.</code> <p>If a custom built kernel has a ramdisk image that does not follow the appropriate naming scheme, the name of the image can be changed using the following procedure:</p> <ol style="list-style-type: none"> 1) Change the name of the ramdisk image to match the SLES naming scheme. 2) Update any file links to the ramdisk image. 3) Edit the boot loader configuration file (for example, <code>/etc/lilo.conf</code>, <code>/etc/yaboot.conf</code>, <code>/boot/grub/grub.conf</code>, <code>/boot/grub/menu.lst</code>), find any references to the old ramdisk image name, and replace them with the new name. 4) Reboot the system to verify the changes. 5) Install the Emulex LPFC Linux driver kit.
The Linux SCSI subsystem sees only eight LUNs when more are present.	<p>Some SCSI drivers do not scan past eight LUNs when the target reports itself as a SCSI-2 device.</p> <p>To resolve this situation, force a SCSI bus scan with the following command:</p> <pre>/usr/sbin/lpfc/lun_scan.</pre> <p>SUSE supplies a <code>/bin/rescan-scsi-bus.sh</code> script, which can be changed to scan everything.</p>

Ethernet Driver Situations and their Resolutions

Table 4-2 lists the Ethernet Driver situations and their resolutions for the OneConnect UCNA.

Table 4-2 Ethernet Driver Situations and their Resolutions

Situation	Resolution
The Ethernet driver works but the transmit and receive data rates are not near a 10 Gb/s line rate.	There could be several reasons for poor performance. The driver logs a warning message if the card is found in a suboptimal slot. If you see this message, in /var/log/messages, move the card to the proper slot. For best performance practices, see “Network Performance Tuning” on page 53.
ethtool configuration settings are not restored after system reboot.	The ethtool settings are not designed to persist across reboot. For persistence, configuration commands should be invoked from a boot script that is executed at system start such as /etc/rc.local.
When MILI and SNMP daemons start, they trigger warning messages within SELinux for certain operations.	<p>This is a known issue and no solution is available. However, to avoid SELinux warning messages, you can disable SELinux. To disable SELinux, open a terminal and enter the following command at the prompt:</p> <pre>echo 0 > /selinux/enforce</pre> <p>To enable SELinux, use the following command:</p> <pre>echo 1 > /selinux/enforce</pre> <p>Also see “Analyzing Performance Issues” on page 58.</p>

iSCSI Driver Situations and their Resolutions

Table 4-3 lists the iSCSI driver situations and their resolutions for the OneConnect UCNA.

Table 4-3 iSCSI Driver Situations and their Resolutions

Situation	Resolution
When you log out of a target while an I/O is running and you log into the target again, you will get an error trace in “/var/log/messages” beginning with the error message, “trying to free buffer.”	No solution available.
With SLES11 SP1, you cannot update an iSCSI v2.0 driver to an iSCSI v2.3 or iSCSI v2.4 driver using the ./elx_iscsi_install.sh script provided with the package.	<p>The script indicates the installation is complete, but modinfo or the OneCommand Manager application still displays the older version of the driver. A system reboot does not update the driver either.</p> <p>To work around this issue, follow these steps:</p> <ol style="list-style-type: none"> 1) Determine whether your current iSCSI driver is a v2.0 driver by running the following command: <pre>modinfo be2iscsi</pre> <p>The version 2.0 driver usually has a format of “2.102.xxx.x”.</p> 2) Find the location of the driver module under /lib by typing: <pre>find /lib -name be2iscsi.ko -print</pre> <p>For example, the output may look like this:</p> <pre>/lib/modules/2.6.32.12-0.7- default/updates/be2iscsi.ko /*this is the module file to remove 2.102.348.0*/ /lib/modules/2.6.32.12-0.7- default/kernel/drivers/scsi/be2iscsi/be2 iscsi.ko /*this is the open be2iscsi, do not remove*/</pre> 3) Remove the old module by typing: <pre>rm -f /lib/modules/2.6.32.12-0.7- default/updates/be2iscsi.ko</pre> <p>Install the latest driver via rpm or elx_iscsi_install.sh.</p>

Log Messages

FC and FCoE Driver Log Messages

Retrieving FC and FCoE Driver Log Messages

LPFC error log messages are logged in the `/var/log/messages` file.

An example of an LPFC message:

```
Jul 2 04:23:34 daffy kernel: lpfc 0000:03:06.0: 0:1305 Link Down
Event x2f2 received Data: x2f2 x20 x110
```

In this example:

- `lpfc 0000:03:06.0` – identifies the PCI location of the particular LPFC hardware port.
- `0:` – indicates Emulex adapter 0
- `1305` – indicates a log message number of 1305.

Note: If “Data:” is present in a log message, any information following “Data:” is intended for Emulex technical support/engineering use only.

LPFC Error Log Messages and their Descriptions

Table 4-4 lists LPFC error log messages and their descriptions.

Table 4-4 LPFC Error Log Messages and their Descriptions

0111: Dropping received ELS cmd
The driver decided to drop an ELS Response ring entry.
Data: (1) ulpStatus (2) ulpWord[4] (3) ulpTimeout
Severity: Error
Log: Always
Action: This error could indicate a software driver or firmware issue. If this issue persists, report these errors to Technical Support.
0113: An FLOGI ELS command <elsCmd> was received from DID <did> in Loop Mode
While in Loop Mode an unknown or unsupported ELS command was received.
Data: None
Severity: Error
Log: Always
Action: Check device DID.
0115: Unknown ELS command <elsCmd> received from NPORT <did>
Received an unsupported ELS command from a remote N_Port.
Data: None
Severity: Error
Log: Always
Action: Check the remote N_Port for a potential issue.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0125: FDISC Failed (x%x). Fabric out of resources
The fabric rejected an FDISC because the switch cannot support additional virtual ports.
Data: lsRjtError
Severity: Error
Log: Always
Action: Reconfigure the switch to support more NPIV logins. If this issue persists, contact Technical Support.

0126: FDISC failed ulpStatus ulpWord4
Data: lsRjtError
Severity: Error
Log: Always
Action: Reconfigure the switch to support more NPIV logins. If this issue persists, contact Technical Support.

0127: ELS timeout
An ELS IOCB command was posted to a ring and did not complete within ULP timeout seconds.
Data: (1) elscmd (2) remote_id (3) ulpcommand (4) ulploTag
Severity: Error
Log: Always
Action: If no ELS command is going through the adapter, reboot the system; If the issue persists, contact Technical Support.

0133: PLOGI: no memory for reg_login
Memory allocation error.
Data: (1) nlp_DID (2) nlp_state (3) nlp_flag (4) nlp_rpi
Severity: Error
Log: LOG_ELS
Action: Memory allocation error. Check system resources. Unload unused modules.

0134: PLOGI cannot issue reg_login
The ELS PLOGI mailbox command has failed.
Data: (1) nlp_DID (2) nlp_state (3) nlp_flag (4) nlp_rpi
Severity: Error
Log: LOG_ELS
Action: Check the port and switch configuration.

0135: cannot format reg_login
Could not allocate an RPI or DMA buffer for the mailbox command.
Data: (1) nlp_DID (2) nlp_state (3) nlp_flag (4) nlp_rpi
Severity: Error
Log: LOG_ELS
Action: None required.

0136: PLOGI completes to NPort <DID> completion
A PLOGI has completed for which there is no NDLP.
Data: (1) ulpStatus (2) ulpWord[4]
Severity: Error
Log: LOG_ELS
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0137: No retry ELS command <ELS_CMD> to remote
Data: (1) ulpStatus (2) ulpWord[4]
Severity: Error
Log: LOG_ELS
Action: None required.
0138: ELS rsp: Cannot issue reg_login for <DID>
REG_LOGIN mailbox command failed.
Data: (1) nlp_DID (2) nlp_state (3) nlp_flag (4) nlp_rpi
Severity: Error
Log: LOG_ELS
Action: None required.
0139: Ignoring ELS cmd tag <ioTag> completion Data
This ELS command was aborted.
Data: (1) ulpStatus (2) ulpWord[4] (3) ulpTimeout
Severity: Error
Log: LOG_ELS
Action: None required.
0140: PLOGI Reject: invalid name
Invalid node WWN provided.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
0141: PLOGI Reject: invalid pname
Invalid port WWN provided.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
0142: PLOGI RSP: Invalid WWN
The PLOGI sent to the port by a remote port had an invalid WWN.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
0143: SLI4 Adapter Hardware Error Data: <status0>/<status1>
The HBA has encountered an unrecoverable error.
Data: None
Severity: Error
Log: LOG_INIT
Action: Use hbacmd to retrieve a dump file.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0144: Not a valid WCQE code: <Completion Code>
The completion queue handler detected an invalid type.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.
0147: Failed to allocate memory for RSCN event
Memory could not be allocated to send the RSCN event to the management application.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
0148: Failed to allocate memory for LOGO event
Memory could not be allocated to send the LOGO event to the FC transport.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
0149: Failed to allocate memory for ELS event
Memory could not be allocated to send the ELS event to the FC transport.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
0154: Authentication not complete
Authentication was restarted because the previous authentication did not complete.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: Check the switch configuration.
0200: CONFIG_LINK bad hba state <hba_state>
A CONFIG_LINK mailbox command completed and the driver was not in the right state.
Data: None
Severity: Error
Log: Always
Action: Software driver error. If this issue persists, report these errors to Technical Support.
0203: Devloss timeout on WWPN <address> NPort <nlp_DID>
A remote N_Port that was discovered by the driver disappeared for more than lpfc_devloss_tmo seconds.
Data: (1) nlp_flag (2) nlp_state (3) nlp_rpi
Severity: Error
Log: Always
Action: If the device generating this message is not a target to which the HBA is connected, this error will not affect the data integrity of the I/O between the HBA and the attached storage and can be ignored.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0206: Device discovery completion error
This indicates that an uncorrectable error was encountered during device (re)discovery after a link up. FC devices will not be accessible if this message is displayed.
Data: None
Severity: Error
Log: Always
Action: Reboot the system. If this issue persists, report the error to Technical Support. Run with verbose mode enabled for more information.

0207: Device <DID> (<WWN>) sent invalid service parameters. Ignoring device.
Invalid service parameters were received from DID. Ignoring this remote port.
Data: DID, WWN
Severity: Error
Log: Always
Action: Verify the remote port's configuration. If this issue persists, report the error to Technical Support. Run with verbose mode on for more details.

0217: Block sgl registration required DMAsize <reqlen> great than a page
The request to post SGL pages does not fit on a page.
Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.

0221: FAN timeout
A link up event was received without the login bit set, so the driver waits E_D_TOV for the Fabric to send a FAN. If no FAN is received, a FLOGI will be sent after the timeout.
Data: None
Severity: Warning
Log: LOG_DISCOVERY verbose
Action: None required. The driver recovers from this condition by issuing a FLOGI to the fabric.

0222: Initial FLOG/FDISKI timeout
The driver sent the initial FLOGI or FDISK to the fabric and never got a response back.
Data: None
Severity: Error
Log: Always
Action: Check Fabric configuration. The driver recovers from this and continues with device discovery.

0223: Timeout while waiting for NameServer login
Our login request to the NameServer was not acknowledged within R_A_TOV.
Data: None
Severity: Error
Log: Always
Action: Check the fabric configuration. The driver recovers from this and continues with device discovery.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0224: NameServer Query timeout
Node authentication timeout, node Discovery timeout. A NameServer Query to the Fabric or discovery of reported remote N_Ports is not acknowledged within R_A_TOV.
Data: (1) fc_ns_retry (2) fc_max_ns_retry
Severity: Error
Log: Always
Action: Check Fabric configuration. The driver recovers from this and continues with device discovery.

0226: Device discovery completion error
This indicates that an uncorrectable error was encountered during device (re)discovery after a link up. FC devices will not be accessible if this message is displayed.
Data: None
Severity: Error
Log: Always
Action: Reboot the system. If this issue persists, report the error to Technical Support. Run with verbose mode on for more details.

0227: Node Authentication timeout
The driver has lost track of what N_Ports are being authenticated.
Data: None
Severity: Error
Log: Always
Action: None required. The driver should recover from this event.

0228: CLEAR LA timeout
The driver issued a CLEAR_LA that never completed.
Data: None
Severity: Error
Log: Always
Action: None required. The driver should recover from this event.

0230: Unexpected timeout, hba linkstate <link_state>
Discovery has timed out and the HBA state is not ready.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: None required.

0231: RSCN timeout
The driver has lost track of what N_Ports have RSCNs pending.
Data: (1) fc_ns_retry (2) lpfc_max_ns_retry
Severity: Error
Log: Always
Action: None required. The driver should recover from this event.

0233: Nodelist not empty
Driver unloaded or hotplug detected a node still in use.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0237: Pending Link Event during Discovery: State <hba_state>
Received link event during discovery. Causes discovery restart.
Data: None
Severity: Warning
Log: LOG_DISCOVERY verbose
Action: None required, unless this issue persists. If persistent, check cabling.
0241: NameServer rsp error
The driver received a NameServer response containing a status error.
Data: (1) CommandResponse.bits.CmdRsp (2) ReasonCode (3) Explanation (4) fc_flag
Severity: Error
Log: LOG_DISCOVERY verbose
Action: Check the fabric configuration. The driver recovers from this and continues with device discovery.
0246: RegLogin failed
The firmware returned a failure for the specified RegLogin.
Data: (1) Did (2) mbxStatus (3) hbaState
Severity: Error
Log: Always
Action: This message indicates that the firmware could not do RegLogin for the specified DID. There may be a limitation on how many nodes an HBA can see.
0249: Cannot issue Register Fabric login: Err %d\
Could not issue the fabric reg login, the err value is unique for each possible failure.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
0251: NameServer login: no memory
Could not allocate memory for the NDLP structure.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
0252: Cannot issue NameServer login
Could not issue an ELS PLOGI to the NameServer DID.
Data: None
Severity: Error
Log: LOG_ELS
Action: Check the port connection and the switch configuration.
0253: Register VPI: Can't send mbox\
Could not issue the REG_LOGIN command for this VPort.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0254: Register VPI: no memory" goto mbox_err_exit
Could not allocate memory for the REG_LOGIN mailbox command.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

0255: Issue FDISC: no IOCB
All of the pre-allocated IOCBs are in use.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.

0256: Issue FDISC: Cannot send IOCB\
Unable to send the fabric IOCB.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.

0257: GID_FT Query error
The GID_FT CT request for the NameServer has failed.
Data: None
Severity: Error
Log: LOG_ELS
Action: Check the switch configuration.

0258: Register Fabric login error:
The REG_LOGIN for the fabric has failed.
Data: None
Severity: Error
Log: LOG_MBOX
Action: Check the port connection and the switch configuration.

0259: No NPIV Fabric support
The switch to which the port is connected does not support NPIV.
Data: None
Severity: Error
Log: LOG_ELS
Action: Check the switch configuration.

0260: Register NameServer error:
The REG_LOGIN mailbox command has failed for the NameServer.
Data: None
Severity: Error
Log: LOG_ELS
Action: Check the switch configuration

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0261: Cannot register NameServer login:
Either a memory allocation issue or an invalid parameter was sent to the REG_LOGIN.
Data: None
Severity: Error
Log: LOG_ELS
Action: At least one message (0142 0121 0133 0134 0135) should precede this message.

0262: No NPIV Fabric support
The switch to which the port is connected does not support NPIV.
Data: None
Severity: Error
Log: LOG_ELS
Action: Check the switch configuration.

0263: Discovery Mailbox error: state:
Either the driver could not allocate resources or it could not send sparam_mbox or cfglink_mbox.
Data: (1) address of sparam_mbox command (2) address of cfglink_mbox command
Severity: Error
Log: LOG_MBOX
Action: Attempt to unload and reload the driver when it is convenient.

0264: No NPIV Fabric support
The switch to which the port is connected does not support NPIV.
Data: None
Severity: Error
Log: LOG_ELS
Action: Check the switch configuration.

0266: Issue NameServer Req <cmdcode> err <rc> Data: <fc_flag> <fc_rscn_id_cnt>
The driver was unable to send the NameServer CT command.
Data: (1) vports fc_flag (2) vports fc_rscn_id_cnt
Severity: Error
Log: LOG_DISCOVERY
Action: Check the port and switch configurations.

0267: NameServer GFF Rsp <did> Error (<ulpStatus> <un.ulpWord[4]>) Data: <fc_flag> <fc_rscn_id_cnt>
The NameServer GFF CT request failed.
Data: (1) vports fc_flag (2) vports fc_rscn_id_cnt
Severity: Error
Log: LOG_DISCOVERY
Action: Check the port and switch configurations.

0268: NS cmd <cmdcode> Error (<ulpStatus> <un.ulpWord[4]>)
The NameServer CT request failed.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: Check the port and switch configurations.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0271: Illegal State Transition: node <nlp_DID> event <evt>, state <nlp_state> Data: <nlp_rpi> <nlp_flag>
The current node state does not have a handler for this event.
Data: (1) nlp_rpi (2) nlp_flag
Severity: Error
Log: LOG_DISCOVERY
Action: Verify that all targets are still visible to the SCSI mid-layer.

0272: Illegal State Transition: node <nlp_DID> event <evt>, state <nlp_state> Data: <nlp_rpi> <nlp_flag>
The driver is completing a PLOGI but do not have the rcv_plogi flag set.
Data: (1) nlp_rpi (2) nlp_flag
Severity: Error
Log: LOG_DISCOVERY
Action: Verify that all targets are still visible to the SCSI mid-layer.

0273: Unexpected discovery timeout, vport State x%x
The discovery process has timed out.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: Verify that all targets are still visible.

0274: lpfc_nlp_put: ndlp:x%pusgmap:x%x refcnt:%d, void *)ndlp, ndlp->nlp_usg_map, atomic_read(&ndlp->kref.refcount)
Data: None
Severity: Warning
Log: LOG_NODE
Action: None required.

0275: lpfc_nlp_put: ndlp:x%pusgmap:x%x refcnt:%d, void *)ndlp, ndlp->nlp_usg_map, atomic_read(&ndlp->kref.refcount)
A kref_put was called again after the node was already inactive.
Data: None
Severity: Warning
Log: LOG_NODE
Action: None required.

0276: lpfc_nlp_get: ndlp:x%pusgmap:x%x refcnt:%d, void *)ndlp, ndlp->nlp_usg_map, atomic_read(&ndlp->kref.refcount)
A kref_get was attempted on a node that was being released.
Data: None
Severity: Warning
Log: LOG_NODE
Action: None required.

0277: lpfc_enable_node: ndlp:x%pusgmap:x%x refcnt:%d, void *)ndlp, ndlp->nlp_usg_map, atomic_read(&ndlp->kref.refcount)
Enable node was attempted on an inactive node.
Data: None
Severity: Warning
Log: LOG_NODE
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0278: lpfc_enable_node: ndlp:x%pusgmap:x%x refcnt:%d, void *)ndlp, ndlp->nlp_usg_map, atomic_read(&ndlp->kref.refcount)
Enable node was attempted on an inactive node.
Data: None
Severity: Warning
Log: LOG_NODE
Action: None required.

0280: lpfc_cleanup_node: ndlp:x%pusgmap:x%x refcnt:%d, void *)ndlp, ndlp->nlp_usg_map, atomic_read(&ndlp->kref.refcount)
Node clean-up was attempted on a node that has already been marked for memory free.
Data: None
Severity: Warning
Log: LOG_NODE
Action: None required.

0281: lpfc_cleanup_node: ndlp:x%pusgmap:x%x refcnt:%d, void *)ndlp, ndlp->nlp_usg_map, atomic_read(&ndlp->kref.refcount)
Node clean-up was called to prepare the node for release.
Data: None
Severity: Warning
Log: LOG_NODE
Action: None required.

0282: ldid:x%x ndlp:x%pusgmap:x%x refcnt:%d, ndlp->nlp_DID, (void *)ndlp, lpfc_init.c-ndlp->nlp_usg_map,
Driver clean-up has found a node that is still on the node list during driver unload or PCI hotplug removal.
Data: None
Severity: Error
Log: LOG_NODE
Action: None required.

0283: Failed to allocate mbox cmd memory
Mailbox allocation error.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0285: Allocated DMA memory size <alloclen> is less than the requested DMA memorysize<reqlen>
Memory allocation was truncated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0286: lpfc_nlp_state_cleanup failed to allocate statistical data buffer <nlp_DID>	Memory allocation failed for node's statistical data. Data: None Severity: Error Log: LOG_INIT Action: None required.
0287: lpfc_alloc_bucket failed to allocate statistical data buffer DID <nlp_DID>	Memory allocation failed for node's statistical data. Data: None Severity: Error Log: LOG_NODE Action: None required.
0288: Unknown FCoE event type <event_type> event tag <event_tag>	The firmware has detected an unknown FCoE event. Data: None Severity: Error Log: LOG_SLI Action: Check the FCoE switch configuration and the HBA DCBX mode.
0289: Issue Register VFI failed: Err <rc>	The driver could not register the Virtual Fabric Index for the FCFL. Data: None Severity: Error Log: LOG_ELS Action: Check the switch and port configurations.
0290: The SLI4 DCBX asynchronous event is not handled yet	The SLI-4 DCBX asynchronous event is not handled yet. Data: None Severity: Error Log: LOG_SLI Action: None required.
0291: Allocated DMA memory size (x%x) is less than the requested DMA memory size (x%x)	The asynchronous DCBX events are not handled in the driver. Data: None Severity: Error Log: LOG_INIT Action: Check the switch configuration.
0293: PM resume failed to start worker thread: error=<error>	The PCI resume (hotplug) could not start the worker thread for the driver. Data: None Severity: Error Log: LOG_INIT Action: Unload and reload the driver.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0294: PM resume Failed to enable interrupt
The PCI resume (hotplug) could not get an interrupt vector.
Data: None
Severity: Error
Log: LOG_INIT
Action: Unload and reload the driver.
0297: Invalid device group<pci_dev_grp>
While unloading the driver, the driver detect a PCI device that it should not have claimed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0299: Invalid SLI revision <sli_rev>
While processing a host attention or unrecoverable error, the driver detected an invalid SLI revision.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0300: LATT: Cannot issue READ_LA: Data:<rc>
The link attention handler could not issue a READ_LA mailbox command.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.
0301: READ_SPARAM: no buffers
The driver attempted to issue a READ_SPARAM mailbox command to the adapter, but there were no buffers available.
Data: None
Severity: Warning
Log: LOG_MBOX verbose
Action: This message indicates: (1) Kernel virtual memory is depleted. Check that the system meets minimum RAM requirements for the Emulex FC adapter. Try closing other applications to free some memory. (2) A possible driver buffer management issue. If this issue persists, report the error to Technical Support.
0302: REG_LOGIN: no buffers
The driver attempted to issue a REG_LOGIN mailbox command to the adapter, but there no buffers were available.
Data: (1) Did, (2) flag
Severity: Warning
Log: LOG_MBOX verbose
Action: This message indicates: (1) Kernel virtual memory is depleted. Check that the system meets minimum RAM requirements for the Emulex FC adapter. Try closing other applications to free some memory. (2) A possible driver buffer management issue. If this issue persists, report the error to Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0313: Ring <ringno> handler: unexpected Rctl <Rctl> Type <Type> received
The RCTL/Type of a received frame did not match any for the configured masks for the specified ring.
Data: None
Severity: Warning
Log: LOG_SLI verbose
Action: This error could indicate a software driver, firmware, or hardware issue. Report these errors to Technical Support.

0303: Ring <ringno> handler: portRspPut <portRspPut> is bigger then rsp ring <portRspMax>
The port rsp ring put index is larger than the size of the rsp ring.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a software driver, firmware or hardware issue. Report these errors to Technical Support.

0304: Stray mailbox interrupt, mbxCommand <mbxcommand> mbxStatus <mbxstatus>
Received a mailbox completion interrupt and there are no outstanding mailbox commands.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0306: CONFIG_LINK mbxStatus error <mbxStatus> HBA state <hba_state>
The driver issued a CONFIG_LINK mailbox command to the HBA that failed.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a firmware or hardware issue. Report these errors to Technical Support.

0310: Mailbox command <mbxcommand> timeout
A mailbox command was posted to the adapter and did not complete within 30 seconds.
Data: (1) hba_state (2) sli_flag (3) mbox_active
Severity: Error
Log: Always
Action: This error could indicate a software driver or firmware issue. If no I/O is going through the adapter, reboot the system. If this issue persists, report the error to Technical Support.

0312: Ring <ringno> handler: portRspPut <rspPutInx> is bigger then rsp ring <numRiocb>
The IOCB command rings put pointer is ahead of the get pointer.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0313: Ring <ringno> handler: unexpected Rctl <Rctl> Type <Type> received
The RCTL/Type of a received frame did not match any for the configured masks for the specified ring.
Data: None
Severity: Warning
Log: LOG_SLI verbose
Action: This error could indicate a software driver, firmware or hardware issue. Report these errors to Technical Support.

0315: Ring <ringno> issue: portCmdGet <local_getidx> is bigger then cmd ring <max_cmd_idx>
The port command ring get index is greater than the size of the command ring.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a software driver, firmware or hardware issue. Report these errors to Technical Support.

0317: iotag <ulp_IoTag> is out of range: max iotag <max_iotag> wd0 <wd0>
The IoTag in the completed IOCB is out of range.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a software driver, firmware or hardware issue. Report these errors to Technical Support.

0318: Failed to allocate IOTAG. last IOTAG is <last_allocated_iotag>
The driver cannot allocate an IoTag. Display the last value used.
Data: None
Severity: Error
Log: Always
Action: This message indicates the adapter HBA I/O queue is full. Typically this happens when heavy I/O is running on a low-end (3 digit) adapter. We suggest you upgrade to a higher-end adapter.

0319: READ_SPARAM mbxStatus error <mbxStatus> hba state <hba_state>
The driver issued a READ_SPARAM mailbox command to the HBA that failed.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a firmware or hardware issue. Report these errors to Technical Support.

0320: CLEAR_LA mbxStatus error <mbxStatus> hba state <hba_state>
The driver issued a CLEAR_LA mailbox command to the HBA that failed.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a firmware or hardware issue. Report these errors to Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0322: Ring <ringno> handler: unexpected completion IoTag <IoTag>
The driver could not find a matching command for the completion received on the specified ring.
Data: (1) ulpStatus, (2) ulpWord[4], (3) ulpCommand, (4) ulpContext
Severity: Warning
Log: LOG_SLI verbose
Action: This error could indicate a software driver or firmware issue. If this issue persists, report these errors to Technical Support.

0323: Unknown Mailbox command <mbxCommand> Cmpl
A unknown mailbox command completed.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a software driver, firmware or hardware issue. Report these errors to Technical Support.

0324: Config port initialization error, mbxCmd <mbxCommand> READ_NVPARM, mbxStatus <mbxStatus>
READ_NVPARMS mailbox command failed during port configuration.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a software driver, firmware or hardware issue. Report these errors to Technical Support.

0328: Rsp Ring <ring number> error: IOCB Data:
The firmware has returned an error for this IOCB.
Data: (1) <iocb word[0]:iocb word[7]>, (2) <rsp word[0]:rsp[word[7]>
Severity: Warning
Log: LOG_SLI
Action: None required.

0330: IOCB wake NOT set
The completion handler associated with the IOCB was never called.
Data: (1) timeout (2) timeleft/jiffies
Severity: Error
Log: Always
Action: This error could indicate a software driver, firmware or hardware issue. If this issue persists, report the error to Technical Support.

0334: Unknown IOCB command
Received an unknown IOCB command completion.
Data: (1) type (2) ulpCommand (3) ulpStatus (4) ulploTag (5) ulpContext)
Severity: Error
Log: Always
Action: This error could indicate a software driver or firmware issue. If this issue persists, report these errors to Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0335: Unknown IOCB command
Received an unknown IOCB command completion.
Data: (1) ulpCommand (2) ulpStatus (3) ulploTag (4) ulpContext)
Severity: Error
Log: Always
Action: This error could indicate a software driver or firmware issue. If this issue persists, report these errors to Technical Support

0336: Rsp Ring <ringno> error: IOCB
An IOCB error has occurred on the specified ring.
Data: (1) ulpWord[0], (2) ulpWord[1], (3) ulpWord[2], (4) ulpWord[3], (5) ulpWord[4], (6) ulpWord[5], (7) irsp+6, (8) irsp+7
Severity: Warning
Log: LOG_SLI verbose
Action: If this issue persists, check the targets. If the targets are okay, report the error to Technical Support.

0340: Adapter temperature is OK now
Adapter temperature has reverted to normal range.
Data: Temperature in Celsius
Severity: Error
Log: LOG_TEMP verbose
Action: No action needed, informational

0341: Ring <ringno> Cannot find buffer for an unsolicited iocb tag <un.ulpWord[3]>
There are no more pre-allocated buffers available to handle unsolicited buffers.
Data: None
Severity: Error
Log: LOG_SLI
Action: Ensure this port is not being managed by multiple ports.

0342: Ring <ringno> Cannot find buffer for an unsolicited iocb tag <unsl3.sli3Words>
This is a multiple IOCB unsolicited command and sufficient buffer space cannot be allocated for it.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0343: Ring <ringno> Cannot find buffer for an unsolicited iocb tag <un.ulpWord[3]>
There are no more pre-allocated buffers available to handle unsolicited buffers.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0344: Ring <ringno> Cannot find buffer for an unsolicited iocb tag <unsl3.sli3Words[7]>
There are no more pre-allocated buffers available to handle unsolicited buffers.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0345: Resetting board due to mailbox timeout iocb. tag 0x%x
A mailbox command failed to complete. The driver is resetting the port.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: If the mailbox command fails again, set the lpfc_log_verbose to LOG_MBOX and retry.

0346: Ring <ring number> handler: unexpected ASYNC_STATUS evt_code <evt code> W0 <hex w0> W1 <hex w1> W2 <hex w2> W3 <hex w3> W4 <hex w4> W5 <hex w5> W6 <hex w6> W7 <hex w7> W8 <hex w8> W9 <hex w9> W10 <hex w10> W11<hex w11>
The HBA received an asynchronous event that was not a temperature event.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0347: Adapter is very hot, please take corrective action
Adapter temperature is above normal range.
Data: Temperature in Celsius
Severity: Error
Log: LOG_TEMP verbose
Action: Shutdown and remove the HBA. Contact Technical Support.

0348: NameServer login: node freed
The enable mode failed to free up the NameServer login.
Data: None
Severity: Error
Log: LOG_ELSI
Action: None required.

0349: rc should be MBX_SUCCESS
The next mailbox command on the mailbox queue has failed.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

0350: rc should have been MBX_BUSY
Attempting to unregister a default RPI from an interrupt context and the mailbox state is not busy.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0351: Config MSI mailbox command failed, mbxCmd <u.mb.mbxComm>, mbxStatus <u.mb.mbxStatus>
The mailbox command sent to the firmware to configure the adapter to use MSI-X has failed.
Data: None
Severity: Warning
Log: LOG_MBOX
Action: Ensure the hardware platform supports MSI-X.

0352: Config MSI mailbox command failed, mbxCmd <u.mb.mbxCommand>, mbxStatus <u.mb.mbxStatus>
The mailbox command sent to the firmware to configure the HBA to use MSI-X has failed.
Data: None
Severity: Error
Log: LOG_MBOX
Action: Ensure the hardware platform supports MSI-X.

0353: Active Mailbox cleared - mailbox timeout exiting
The mailbox timeout handler has determined that the driver is in the process of completing this mailbox command.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

0357: MSI-X interrupt with no EQE
SLI-4 adapter interrupt on the slow path but there is no associated EQE.
Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

0358: MSI-X interrupt with no EQE
SLI-4 adapter interrupt on the fast path but there is no associated EQE.
Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

0359:Not a valid slow-path completion " event: majorcode=x%x, minorcode=x%x\n", bf_get(lpfc_eqe_major_code, eqe), bf_get(lpfc_eqe_minor_code, eqe));
SLI-4: The EQE is invalid.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0360:Unsupported EQ count. <entry_count>
Cannot create an event queue of this size.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0361:Unsupported CQ count. <entry_count>
Cannot create an completion queue of this size.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0362:Unsupported MQ count. <entry_count>
Cannot create MQ of this size.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0364:Invalid param:
SLI-4: The post-SGL function was passed an invalid XRI.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0365:Slow-path CQ identifier <CQID> does not exist:
The Completion Queue ID passed in the event queue entry does not reference a valid completion queue.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0366:Not a valid fast-path completion event: majorcode=<major code hex>, minor-code=<minor code hex>
The major or minor code in the Event Queue field is invalid.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0367: Fast-path completion queue does not exist
The fast path completion queue referenced by the CQID does not exist.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0368: Mis-matched fast-path completion queue identifier: eqcqid=%d, fcpcqid=%d
The CQID in the event queue entry does not match the fcp_cqid that was passed into the routine.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0369: No entry from fast-path completion queue fcpcqid=<queue_id>
There were no completions in the completion queue referenced by fcp_cqid.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0370: Invalid completion queue type <type>
The event queue entry is not for a mailbox or a work queue entry.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0371: No entry from the CQ: identifier <queue_id>, type <type>
There was no completion queue event for this event queue entry.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0372: iotag <iotag> is out of range: max iotag (<sli.last_iotag>)
The IOCB lookup cannot be performed because the iocb_tag is out of range.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0373: FCP complete error: status=<status> hw_status=<hw status>, total_data_specified=<total data transferred>, parameter=<rsp word[4]>, word3=<wcqe word 3>
Logs the FCP failure. Status and parameter are equivalent to ulpStatus and ulpWord[4].
Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

0374: FCP complete with no corresponding cmdiocb: iotag <iocb iotag>
There was no IOCB on the in-progress list that matched this iotag.
Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

0375: FCP cmdiocb not callback function iotag: <iocb iotag>
The IOCB found for this iotag does not have a completion handler set in it.
Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0377: Error <rc> parsing vpd. Using defaults. Could not parse the VPD data, so the driver is using the default values. Data: None Severity: Error Log: Always Action: None required.
0378: No support for fcpi mode. Could not configure the port to run in FCP initiator mode. Data: None Severity: Warning Log: LOG_MBOX, LOG_SLI Action: None required.
0379: Feature Mismatch Data: <req_ftr word2 hex> <req_ftr word3 hex> <cfg_enable_npiv> <max vpi hex> The features passed in to the driver as module parameters do not match what the firmware can do. Setting to default values. Data: None Severity: Warning Log: LOG_MBOX, LOG_SLI Action: None required.
0381: Error %d during queue setup. Could not set up all the queues that the driver requires to exchange I/Os with the HBA. Data: None Severity: Error Log: LOG_MBOX, LOG_SLI Action: Reload the driver.
0382: READ_SPARAM command failed status <issue status>, mbxStatus <mailbox status> The READ_SPARAM mailbox command has failed during initialization. The HBA has been set to error state. Data: None Severity: Error Log: LOG_MBOX, LOG_SLI: Action: Take a dump with hbacmd and then try reloading the driver.
0383: Error <rc> during scsi sgl post operation The SGL entries could not be registered with the adapter. Data: None Severity: Warning Log: LOG_MBOX, LOG_SLI Action: Reset the adapter using hbacmd.
0384: There is pending active mailbox cmd The mailbox commands have overlapped. This command should have been added to the mailbox queue. Data: None Severity: Error Log: LOG_MBOX, LOG_SLI Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0385: rc should have been MBX_BUSY
The completion handler for REG_LOGIN detected the IMMED_UNREG flag and tried to issue the UNREG_LOGIN command from an interrupt level. The mailbox status should still be busy.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.
0386: ELS complete with no corresponding cmdiocb: iotag <iotag>
The completion that the ISR is handling cannot find a tag associated with the IOTAG.
Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.
0387:Failed to allocate an iocbq
Failed to get an IOCBQ from the list of available IOCBQs.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.
0388:Not a valid WCQE code: x<hex cqe_code>
The event code is invalid. This event will be dropped.
Data: None
Severity: Error
Log: LOG_SLI
Action: Ensure the adapter's firmware is current.
0391:Error during rpi post operation
The driver was trying to post pages to the firmware to be used to keep target login information and encountered a failure.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: Unload and reload the driver.
0393:Error <rc> during rpi post operation
The driver was trying to post pages to the firmware to keep target login information and encountered a failure.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: Unload and reload the driver.
0394: Failed to allocate CQ_EVENT entry
The asynchronous event handler was unable to allocate an event queue entry to which to transfer the asynchronous event.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: This could be a V-LINK clear from the switch or a fatal error from the firmware. Perform a dump from the OneCommand Manager application.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0395: The mboxq allocation failed
The asynchronous link event handler could not allocate a mailbox command to issue the READ_LA (read link attention) mailbox command.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0396:The lpfc_dmabuf allocation failed
The asynchronous link event handler could not allocate a mailbox command to issue the READ_LA mailbox command.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0397:The mbuf allocation failed
The asynchronous link event handler could not allocate DMA-able memory for the READ_LA mailbox command.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

0398 Invalid link fault code: <hex link_fault>
The attempt to read the link attention register has returned an unknown value.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0399 Invalid link attention type: <hex link_type>
The READ_LA mailbox command has returned an invalid link type.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required

0400: lpfc_nODEV_tmo attribute cannot be set to <val>, allowed range is [<LPFC_MIN_DEVLOSS_TMO>, <LPFC_MAX_DEVLOSS_TMO>
The attempt to set the devloss timeout value failed because the value is out of the allowable range.
Data: None
Severity: Error
Log: LOG_INIT
Action: Use a value between the minimum and maximum values.

0401: Ignoring change to nodev_tmo because devloss_tmo is set
Attempting to change the nodev timeout when the devloss has already been set.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0402: Cannot find virtual addr for buffer tag on ring <ringno>
A DMA buffer is unavailable for this unsolicited command.
Data: (1) tag (2) next (3) prev (4) postbufq_cnt
Severity: Error
Log: LOG_INIT
Action: None required.

0403: lpfc_nodev_tmo attribute cannot be set to <val>, allowed range is [<code><LPFC_MIN_DEVLOSS_TMO></code> , <code><LPFC_MAX_DEVLOSS_TMO></code>]
Attempt to set the nodev timeout value is outside the range of the devloss timeout range.
Data: None
Severity: Error
Log: LOG_INIT
Action: Set the nodev timeout between the minimum and maximum timeout range.

0404: lpfc_devloss_tmo attribute cannot be set to <val>, allowed range is [<code><LPFC_MIN_DEVLOSS_TMO></code> , <code><LPFC_MAX_DEVLOSS_TMO></code>]
Attempt to set the devloss timeout value is outside the allowed range.
Data: None
Severity: Error
Log: LOG_INIT
Action: Set the devloss timeout between the minimum and maximum devloss range.

0405: lpfc_link_speed attribute cannot be set to %d, allowed values are [<code>"LPFC_LINK_SPEED_STRING"</code>]
Attempt to set the link speed value is outside the allowed range.
Data: None
Severity: Error
Log: LOG_INIT
Action: Set the link speed between 0 and the maximum.

0406: Adapter maximum temperature exceeded <temperature>, taking this port offline
The driver has received an error for the HBA indicating that the maximum allowable temperature has been exceeded.
Data: (1) work_hs (2) work_status[0] (3) work_status[1]
Severity: Error
Log: LOG_INIT
Action: Ensure the server fans are not blocked. Shut down the server if the airflow is restricted.

0407: Ignoring nodev_tmo module parameter because devloss_tmo is set.
Both module parameters (nodev and devloss) were set so the driver is ignoring the nodev parameter.
Data: None
Severity: Error
Log: LOG_INIT
Action: Only one of these parameters must be set.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0410: Cannot find virtual addr for mapped buf on ring <ringno>
The driver cannot find the specified buffer in its mapping table. Therefore, it cannot find the virtual address needed to access the data.
Data: (1) phys (2) next (3) prev (4) postbufq_cnt
Severity: Error
Log: Always
Action: This error could indicate a software driver or firmware issue. If this issue persists report these errors to Technical Support.

0421: MSI-X slow-path request_irq failed <rc>
The kernel API to request an IRQ has failed.
Data: None
Severity: Warning
Log: LOG_INIT
Action: Use module parameter lpfc_use_msi=0 (INTx).

0422: lpfc_restrict_login attribute cannot be set to <val>, allowed range is [0, 1]
Attempt to set the restrict login parameter to something other than on or off.
Data: None
Severity: Error
Log: LOG_INIT
Action: Use 0 (Off) or 1 (On)

0423: lpfc_"#attr" attribute cannot be set to %d, allowed range is ["#minval", "#maxval"]
This is a compile time macro that is used by several module parameters during initialization. Each module parameter has its own minimum and maximum values that are displayed.
Data: None
Severity: Error
Log: LOG_INIT
Action: Set the module parameter between the minimum and maximum values.

0424:lpfc_"#attr" attribute cannot be set to %d, allowed range is ["#minval", "#maxval"]
This is a compile time macro that is used by several module parameters to set the value.
Data: None
Severity: Error
Log: LOG_INIT
Action: Set the module parameter between the minimum and maximum values.

0425:lpfc_restrict_login attribute cannot be set to %d, allowed range is [0, 1]
The module parameter lpfc_restrict_login can only be set to 0 (off) or 1 (on).
Data: None
Severity: Error
Log: LOG_INIT
Action: Set lpfc_restrict_login=[0,1].

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0426: Failed to enable interrupt
The driver failed to start the interrupt.
Data: None
Severity: Error
Log: LOG_INIT
Action: Unload and reload the driver.

0427: Cannot re-enable interrupt after slot reset
The driver was unable to enable the interrupt after an HBA reset.
Data: None
Severity: Error
Log: LOG_INIT
Action: Unload and reload the driver.

0429: MSI-X fast-path request_irq failed (<rc>)
The driver received an error for the request_irq_call.
Data: None
Severity: Warning
Log: LOG_INIT
Action: Unload and reload the driver.

0430: PM resume Failed to enable interrupt
The driver's power management resume function could not enable the interrupt.
Data: None
Severity: Error
Log: LOG_INIT
Action: Perform another PM suspend and resume or HBA reset.

0431: Failed to enable interrupt.
The driver failed to start the interrupt.
Data: None
Severity: Error
Log: LOG_INIT
Action: Unload and reload the driver.

0433: Wakeup on signal: rc=<rc>
A signal other than the LPFC_DATA_READY was received on the worker thread.
Data: None
Severity: Error
Log: LOG_ELS
Action: Unload and reload the driver.

0434: PM resume failed to start worker thread: error=<error>
The driver's power management resume function could not start the worker thread.
Data: None
Severity: Error
Log: LOG_INIT
Action: Unload and reload the driver.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0435: Adapter failed to get Option ROM version status <rc>
The driver could not read the HBA's option ROM.
Data: None
Severity: Error
Log: LOG_INIT
Action: Reset the HBA. Ensure the adapter's firmware is current.

0436: Adapter failed to init, timeout, status reg <status>
The adapter failed during power-up diagnostics after it was reset.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0437: Adapter failed to init, chipset, status reg <status>
The adapter failed during power-up diagnostics after it was reset.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0438: Adapter failed to init, chipset, status reg <status>
The adapter failed during power-up diagnostics after it was reset.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0439: Adapter failed to init, mbxCmd <mbxCommand> READ_REV, mbxStatus <mbxStatus>
Adapter initialization failed when issuing a READ_REV mailbox command.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0440: Adapter failed to init, READ_REV has missing revision information
A firmware revision initialization error was detected.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. Install the latest firmware revision. If this issue persists, report the error to Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0442: Adapter failed to init, mbxCmd <mbxCommand> CONFIG_PORT, mbxStatus <mbxStatus>
Adapter initialization failed when issuing a CONFIG_PORT mailbox command.
Data: (1) hbainit
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0443: Adapter failed to set maximum DMA length mbxStatus <u.mb.mbxStatus>
Cannot set the maximum DMA length to reflect cfg_pci_max_read.
Data: None
Severity: Error
Log: LOG_INIT
Action: Set module parameter lpfc_pci_max_read to 512, 1024, 2048, or 4096.

0446: Adapter failed to init, mbxCmd <mbxCommand> CFG_RING, mbxStatus <mbxStatus>, ring <num>
Adapter initialization failed when issuing a CFG_RING mailbox command.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0447: Adapter failed init, mbxCmd <mbxCommand> CONFIG_LINK mbxStatus <mbxStatus>
Adapter initialization failed when issuing a CONFIG_LINK mailbox command.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0448: Adapter failed to init, mbxCmd <mbxCommand> READ_SPARM, mbxStatus <mbxStatus>
Adapter initialization failed when issuing a READ_SPARM mailbox command.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0449: lpfc_%attr attribute cannot be initialized to %d, allowed range is [%min, %max]
The sysfs attribute value written exceeds attribute range.
Data: (1) attribute name (2) value written (3) minimum value (3) maximum value
Severity: Error
Log: Always
Action: Write a value within the supported range.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0450: lpfc_attr attribute cannot be set to %d, allowed range is [%min, %max]
The sysfs attribute value written exceeds attribute range.
Data: (1) attribute name (2) value written (3) minimum value (3) maximum value
Severity: Error
Log: Always
Action: Write a value within the supported range.

0451: Enable interrupt handler failed
The driver attempted to register the HBA interrupt service routine with the host operating system, but failed.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or driver issue. If this issue persists, report the error to Technical Support.

0453: Adapter failed to init, mbxCmd <mbxCommand> READ_CONFIG, mbxStatus <mbxStatus>
Adapter initialization failed when issuing a READ_CONFIG mailbox command.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0454: Adapter failed to init, mbxCmd <mbxCommand> INIT_LINK, mbxStatus <mbxStatus>
Adapter initialization failed when issuing an INIT_LINK mailbox command.
Data: None
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

0456: Adapter failed to issue ASYNCEVT_ENABLE mbox status x%x
The mailbox command to enable an asynchronous event notification failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: Ensure the adapter firmware is current. Reload the driver.

0457: Adapter Hardware Error
The driver received an interrupt indicating a possible hardware issue.
Data: (1) status (2) status1 (3) status2
Severity: Error
Log: Always
Action: This error could indicate a hardware or firmware issue. If this issue persists, report the error to Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0458: Bring adapter online
The FC driver has received a request to bring the adapter online. This may occur when running lputil.
Data: None
Severity: Warning
Log: LOG_INIT verbose
Action: None required.

0459: Adapter heartbeat failure, taking this port offline.
The Heartbeat mailbox command failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: Ensure the adapter firmware is current. Reload the driver.

0460: Bring adapter offline
The FC driver has received a request to bring the adapter offline. This may occur when running lputil.
Data: None
Severity: Warning
Log: LOG_INIT verbose
Action: None required.

0466: Outstanding IO when bringing Adapter offline
The I/O is still pending while attempting to stop the driver.
Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.

0467: lpfc_topology attribute cannot be set to %d, allowed range is [0, 6], phba->brd_no, val.
The lpfc_topology module parameter is invalid.
Data: None
Severity: Error
Log: LOG_INIT
Action: Use a value in the valid range.

0468: lpfc_restrict_login must be 0 for Physical ports. "vport->cfg_restrict_login = 0;
Cannot restrict the login for the physical port.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0469: lpfc_link_speed attribute cannot be set to %d, allowed range is [0, 8]
The link speed module parameter is invalid.
Data: None
Severity: Error
Log: LOG_INIT
Action: Use a link speed parameter in the valid range.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0472: PCI channel I/O permanent failure
The PCI bus has detected an error.
Data: None
Severity: Error
Log: LOG_INIT
Action: Issue an HBA reset.

0474: Unable to allocate memory for issuing MBOX_CONFIG_MSI command
Mailbox memory pool allocation error.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0475: Not configured for supporting MSI-X cfg_use_msi: 0x%x
The lpfc_use_msi module parameter should have been set to 2.
Data: None
Severity: Error
Log: LOG_INIT
Action: Set module parameter lpfc_use_msi = 2.

0476: HBA not supporting SLI-3 or later SLI Revision: <sli_rev>
The HBA does not support SLI-3 or SLI-4.
Data: None
Severity: Error
Log: LOG_INIT
Action: This HBA does not support msi. Set lpfc_use_msi=0.

0478: MSI request_irq failed (<rc>).
The request_irq kernel API has failed.
Data: None
Severity: Warning
Log: LOG_INIT
Action: Set lpfc_use_msi=0.

0479: Deferred Adapter Hardware Error
An adapter hardware error was sent to the driver.
Data: (1) work_hs, (2) work_status[0], (3) work_status[1]
Severity: Error
Log: LOG_INIT
Action: Perform a dump using hbacmd.

0483:Invalid link-attention link speed: x%x", bf_get(lpfc_acqe_link_speed, acqe_link).
The link speed reported in the link attention interrupt is invalid.
Data: None
Severity: Error
Log: LOG_INIT
Action: Check the switch configuration.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0485: MSI-X slow-path request_irq failed (<rc>).
The request_irq kernel API has failed.
Data: None
Severity: Warning
Log: LOG_INIT
Action: Set module parameter lpfc_use_msi=0.

0486: MSI-X fast-path (<index>) request_irq failed (<rc>).
The request_irq kernel API has failed.
Data: None
Severity: Warning
Log: LOG_INIT
Action: Set module parameter lpfc_use_msi=0.

0490: MSI request_irq failed (<rc>).
The request_irq kernel API has failed.
Data: None
Severity: Warning
Log: LOG_INIT
Action: Set module parameter lpfc_use_msi=0.

0492: Unable to allocate memory for issuing SLI_CONFIG_SPECIAL mailbox command
Mailbox memory pool allocation error.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0493: SLI_CONFIG_SPECIAL mailbox failed with status<rc>
Mailbox command failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: Ensure the adapter's firmware is current. Unload and reload the driver.

0494: Unable to allocate memory for issuing "SLI_FUNCTION_RESET mailbox command"
Mailbox memory pool allocation error.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0495: SLI_FUNCTION_RESET mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>.
Mailbox command failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: Reset the HBA.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0496: Failed allocate slow-path EQ
The event queue for the slow path was not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: Unload and reload the driver.
0497: Failed allocate fast-path EQ
The event queue for the fast path was not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: Unload and reload the driver.
0499: Failed allocate fast-path FCP CQ (<fcp_cqid>).
The completion queue event for the fast path could not be allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: Unload and reload the driver.:
0500: Failed allocate slow-path mailbox CQ
Failed to allocate slow-path mailbox CQ.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0501: Failed allocate slow-path ELS CQ
Failed to allocate slow-path ELS CQ.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0503: Failed allocate fast-path FCP
Failed to allocate fast-path FCP.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0504: Failed allocate slow-path ELS WQ
Failed to allocate slow-path ELS WQ.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0505: Failed allocate slow-path ELS MQ
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0506: Failed allocate receive HRQ\n
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0507: Failed allocate receive DRQ
Failed to allocate receive DRQ.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0520: Slow-path EQ not allocated
The slow-path EQ not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0522: Fast-path EQ <fcp_eqidx> not allocated
The fast-path EQ is not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0523: Failed setup of fast-path EQ <fcp_eqidx>, rc = <rc>
The fast-path EQ setup failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0526: Fast-path FCP CQ <fcp_cqidx> not allocated
The fast-path FCP is not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0527: Failed setup of fast-path FCP CQ <fcp_cqidx>, rc = <rc>
The fast-path FCP CQ setup failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0528: Mailbox CQ not allocated
The mailbox CQ is not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0530: ELS CQ not allocated
The ELS CQ is not allocated
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0534: Fast-path FCP WQ <fcp_wqid> not allocated
The fast-path FCP WQ is not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0535: Failed setup of fast-path FCP WQ <fcp_wqid>, rc = <rc>
The fast-path FCP WQ setup failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0536: Slow-path ELS WQ not allocated
The slow-path ELS WQ is not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0538: Slow-path MQ not allocated
The slow-path MQ is not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

0540: Receive Queue not allocated
The Receive Queue is not allocated.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0542: lpfc_create_static_vport failed to allocate mailbox memory
Failed to allocate mailbox memory for VPort creation.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0543: lpfc_create_static_vport failed to allocate vport_info\n"))
Failed to allocate vport_info.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0544: lpfc_create_static_vport failed to issue dump mailbox command ret <rc> status <mbxStatus>
Failed to issue a dump mailbox command for static VPort creation.
Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.
0545: lpfc_create_static_vport bad information header 0x%x 0x%x\n", le32_to_cpu(vport_info->signature), le32_to_cpu(vport_info->rev) & VPORT_INFO_REV_MASK);
Invalid information header; the signature or revision is invalid.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0546: lpfc_create_static_vport failed to create vport
Failed to create a VPort.
Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.
0560: lpfc_enable_auth attribute cannot be set to <val>, allowed range is [0, 1]
The lpfc_enable_auth attribute can only be 0 or 1.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
0582: Error <rc> during sgl post operation
The SGL post operation failed.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_IP verbose
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0602: Failed to allocate CQ_EVENT entry
Failed to allocate a CQ_EVENT entry.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.
0603: Invalid work queue CQE subtype (x%x)\n", cq-<subtype>
Invalid work queue CQE.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.
0700: Bus Reset on target <i> failed
The bus reset for the specified target failed.
Data: None
Severity: Error
Log: LOG_FCP
Action: None required.
0704: At limitation of <total> preallocated command buffers
The maximum number of command buffers have already been allocated.
Data: None
Severity: Warning
Log: LOG_FCP verbose
Action: None required.
0705: Allocation request of <num> command buffers will exceed max of <hba_queue_depth>. Reducing allocation request to <size>
The number of command buffers requested will exceed the maximum so a smaller quantity will be allocated.
Data: None
Severity: Warning
Log: LOG_FCP verbose
Action: None required.
0708: Allocation request of <num_to_alloc> command buffers did not succeed. Allocated <num_allocated> buffers.
The allocation request for the specified command buffers did not succeed. However, the specified number of buffers has been allocated.
Data: None
Severity: Warning
Log: LOG_FCP
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0711: detected queue full - lun queue depth adjusted to%d
The driver detected a queue full status on a SCSI command response. New LUN queue depth is reported.
Data: (1) New LUN queue depth
Severity: Warning
Log: LOG_FCP verbose
Action: This may indicate an oversubscribed target array. Check your SAN configuration and I/O workload.

0713: SCSI layer issued Device Reset (%d, %d)
A device reset was issued.
Data: None
Severity: Error
Log: LOG_FCP
Action: None required.

0714: SCSI layer issued bus reset
The SCSI layer is requesting the driver to abort all I/Os to all targets on this HBA.
Data: (1) ret
Severity: Error
Log: Always
Action: Check the state of the targets in question.

0720: FCP command <cmnd[0]> residual overrun error
A residual overrun error has occurred while processing the specified FCP command.
Data: (1) request_bufflen (2) resid
Severity: Warning
Log: LOG_FCP verbose
Action: If this issue persists, check the targets for errors.

0721: Device Reset rport failure: rdata <rdata>
The reset of the R_Port failed.
Data: None
Severity: Error
Log: LOG_FCP
Action: None required.

0722: Target Reset rport failure: rdata <rdata>
The reset of the target failed.
Data: None
Severity: Error
Log: LOG_FCP
Action: None required.

0723: SCSI layer issued Target Reset (%d, %d)
The SCSI layer issued a target reset.
Data: None
Severity: Error
Log: LOG_FCP
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0724: I/O flush failure for context <"LUN", "TGT", "HOST", "Unknown">: cnt <cnt>
The I/O flush to the LUN, target, or host has failed.
Data: None
Severity: Error
Log: LOG_FCP
Action: None required.

0727: TMF <cmd> to TGT <TGT#> LUN <LUN#> failed (<ulpStatus>, <ulpWord[4]>)
The task management command failed.
Data: None
Severity: Error
Log: LOG_FCP
Action: None required

0729: FCP cmd <cmd> failed <target>/<lun> status: <status> result: <result>
The specified device failed an FCP command.
Data: (1) ulpContext (2) iotag
Severity: Warning
Log: LOG_FCP verbose
Action: Check the state of the target in question.

0730: FCP command failed: RSP
The FCP command failed with a response error.
Data: (1) resp_info (2) scsi_status (3) ResId (4) SnsLen (5) RspLen (6)rsplInfo3
Severity: Warning
Log: LOG_FCP verbose
Action: Check the state of the target in question.

0734: FCP read check error
The issued FCP command returned a read check error.
Data: (1) fcpDI (2) rspResId (3) fcpi_parm (4) cmd[0]
Severity: Warning
Log: LOG_FCP verbose
Action: Check the state of the target in question.

0735: FCP Read Check Error and Underrun Data
HBA reported under run from storage array.
Data: (1) vpi (2) fcpDI (3) res_id (4) fcpi_parm
Severity: Warning
Log: LOG_FCP_ERROR verbose
Action: No action needed, informational.

0748: Abort handler timed out waiting for abort to complete:ret <status> D <target id> LUN <lun id>
The abort handler timed out waiting for abort to complete.
Data: None
Severity: Error
Log: Always
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

0749: SCSI layer issued abort device
The SCSI layer aborted a device.
Data: (1) ret, (2) id, (3) lun, (4) snum
Severity: Warning
Log: LOG_FCP verbose
Action: None required.

0915 Register VPI failed:<mbxStatus>
Could not register the VPI.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

1005: AUTHENTICATION_FAILURE Nport:<port>
The system detected DHCHAP authentication failure on a port.
Data: (1) nlp_DID
Severity: Error
Log: LOG_SECURITY
Action: Verify authentication settings and keys on local and remote port.

1006: Bad Name tag in auth message < message >
DHCHAP Authentication process failed when invalid tag was detected.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1007: Bad Name length in auth message < message >
DHCHAP Authentication process failed when invalid name was detected.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1008: Bad Number of Protocols <message>
DHCHAP Authentication process failed due to unexpected protocol number.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1009: Bad param type <message>
DHCHAP Authentication process failed when invalid protocol was detected.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1010: Bad Tag 1 <message>
DHCHAP Authentication process failed when bad Tag was detected.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1011: Auth_neg no hash function chosen
DHCHAP Authentication process failed when an incorrect hash function was specified.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1012: Auth_negotiate Bad Tag <message>
DHCHAP Authentication process failed due to bad Tag for auto negotiation.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1013: Auth_negotiate no DH_group found
DHCHAP Authentication process failed when incorrect or missing DH Group was detected.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1014: dhchap challenge bad name tag <message>
DHCHAP Authentication process failed when incorrect Challenge name tag was detected.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1015: dhchap challenge bad name length <message>
DHCHAP Authentication process failed due to unexpected Challenge name length.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1016: dhchap challenge Hash ID not Supported <message>
DHCHAP Authentication process failed due to uncorroborated Challenge Hash ID.
Data: (1) message
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1017: dhchap challenge could not find DH Group
DHCHAP Authentication process failed due to uncorroborated Challenge Group.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1018: dhchap challenge No Public key for non-NULL DH Group
There is no Public key for the non-NULL DH Group.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: None required.

1019: Request tranid <tran_id> timed out
A transaction with storage array could not complete due to timeout.
Data: (1) tran_id
Severity: Warning
Log: LOG_SECURITY verbose
Action: Software driver warning. If this issue persists, report these errors to Technical Support.

1021: ERROR: attempted to queue security work, when no workqueue created
Driver encountered missing queue required for processing security information.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report these errors to Technical Support.

1022: Security request does not exist
A security request operation failed because there was no match found for such request.
Data: None
Severity: Warning
Log: LOG_SECURITY
Action: Software driver warning. If this issue persists, report these errors to Technical Support.

1023: Warning - data may have been truncated. Data: <data> reqdl: <data_len> mesdl:<data_len>
A security message exchange operation failed because the response was missing or unreliable.
Data: None
Severity: Warning
Log: LOG_SECURITY
Action: Software driver warning. If this issue persists, report these errors to Technical Support.

1028: Start Authentication: No buffers
The authentication failed because some memory resources were not allocated.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1029: Reauthentication Failure
The driver encountered errors and there was a failure to re-authenticate.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.
1031: Start Authentication: Get config failed
The authentication failed due to some error during port configuration.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.
1032: Start Authentication: get config timed out
The node authentication was aborted because waiting for port configuration to complete, timed out.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.
1033: Received auth_negotiate from Nport: < nlp_DID>
Unsolicited authentication negotiation message received from a port.
Data: (1) nlp_DID
Severity: Warning
Log: LOG_SECURITY
Action: No action, this message is informational.
1034: Not Expecting Challenge - Rejecting Challenge
Unsolicited authentication challenge received from a port, was rejected.
Data: None
Severity: Warning
Log: LOG_SECURITY
Action: Software driver warning. If this issue persists, report errors to the Technical Support.
elx_mag1036: Authentication transaction reject - re-auth request reason <reason> exp <explanation>
An authentication was rejected and requested again due to reason as displayed with explanation.
Data: (1) reason (2) explanation.
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.
1037: Authentication transaction reject - restarting authentication, reason <reason> exp <explanation>
An authentication process was rejected then restarted and authentication requested again due to reason as displayed with explanation.
Data: (1) reason (2) explanation.
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1039: Not Expecting Reply - rejecting. State <state>
An unanticipated reply was received during authentication and was subsequently rejected.
Data: (1) auth_state.
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.
1040: Bad Reply trans_id- rejecting. Trans_id < trans_id > Expecting: < trans_id>
Unexpected transaction id was received during authentication and was subsequently rejected.
Data: (1) auth_state
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.
1043: Authentication LS_RJT
The authentication request was rejected.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
1045: Issue AUTH_NEG failed Status:%x
The authentication negotiation failed.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
1048: Issue AUTH_REJECT failed
Could not issue the reject for the authentication request.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.
1049: Authentication is enabled but authentication service is not running
Discovery failed because DHCHAP Authentication was enabled while no authentication service was established.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Start the authentication daemon (fcauthd).
elx_msg1050: Authentication mode is disabled, but is required by the fabric
Discovery failed because the switch fabric required authentication, but authentication was not configured or the authentication mode for this port pair is disabled.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Configure the driver to authenticate with the switch or disable authentication on the switch to this port.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1053: Start Authentication: Security service offline
The authentication failed because security service was unavailable.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1055: Authentication parameter is disabled, but is required by the fabric
FLOGI failed because the fabric has indicated that authentication is required, but authentication has not yet been configured or enabled on the HBA.
Data: None
Severity: Error
Log: LOG_SECURITY
Action: Configure authentication on this HBA.

1057: Authentication transaction reject. reason <reason> exp <explanation>
An authentication was rejected and requested again due to the reason as displayed with the explanation.
Data: (1) reason (2) explanation.
Severity: Error
Log: LOG_SECURITY
Action: Software driver Error. If this issue persists, report errors to the Technical Support.

1058: Waiting for authentication service
There was a delay when the authentication service was not initially available as expected.
Data: None
Severity: Warning
Log: LOG_SECURITY
Action: Software driver warning. If this issue persists, report these errors to Technical Support.

1059: Authentication became available
The authentication service came online but was not initially available as expected.
Data: None
Severity: Warning
Log: LOG_SECURITY
Action: Software driver warning. If this issue persists, report these errors to Technical Support.

1201: Failed to allocate dfc_host
Failed to allocate memory for the dfc_host_struct.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.

1210: Invalid cmd size: cmd <cmd> cmdsz <cmdsized> rspsz <rspsized>
The management command for LPFC 2100 has failed.
Data: None
Severity: Error
Log: LOG_LIBDFC
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1211: genreq alloc failed\n");
Resource allocation failure.
Data: (1) return code
Log: LOG_LIBDFC
Severity: Error
Action: Kernel memory resources too low.
1213: FCoE cmd overflow: off <#> + cnt <#> > cmdsz <#>
Application has tried to read more data than originally requested.
Data: (1) response offset (2) size (3) cmd size
Log: LOG_LIBDFC
Severity: Error
Action: Application may have sent a invalid command.
1214: Cannot issue FCoE cmd SLI not active: <#> rc= -EACCESS
The SLI layer has not been initialized.
Data: (1) offset
Log: LOG_LIBDFC
Severity: Error
Action: Restart the HBA.
1215: Cannot issue FCoE cmd: not ready or not in maint mode"
Either the external link is unplugged, link down, and the FCoE is not in maintenance mode.
Data: (1) current offset (2) return code.
Log: LOG_LIBDFC
Severity: Error
Action: Plug external cable in or set FCoE in maintenance mode.
1216: FCoE IOCB failed: off <#> rc <#>
FCoE command generated by the application has failed.
Data: (1) offset (2) return code.
Log: LOG_LIBDFC
Severity: Error
Action: Application should retry the command.
1223: menlo_write: couldn't alloc genreq
Resource allocation failure.
Data: None
Log: LOG_LIBDFC
Severity: Error
Action: Kernel memory resources too low.
1224: FCoE iocb failed off <#> rc=<#>",
FCoE command failed in SLI.
Data: (1) offset (2) return code
Log: LOG_LIBDFC
Severity: Informational.
Action: Retry the command, if it fails again, reset HBA when convenient.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1227: FCoE IOCB TMO: handler set for <context3>
The management command for the LPFC 2100 has timed out.
Data: None
Severity: Warning
Log: LOG_LIBDFC
Action: None required.
1228: FCoE IOCB TMO: handler set for <context3>
A management IOCB for the LPFC 2100 has timed out
Data: None.
Severity: Warning
Log: LOG_LIBDFC
Action: None required.
1229: Waiting for menlo mnt
Waiting for the LPFC 2100 to enter maintenance method.
Data: None.
Severity: Warning
Log: LOG_LIBDFC
Action: None required.
1230: Could not find buffer for FCoE cmd:off <#> indmp <addr> off <#>
Could not find resources associated with this FCoE command.
Data: (1) current offset (2) buffer desc pointer (3) size.
Severity: Error
Log: LOG_LIBDFC
Action: Try reloading the driver when convenient.
1231: bad bpl:
A invalids buffer list was detected upon completion.
Data: None.
Severity: Error
Log: LOG_LIBDFC
Action: None required.
1235: Could not find buffer for FCoE cmd: off:<#> poff:<#> cnt:<#> mlastcnt:<#> addl:<x> addh:<x> mdsz:<#>
FCoE command failed because it could not find the resource.
Data: (1) current offset (2)previous offset (3) count (4) last count (5) address low (6) address high
Severity: Error
Log: LOG_LIBDFC
Action: No action needed, informational.
1238: FCoE IOCB failed: off <#> rc=<#>
The command generated by the driver to check the FCoE has failed.
Data: (1) offset (2) return code
Log: LOG_LIBDFC
Severity: Error
Action: Make sure link is up or the adapter has set menlo in maintenance mode.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1240: Unable to allocate command buffer memory
Could not allocate memory for the command buffer.
Data: None.
Severity: Error
Log: LOG_LINK_EVENT
Action: None required.
1243: Menlo command error. code=%d.\n", mlrsp->code
The Menlo maintenance command failed.
Data: None.
Severity: Error
Log: LOG_LINK_EVENT
Action: None required.
1244: Unable to allocate response buffer memory.
Could not allocate memory for the management command response.
Data: None.
Severity: Error
Log: LOG_LINK_EVENT
Action: None required.
1246: FCoE chip is running golden firmware. Update FCoE chip firmware immediately <fw_type>
The FCoE is running the golden firmware.
Data: (1) firmware-type
Severity: Error
Log: LOG_LINK_EVENT
Action: Try resetting the FCoE to operational mode and disable maintenance mode.
1247: FCoE chip is running diagnostic firmware. Operational use suspended. <fw_type>
The FCoE is running a diagnostic.
DATA:(1) firmware-type
Severity: Error
Log: LOG_LINK_EVENT
Action: Try resetting the FCoE to operational mode.
1248: FCoE chip is running unknown firmware. <fw_type>
The FCoE is running an unknown firmware version.
Data: (1) firmware-type
Severity: Error
Log: LOG_LINK_EVENT
Action: Try resetting the FCoE to operational mode. Try loading latest FCoE firmware.
1249: Invalid FRU data found on adapter. Return adapter to Emulex for repair.
The FRU data on the FCoE chip is invalid.
Data: (1) firmware-type
Severity: Error
Log: LOG_LINK_EVENT
Action: Try resetting the FCoE to operational mode. Try loading latest FCoE firmware or send the HBA back to Emulex for repair.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1250: Menlo command error. code=<#>
The IOCB driver sent to check FCoE state has bad header size.
Data: (1) return code
Log: LOG_LINK_EVENT
Severity: Error
Action: Try resetting the FCoE to operational mode.
1251: Menlo command error. code=<#>
The IOCB driver sent to check FCoE state has failed, no resources.
Data: (1) return code
Log: LOG_LINK_EVENT
Severity: Error
Action: Try resetting the FCoE to operational mode.
1252: Menlo command error. code=<#>
The IOCB driver sent to check FCoE state has failed.
Data: (1) return code
Log: LOG_LINK_EVENT
Severity: Error
Action: Try resetting the FCoE to operational mode.
1257: lpfc_menlo_issue_iocb: handler set for <context3>.
Data: None
Log: LOG_LIBDFC
Severity: Warning
Action: None required.
1259: mbox: Issued mailbox cmd <u.mb.mbxCommand> while in stopped state.
Only the dump mailbox command and reset adapter mailbox command are allowed when in the stopped state.
Data: None
Severity: Warning
Log: LOG_MBOX
Action: None required.
1262: Failed to allocate dfc_host
Could not allocate memory the dfc_host_struct.
Data: None
Log: LOG_LIBDFC
Severity: Error
Action: None required.
1268: Find ndlp returned NULL for oxid:x%x SID:x%x, oxid, sid.(int)off, rc.
Could not find the node for this DID.
Data: None
Severity: Warning
Log: LOG_ELS
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1302: Invalid speed for this board: Reset link speed to auto: <cfg_link_speed>
The driver is re-initializing the link speed to auto-detect.
Data: None
Severity: Warning
Log: LOG_LINK_EVENT verbose
Action: None required.
1303: Link Up Event <eventTag> received
A link up event was received. It is possible for multiple link events to be received together.
Data: (1) fc_eventTag (2) granted_AL_PA (3) UlnkSpeed (4) alpa_map[0]
Detail: If link events received, log (1) last event number received, (2) ALPA granted, (3) Link speed (4) number of entries in the loop init LILP ALPA map. An ALPA map message is also recorded if LINK_EVENT verbose mode is set. Each ALPA map message contains 16 AL_PAs.
Severity: Error
Log: Always
Action: If numerous link events are occurring, check the physical connections to the FC network.
1304: Link Up Event ALPA map
A link up event was received.
Data: (1) wd1, (2) wd2, (3) wd3, (4) wd4
Severity: Warning
Log: LOG_LINK_EVENT verbose
Action: If numerous link events are occurring, check the physical connections to the FC network.
1305: Link Down Event <eventTag> received
A link down event was received.
Data: (1) fc_eventTag (2) hba_state (3) fc_flag
Severity: Error
Log: Always
Action: If numerous link events are occurring, check the physical connections to the FC network.
1306: Link Up Event in loop back mode x%x received Data: x%x x%x x%x x%x
Link up notification; configured for loopback.
Data: (1) fc_eventTag (2) granted_AL_PA (3) UlnkSpeed (4) alpa_map[0]
Severity: Error
Log: LOG_LINK_EVENT
Action: None required.
1308: Menlo Maint Mode Link up Event x%x rcvd Data: x%x x%x x%x
Link down notification; configured for loopback.
Data: (1) fc_eventTag (2) port_state (3) vport fc_flag
Severity: Error
Log: LOG_LINK_EVENT
Action: None required.
1309: Link Up Event npiv not supported in loop topology
NPIV is not supported in loop topology.
Data: None
Severity: Error
Log: LOG_LINK_EVENT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1310: Menlo Maint Mode Link up Event <eventTag> rcvd
The link is up in maintenance mode; only management commands are allowed.
Data: (1) fc_eventTag (2) port_state (3) vport fc_flag
Severity: Error
Log: LOG_LINK_EVENT
Action: None required.
1312: Link Down Event <eventTag> received
Maintenance mode link up notification received without entering link down.
Data: (1) fc_eventTag (2) port_state (3) vport fc_flag
Severity: Error
Log: LOG_LINK_EVENT
Action: None required.
1400: Failed to initialize sgl list.
Failed to initialize SGL list during initialization.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
1401: Failed to enable pci device.
Failed to enable PCI device during initialization.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
1402: Failed to set up pci memory space.
PCI initialization failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
1403: Failed to set up driver resource.
Driver resource initialization failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
1404: Failed to set up driver resource.
Driver resource initialization failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1405: Failed to initialize iocb list.
Driver resource initialization failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1406: Failed to set up driver resource.
Initialization failed to set up driver resource.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1407: Failed to create scsi host.
Initialization failed to create SCSI host.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1408: Failure HBA POST Status: sta_reg=0x%x, "perr=x%x, sfi=x%x, nip=x%x, ipc=x%x, xrom=x%x, "dl=x%x, pstatus=x%x\n", sta_reg.word0, bf_get(lpfc_hst_state_perr, &sta_reg),
The HBA's power on self test has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1409: Failed to enable pci device.
Failed to enable PCI device during initialization.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1410: Failed to set up pci memory space.
Initialization failed to set up PCI memory space.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1411: Failed to set up driver resource.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1412: Failed to set up driver resource.
Initialization failed to set up driver resource.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1413: Failed to initialize iocb list.
Initialization failed to initialize the IOCB list.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1414: Failed to set up driver resource.
Initialization failed to set up driver resource.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1415: Failed to create scsi host.
Initialization failed to create SCSI host.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1416: Failed to allocate sysfs attr
Initialization failed to sysfs attribute.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1418: Invalid HBA PCI-device group: <dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1419: Invalid HBA PCI-device group: <dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1420: Invalid HBA PCI-device group:<dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1421: Failed to set up hba
Initialization failed to set up the HBA.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1422: HBA Unrecoverable error: uerr_lo_reg=<ue lo>, uerr_hi_reg=<ue hi>, online0_reg=<Online0>, online1_reg=<Online1>
The HBA has notified the driver that it has encountered an unrecoverable error.
Data: None
Severity: Error
Log: LOG_INIT
Action: A dump from the OneCommand Manager application should be taken. Then, the driver should be unloaded and reloaded.

1423: HBA Unrecoverable error: uerr_lo_reg=<ue lo>, uerr_hi_reg=<ue hi>, online0_reg=<Online0>, online1_reg=<Online1>
The HBA has notified the driver that it has encountered an unrecoverable error.
Data: None
Severity: Error
Log: LOG_INIT
Action: A dump from the OneCommand Manager application should be taken. Then, the driver should be unloaded and reloaded.

1424: Invalid PCI device group:<pci_dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1425: Invalid PCI device group: <pci_dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1426: Invalid PCI device group: <pci_dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1427: Invalid PCI device group: <pci_dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1428: Invalid PCI device group: <pci_dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1429: Invalid PCI device group: <pci_dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1430: Failed to initialize sql list
Failed to initialize SQL list.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1431: Invalid HBA PCI-device group: <pci_dev_grp>
Invalid HBA PCI-device group detected.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1432: Failed to initialize rpi headers.
Failed to initialize RPI headers.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

1476: Failed to allocate sysfs attr
Failed to allocate sysfs attributes.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1477: Failed to set up hba
Failed to set up the HBA.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
1603: Loopback test did not receive expected data length. actual length <len>expected length <full_size>.
The loopback test did not receive the same amount of data that it transmitted.
Data: None
Severity: Error
Log: LOG_LIBDFC
Action: None required.
1800: Could not issue unreg_vpi
Driver attempt to unregister VPI failed.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: Software driver error. If this issue persists, report these errors to Technical Support.
1801: Create vport work array FAILED: cannot do scsi_host_get
The driver was unable to get a reference to a SCSI host.
Data: None
Severity: Warning
Log: LOG_VPORT verbose
Action: Software driver warning. If this issue persists, report these errors to Technical Support.
1816: FLOGI NPIV supported, response data <port>
The fabric reports support for NPIV upon FLOGI.
Data: (1) response_multiple_NPort
Severity: Warning
Log: LOG_VPORT verbose
Action: No action needed, informational.
1817: Fabric does not support NPIV - configuring single port mode
The fabric reports no support for NPIV upon FLOGI.
Data: None
Severity: Warning
Log: LOG_VPORT verbose
Action: No action needed, informational.
1818: VPort failed init, mbxCmd <mailbox command> READ_SPARM mbxStatus <mailbox status> , rc = <status>
A pending mailbox command that was issued to initialize the port failed.
Data: (1) mbxCmd (2) mbxStatus (3) rc
Severity: Error
Log: LOG_VPORT verbose
Action: Software driver error. If this issue persists, report these errors to Technical Support.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1819: Unrecognized lpfc_sli_mode parameter: <mode>
The user has attempted to set the SLI mode to an invalid value. The valid values for the SLI mode are 0, 2, and 3.
Data: (1) lpfc_sli_mode
Severity: Error
Log: LOG_VPORT verbose
Action: The lpfc_sli_mode driver parameter setting must be corrected. Valid values are 0, 2, and 3.

1820: Unable to select SLI-3. Not supported by adapter.
The HBA is incapable of operating in a given mode.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: SLI-3 mode is only available on some HBAs. Do not attempt to force the SLI mode to 3 on HBAs that do not support SLI-3 mode. This is an informational message. HBAs that do not support SLI-3 will be configured to run in SLI-2 mode, but it is recommended to use the auto setting (0).

1821: Create VPORT failed. Invalid WWN format
The port could not be created due to an invalid WWNN or WWPN format.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: Provide a valid WWN when creating Vports.

1822: Invalid <name>: <xx: xx: xx: xx: xx: xx: xx: xx>
An invalid WWN was used when creating a VPort.
Data: (1) type_name (2) wwn[1] (3) wwn[3] (3) wwn[5] (4) wwn[7]
Severity: Error
Log: LOG_VPORT verbose
Action: When creating a VPort you must furnish a valid WWN.

1823: Create VPORT failed. Duplicate WWN on HBA.
The port could not be created because it would duplicate an existing WWNN HBA address. The resources for the port had to be discarded.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: Provide a WWN that is unique.

1824: NPIV enabled: Override lpfc_sli_mode parameter (<mode>) to auto(0)
The lpfc_enable_npiv and lpfc_sli_mode driver parameter settings conflict. The HBA must be configured for SLI-3 mode to support NPIV.
Data: (1) lpfc_sli_mode
Severity: Error
Log: LOG_VPORT verbose
Action: This is an informational message that indicates that the lpfc_enable_npiv and lpfc_sli_mode parameter settings are not compatible. Resolve the parameter conflict by setting the SLI mode to 0 or 3 or, if SLI-2 mode is required then disable NPIV.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1825: Vport Created.
This message is displayed to indicate that a port was created in the system. It is displayed at this level to ensure it is always appears at all log levels.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: No action, informational.

1826: Vport Disabled.
The port had to be disabled in the system.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: No action, informational.

1827: Vport Enabled
The port had to be enabled after possible recovery from some errors.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: No action, informational.

1828: Vport Deleted
A Vport was deleted.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: No action, informational.

1829: CT command failed to delete objects on fabric.
A command issued to the fabric to delete an associated resource for an object, such as for a port, failed.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: Software driver error. If this issue persists, report these errors to Technical Support.

1830: Signal aborted mbxCmd <command>
A pending mailbox command was aborted because the thread received a signal.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: You should retry the attempted command.

1831: Create VPORT Interrupted
The port creation process was unexpectedly interrupted at a critical time and the operation was unsuccessful.
Data: None
Severity: Error
Log: LOG_VPORT verbose
Action: The process was interrupted while creating a VPort. Retry the command.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

1832: No pending MBOX command to handle
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.
1835: Vport discovery quiesce failed: state <port_state> fc_flags <fc_flag> wait msecs <jiffies_to_msecs(jiffies - start_time)>
Could not pause discovery on this VPort.
Data: None
Severity: Error
Log: LOG_VPORT
Action: None required.
1836: Could not issue unreg_login(all_rpis) status <rc>
The unreg_login cannot be issued.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_VPORT
Action: None required.
1837: Vport_delete failed: Cannot delete static vport
Static VPorts cannot be deleted.
Data: None
Severity: Error
Log: LOG_VPORT
Action: None required.
1838: Failed to INIT_VPI on vpi <vpi> status <rc>
Failed to INIT_VPI.
Data: None
Severity: Error
Log: LOG_VPORT
Action: None required.
2000: Failed to allocate mbox for read_FCF cmd
Failed to allocate mailbox for READ_FCF command.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.
2001: Unable to allocate memory for issuing SLI_CONFIG_SPECIAL mailbox command
Unable to allocate memory for issuing the SLI_CONFIG_SPECIAL mailbox command.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2002: Error Could not grow rpi count
An error occurred because the RPI count could not be increased.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2004: Failed to allocate XRI.last XRITAG is <XRI> Max XRI is <MAX_XRI>, Used XRI is <USED_XRI>.
All XRIs are in use.
Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

2005: Unable to deregister pages from HBA: <rc>
The SGL pages could not be unregistered from the firmware.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2007: Only Limited Edition cmd Format supported <iocb.ulpCommand>
The SGL pages could not be unregistered from the firmware.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2008: Error <rc> posting all rpi headers
The RPI headers could not be posted to the firmware.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2009: Failed to allocate mbox for ADD_FCF cmd
Failed to allocate mailbox for ADD_FCF command.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2010: Resume RPI Mailbox failed status <status>, mbxStatus <mbx status>
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2011: Unable to allocate memory for issuing SLI_CONFIG_SPECIAL mailbox command
Unable to allocate memory for issuing SLI_CONFIG_SPECIAL mailbox command.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2012: Mailbox failed , mbxCmd <mbx_cmd> READ_CONFIG, mbxStatus <mbx status>
The READ_CONFIG mailbox command failed.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2013: Could not manually add FCF record 0, status <rc>
Could not add FCF record to the FCF list.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2014: Invalid command <iocb.ulpCommand>
The IOCB command is invalid.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2015: Invalid CT %x command <iocb.ulpCommand>
Invalid Command-Type in the IOCB is not supported.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2017: REG_FCFI mbxStatus error <mbx status> HBA state <port_state>
The REG_FCFI mailbox command has failed.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

2018: REG_VFI mbxStatus error <mbx status> HBA state <port_state>
The REG_VFI mailbox command has failed.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2020: Failed to allocate mbox for ADD_FCF cmd
Failed to allocate mailbox for ADD_FCF command.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2022: VPI Mailbox failed status <status>, mbxStatus <mbxStatus>
The INIT VPI mailbox command has failed.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2401: Failed to allocate memory for ELS XRI management array of size <els_xri_cnt>.
Initialization failed to allocate memory for the ELS XRI management array.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2500: EQ_CREATE mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>
The mailbox command sent to create the event queue has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2501: CQ_CREATE mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>
The mailbox command sent to create the completion queue has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2502: MQ_CREATE mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>
The mailbox command sent to create the mailbox queue has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2503: WWQ_CREATE mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>
The mailbox command sent to create the work queue has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2504: RQ_CREATE mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>

The mailbox command sent to create the receive queue has failed.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2505: EQ_DESTROY mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>

The mailbox command sent to delete the event queue has failed.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2506: CQ_DESTROY mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>

The mailbox command sent to delete the completion queue has failed.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2507: MQ_DESTROY mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>

The mailbox command sent to delete the mailbox queue has failed.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2508: WQ_DESTROY mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>

The mailbox command sent to delete the work queue has failed.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2509: RQ_DESTROY mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>

The mailbox command sent to delete the work queue has failed.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2510: RQ_DESTROY mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>
The mailbox command sent to delete the work queue has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2511: POST_SGL mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>n
The mailbox command sent to post the SGL pages to the firmware has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2512: REMOVE_ALL_SGL_PAGES mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>
The mailbox command sent to delete the SGL pages from the firmware has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2513: POST_SGL_BLOCK mailbox command failed status <shdr_status> add_status <shdr_add_status> mbx status <rc>
The mailbox command sent to post the SGL pages to the firmware has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2514: POST_RPI_HDR mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx status <rc>
The mailbox command sent to post the RPUI header pages to the firmware has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2515: ADD_FCF_RECORD mailbox failed with status <rc>
The mailbox command to add the FCF record has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2516: DEL FCF of default FCF Index failed mbx status <rc>, status <shdr_status> add_status<shdr_add_status>

The mailbox command to delete the FCF record has failed.

Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2517: Unregister FCFI command failed status %d, mbxStatus %x", rc, bf_get(lpfc_mqe_status, &mbx->u.mqe)

The driver was unable to unregister the FCFI from the firmware.

Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2518: Requested to send 0 NOP mailbox cmd

Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.

2519: Unable to allocate memory for NOP mailbox command

Memory allocation for this mailbox command has failed.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2520: NOP mailbox command failed status %x add_status %x mbx status %x, shdr_status, shdr_add_status, rc.

The NOP mailbox command has failed.

Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.

2521: READ_FCF_RECORD mailbox failed with status <shdr_status> add_status <shdr_add_status>, mbx

The READ_FCF_RECORD mailbox command failed.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2523: Allocated DMA memory size (<alloc_len>) is less than the requested DMA memory size (<req_len>)

The ADD_FCF_RECORD mailbox command failed to retrieve the length required from the firmware.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2524: Failed to get the non-embedded SGE virtual address
The READ_FCF_RECORD mailbox command could not retrieve the SGE that was requested.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

2527: Failed to allocate non-embedded SGE array.
Failed to allocate the non-embedded SGE array.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

2528: Mailbox command <vpi> cannot issue
The mailbox command could not be issued because the mailbox interrupt is disabled.
Data: (1) mbxCommand (2) sli_flag (3) flag
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2529: Mailbox command <vpi> cannot issue
Data: (1) mbxCommand (2) sli_flag (3) flag
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2530: Mailbox command <vpi> cannot issue
The SLI layer in the driver is inactive.
Data: (1) mb.mbxCommand (2) sli_flag (3) flag
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2531: Mailbox command <cpi> cannot issue
Data: (1) mb.mbxCommand (2) sli_flag (3) flag
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2532: Mailbox command <vpi> (<mbxCommand>) cannot issue
The mailbox bootstrap code detected that the SLI layer is active.
Data: (1) sli4_mbox_opcode (2) sli_flag, (3) MBX_POLL
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2533: Mailbox command <vpi> (<mbxCommand>) cannot issue
Data: (1) sli4_mbox_opcode (2) sli_flag (3) MBX_NOWAIT
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2535: Unsupported RQ count. (<entry_count>)
The receive queue ring can only be 512, 1024, 2048, or 4096.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2536: Unsupported RQ count. (<entry_count>)
The receive queue ring can only be 512, 1024, 2048, or 4096.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2537: Receive Frame Truncated!
The receive unsolicited handler detected a truncated frame.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2539: Dropped frame rctl:%s type:%s\n
An unsupported frame was received by the port and dropped.
Data: (1) rctl_names[fc_hdr->fh_r_ctl], (2) type_names[fc_hdr->fh_type]
Severity: Error
Log: Always
Action: No action needed, informational.

2540: Ring <ring #> handler: unexpected Rctl <fh_rctl> Type <fh_type>
The received frame has an unsupported RCTL or FH_TYPE.
Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

2541: Mailbox command <vpi> (<mbxCommand>) cannot issue
Data: (1) sli_mbox_opcode (2) sli_flag (3) flag
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2542: Try to issue mailbox command <vpi> (<mbxCommand>) synchronously ahead of async mailbox command queue
Attempting to send a synchronous mailbox command ahead of the asynchronous mailbox commands.
Data: (1) sli4_mbx_opcode or sli_mbox_opcode, (2) sli_flag, (3) flag
Severity: Warning
Log: LOG_MBOX, LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2543: Mailbox command <vpi> (<mbxCommand>) cannot issue
The mailbox command does not have all of the fields set correctly.
Data: (1) sli_mbox_opcode (2) sli_flag (3) flag
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.
2544: Mailbox command <vpi> (<mbxCommand>) cannot issue
The HBA cannot be accessed on the PCI bus.
Data: (1) sli_mbox_opcode (2) sli_flag (3) flag
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.
2546: New FCF found index <index> tag <event_tag>
A new FCF has been found.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: None required.
2547: Read FCF record failed
Could not read the FCF record from the firmware.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: None required.
2548: FCF Table full count <count> tag <event_tag>
The FCF table is full.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.
2549: FCF disconnected from network index <index> tag <event_tag>
The FCF has disconnected from the network.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: None required.
2550: UNREG_FCFI mbxStatus error <u.mb.mbxStatus> HBA state <port_state>.
The UNREG_FCFI mailbox command has failed.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2551: UNREG_FCFI mbox allocation failed HBA state <port_state>
The allocation for the UNREG_FCFI mailbox command has failed.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

2552: UNREG_FCFI issue mbox failed rc <rc> HBA state <port_state>.
The UNREG_FCFI mailbox command has failed.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

2553: lpfc_unregister_unused_fcf failed to read FCF record HBA state.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

2554: Could not allocate memory for fcf record
Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2555: UNREG_VFI mbxStatus error <u.mb.mbxStatus> HBA state <port_state>
The UNREG_VFI mailbox command has failed.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

2556: UNREG_VFI mbox allocation failed HBA state <port_state>
Could not allocate memory for UNREG_VFI mailbox command.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

2557: UNREG_VFI issue mbox failed rc <rc> HBA state <port_state>
Could not issue the UNREG_VFI mailbox command.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2558: ADD_FCF_RECORD mailbox failed with status<shdr_status> add_status <shdr_add_status>
The ADD_FCF_RECORD mailbox command has failed.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2559: Block sgl registration required DMA size <reqlen> great than a page.
Attempting to register more SGEs with the firmware than can fit in a page.
Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.

2560: Failed to allocate mbox cmd memory\n
Failed to allocate mailbox command memory.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2561: Allocated DMA memory size (<alloclen>) is less than the requested DMA memory size (<reqlen>)
Could not get the memory required for the number of XRIs that are attempting to be posted.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2562: No room left for SCSI XRI allocation max_xri=<sli4_hba.max_cfg_param.max_xri>, els_xri=<els_xri_cnt>n
The number of allocated XRIs has reached the max_xri value.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2563: Failed to allocate memory for SCSI XRI management array of size <sli4_hba.scsi_xri_max>.
Initialization could not allocate memory to hold the XRIs.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2564: POST_SGL_BLOCK mailbox command failed status <shdr_status> add_status <shdr_add_status> mbx status <rc>
The list of XRI SGEs failed to be registered with the firmware.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2566: Failed to allocate table entry
Failed to allocate connection table entry.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2567: Config region 23 has bad signature
The driver was unable to read Config Region 23 because it has an invalid signature.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2568: Config region 23 has bad version
The driver was unable to read Config Region 23 because it is an invalid version.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2569: lpfc_dump_fcoe_param: memory allocation failed
Memory allocation has failed.
Data: None
Severity: Warning
Log: LOG_MBOX
Action: None required

2570: Failed to read FCoE parameters
The driver failed to read FCoE parameters.
Data: None
Severity: Error
Log: LOG_MBOX, LOG_INIT
Action: None required.

2572: Failed allocate memory for fast-path per-EQ handle array
Failed to allocate memory for the fast-path per-EQ handle array.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2573: Failed allocate memory for msi-x interrupt vector entries
The driver was unable to allocate memory during initialization of the MSI-X interrupt array.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2574: Not enough EQs (<sli4_hba.max_cfg_param.max_eq>) from the pci function for supporting FCP EQs (<cfg_fcp_eq_count>)

Failed to create the minimum fast-path event queues.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2575: Not enough EQs (<max_eq>) from the pci function for supporting the requested FCP EQs (<cfg_fcp_eq_count>), the actual FCP EQs can be supported: <eq_count>

The driver was not configured with enough fast-path event queues.

Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.

2576: Failed allocate memory for fast-path EQ record array

Failed to allocate memory for the fast-path EQ record array.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2577: Failed allocate memory for fast-path CQ record array

Failed to allocate memory for the fast-path EQ record array.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2578: Failed allocate memory for fast-path WQ record array

Failed to allocate memory for the fast-path EQ record array.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2579: Slow-path wqe consume event carries miss-matched qid: wcqe-qid=<wcqe_qid>, sp-qid=<sp_qid>

The consumed entry does not have the slow path's queueID.

Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

2580: Fast-path wqe consume event carries miss-matched qid: wcqe-qid=<fcp_wqid>.

The consumed entry does not have the fast path's queueID.

Data: None
Severity: Warning
Log: LOG_SLI
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2581: Not enough WQs (<slr4_hba.max_cfg_param.max_wq>) from the pci function for supporting FCP WQs (<cfg_fcp_wq_count>)

The driver was not configured with the minimum number of fast-path work queues.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2582: Not enough WQs (<max_wq>) from the pci function for supporting the requested FCP WQs (<cfg_wq_count>), the actual FCP WQs can be supported: <wq_count>

The driver was not configured with enough fast-path work queues.

Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.

2593: The FCP EQ count(<cfg_fcp_eq_count>) cannot be greater than the FCP WQ count(<cfg_fcp_wq_count>), limiting the FCP EQ count to <cfg_fcp_wq_count>

The fast-path event queue cannot be greater than the fast-path work queue count.

Data: None
Severity: Warning
Log: LOG_INIT
Action: None required.

2597: Mailbox command <vpi> (<mbxCommand>) cannot issue

Synchronou(2) slr_flag (3) flag

Data: None
Severity: Error
Log: LOG_MBOX, LOG_SLI
Action: None required.

2598: Adapter Link is disabled.

The adapter link has been disabled.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2599: Adapter failed to issue DOWN_LINK mbox command rc <rc>

The driver was unable to issue the DOWN_LINK mailbox command.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2600: lpfc_slr_read_serdes_param failed to allocate mailbox memory

Failed to allocate mailbox memory.

Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2605: lpfc_dump_static_vport: memory allocation failed
Failed to allocate mailbox memory.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

2606: No NPIV Fabric support
No NPIV Fabric support.
Data: None
Severity: Error
Log: LOG_ELS
Action: None required.

2607: Failed to allocate init_vpi mailbox
Failed to allocate the INIT_VPI mailbox command.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

2608: Failed to issue Init VPI mailbox
The driver was unable to send an INIT_VPI mailbox command.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

2609: Init VPI mailbox failed <u.mb.mbxStatus>
The INIT_VPI mailbox command failed.
Data: None
Severity: Error
Log: LOG_MBOX
Action: None required.

2610: UNREG_FCFI mbox allocation failed
Failed to allocate mailbox memory.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

2611: UNREG_FCFI issue mbox failed
Could not issue the UNREG_FCFI mailbox command.
Data: None
Severity: Error
Log: LOG_DISCOVERY, LOG_MBOX
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2619: Config region 23 has bad signature
Configuration region 23 has an invalid signature.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2620: Config region 23 has bad version
Configuration region 23 has an invalid version.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2621: Failed to allocate mbox for query firmware config cmd
Failed to allocate mailbox memory.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

2622: Query Firmware Config failed mbx status <rc>, status <shdr_status> add_status <shdr_add_status>
Could not read the firmware configuration.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2623: FCoE Function not supported by firmware. Function mode = <function_mode>>
FCoE is not supported by this firmware.
Data: None
Severity: Error
Log: LOG_SLI
Action: Use the OneCommand Manager application to update to the latest firmware.

2707: Ring <Ring#> handler: Failed to allocate iocb Rctl <fh_rctl> Type <fh_type> received
The driver was unable to allocate memory to send a query config mailbox command.
Data: None
Severity: Error
Log: LOG_SLI
Action: None required.

2717: CT context array entry [<index>] over-run: oxid:<fh_ox_id>, sid:<fh_SID>
All of the array slots to hold buffers that are passed to the application are in use.
Data: None
Severity: Warning
Log: LOG_ELS
Action: None required.

Table 4-4 LPFC Error Log Messages and their Descriptions (Continued)

2718: Clear Virtual Link Received for VPI <index> tag <event_tag>
A Clear virtual link was received from the Fabric for this VPI.
Data: None
Severity: Error
Log: LOG_DISCOVERY
Action: None required.

2719: Invalid response length: tgt <TGT_ID> lun <LUN> cmd <CMD> rsplen <RSPLEN>
The response length for this FCP command is not supported.
Data: None
Severity: Error
Log: LOG_FCP
Action: None required.

2721: ndlp null for oxid %x SID %x\n, icmd->ulpContext, dfchba->ct_ctx[tag].SID);
The Node value for this SID is not in the node list.
Data: None
Severity: Warning
Log: LOG_ELS
Action: None required.

2726: READ_FCF_RECORD Indicates empty FCF table
The driver requested the firmware provide a list of FCF entries to connect to and the firmware responded that the FCF table is empty.
Data: None
Severity: Error
Log: LOG_INIT
Action: None required.

Ethernet Driver Log Messages

Retrieving Ethernet Driver Log Messages

Ethernet driver messages are logged in the `/var/log/messages` file. This log file is an ASCII text file and can be viewed and searched with a text editor such as `vim`. A log file automatically rotates as it gets larger. Rotated log files are named `messages.x`, where “x” is an integer.

To search the log file for error messages, at the command prompt type:

```
# cd /var/log
# vim messages
```

A message is displayed similar to the following:

```
Aug 15 09:57:48 S74 kernel: Invalid MTU requested. Must be between
256 and 8974 bytes
```

Ethernet Driver Log Messages and their Descriptions

When reporting an issue with the OneConnect UCNA, check the kernel message log using the `dmesg` command or the `/var/log/messages` file, and report any of these entries that may be present.

There are three types of Ethernet log messages: error, information, and warning.

Ethernet driver warning messages logged by a OneConnect UCNA start with “be2net <BID>”, where <BID> is the PCI bus identifier string. For example:

```
be2net 0000:0d:00.1: MTU must be between 256 and 9000 bytes.
```

Note: In the following table, <D>, <DD>, or <DDD> refers to decimal values that appear in the log messages.

Table 4-5 Ethernet Driver Log Messages of Type

Ethernet Log Message	Description	Type
<D> pending tx-completions	The driver did not get completions for some transmit requests from the NIC while unloading the driver. This is usually indicates an issue with the NIC.	Error
cmd to write to flash rom failed. type/op <D>/<D>:	The firmware command to install the latest firmware version on the OneConnect UCNA failed with the indicated error code.	Error
Could not create sysfs group	The creation of the “flash_fw” entry under the <code>/sys/class/net/eth<x></code> failed. The driver is fully functional but you cannot install later firmware versions on the OneConnect UCNA.	Error
Could not get crc from flash, not flashing redboot	The driver could not get enough information from the NIC to decide whether or not Redboot (the primary boot program) should be flashed. The driver skips updating this section. This is a very unlikely error.	Error
Could not set PCI DMA Mask	The operating system call to set the DMA mask failed.	Error

Table 4-5 Ethernet Driver Log Messages of Type (Continued)

Ethernet Log Message	Description	Type
Emulex OneConnect 10Gbps NIC initialization failed	Initialization of the UCNA or allocation of a resource for initializing the driver failed. In most cases, this message is accompanied by a more specific error message. Try rebooting the system after power cycling. If the issue persists, this could be a symptom of a hardware issue or corrupted firmware.	Error
Error in cmd completion: status(compl/extd)= <DDD>/<DDD>	A firmware command failed with the indicated status code and extended status code.	Error
Firmware <filename> load error (signature did not match)	The firmware image under /lib/firmware/<filename> did not have the signature of a proper OneConnect firmware image file. The firmware in this file will not be flashed. Copy the proper file and try flashing again.	Error
Firmware flashed successfully	This is an informational message that the firmware on the OneConnect UCNA has been updated.	Information
Firmware load error	Updating the OneConnect UCNA with new firmware failed. Usually this message is accompanied by a detailed messages on the failure.	Error
Flash cookie not found in firmware image	The firmware image under /lib/firmware/<filename> does not have the expected cookie. The firmware in this file will not be flashed. Copy the proper file and try flashing again.	Error
Flashing firmware file <filename>	This is an informational message that the firmware in the OneConnect UCNA is being updated with the firmware image in the file indicated.	Information
INTx Request IRQ failed - err <DDD>	The request for INTx interrupt registration failed. The driver is non-functional if the INTx interrupt cannot be registered.	Error
Invalid mailbox completion	The driver received an unexpected completion status for a firmware command.	Error
Link down Link up	This is an informational message about a change in Link status.	Information
mbox_db poll timed out	A firmware command did not complete in a reasonable time. This is most likely due to a firmware or hardware issue. Try rebooting the system.	Error
Memory allocation failure while flashing	The driver could not allocate memory required for flashing the firmware image. Try flashing later. If this issue persists, try flashing after a system reboot.	Error
Module param rx_frag_size must be 2048/4096/8192. Using 2048	An unsupported receive buffer size was passed for the rx_frag_size module parameter. The driver ignores the specified value and uses the default RX buffer size of 2048.	Warning
MSIX request IRQ failed -err <DDD>	The request for MSI-X interrupt registration failed. The driver will use INTx interrupts.	Warning

Table 4-5 Ethernet Driver Log Messages of Type (Continued)

Ethernet Log Message	Description	Type
MTU changed from <DDD> to <DDD>	This is an informational message that the MTU value changed as requested.	Information
MTU must be between 256 and 9000 bytes	Request to change the MTU was issued with an invalid MTU value. The request failed and MTU will not be changed.	Warning
OneConnect card in a Gen<D> x<D> PCIe slot. It should be in a Gen2 x8 slot for best performance	The driver found a OneConnect UCNA in a PCIe slot that is Gen1 or less than x8 wide. The UCNA will continue to work but will not provide the best performance.	Warning
Pause param set failed	The firmware command to set pause frame settings failed.	Warning
POST failed	The power-on self test of the UCNA failed. This is an indication of a hardware or firmware issue. Try rebooting the system after a reset.	Error

iSCSI Driver Log Messages

Retrieving iSCSI Driver Error Log Messages

For Linux systems, the iSCSI driver generates log messages to the `/var/log/messages` file. The log file is an ASCII text file and can be viewed and searched with your preferred text editor.

To search the log file for error messages, at the command prompt type:

```
# cd /var/log
# vim messages
```

iSCSI Driver Error Log Messages and their Descriptions

All iSCSI driver error log messages are preceded by a prefix of “`scsiX:`” (if the SCSI host controller can be determined, with “`X`” being a number representing the Linux SCSI host controller) and “`OneConnect iSCSI Driver:`”, for example:

```
scsi2: OneConnect iSCSI Driver: Rejected IOCTL since buffer size
limit exceeded

scsi2: OneConnect iSCSI Driver: Subsystem / Opcode = 0x12345678 /
0x11223344

scsi2: OneConnect iSCSI Driver: Payload Length = 0x1000
```

Note: Some error log entries may be followed by additional entries that provide further information.

Note: In the following table, “`0xX`” refers to a hexadecimal value that appears in the log messages.

Table 4-6 lists iSCSI driver error log messages and their descriptions.

Table 4-6 iSCSI Driver Error Log Messages

iSCSI Error Log Message	Description
Did not receive an iSCSI Command window update from Target for at least 25 Secs. Session Handle	Check for any errors reported at the target. The Emulex iSCSI initiator is only supported with certified targets. Check for software updates at the target vendor’s website and the Emulex website. If this fails, contact technical support.
Driver version does not match Firmware. Please run Installer.	This failure indicates that the driver version that is running on the system does not match the version of the firmware flashed on the board. This issue can be addressed by running the installer from the desired version.
Error during iSCSI offload Session Handle / Firmware Error code	This may indicate a target is in error or may point to transient network connectivity issues. It may also indicate a firmware error.

Table 4-6 iSCSI Driver Error Log Messages (Continued)

iSCSI Error Log Message	Description
Extended Timeout Expired. Loss of connection to target exceeded ETO limits. Session ID = 0xX	Check the connection to the target or the state of the target device. If the target is made available, any sessions that existed previously will be reestablished and the devices will be available for I/O.
Initialization Failure	This failure may be due to the firmware not being present or running currently. This failure may also indicate a hardware issue.
Initialization failure during Power Management Bootup	This failure may be due to the firmware not being present or running currently. This failure may also indicate a hardware issue.
Internal API failed during Initialization	This failure may indicate a low memory condition.
Hardware Initialization Failed. Either Hardware/Firmware is not initialized or is malfunctioning.	This failure indicates that the hardware has not been initialized or is malfunctioning. This may also indicate that the firmware is not running correctly.
LinkDown Timeout Expired. Please check the Physical Link to OneConnect.	Check the links to the UCNA. If the link is reestablished, any sessions that existed previously will be reestablished and the devices will be available for I/O.
OSM Hardware Initialization Failure	This failure indicates that the hardware has not been initialized or is malfunctioning. This may also indicate that the firmware is not running correctly.
OSM Resource Allocation Failure	The operating system failed to allocate resources for the device. Check low memory conditions and operating system hardware resource conflicts.
Received a TMF Abort for an I/O that is not present with the driver.	This may indicate a slow connection to the target. Check network connectivity to the target for any errors.
Received invalid iSCSI Command Sequence Number update from Target. Session Handle = 0xX MaxCmdSN = 0xX ExpCmdSN = 0xX	Check for any errors reported at the target. The Emulex iSCSI initiator is only supported with certified targets. Check for software updates at the target vendor's website and the Emulex website. If this fails, contact technical support.
Received unsupported Task Management Function. Task Management Function code = 0xX	The operating system version is not supported.
Rejected IOCTL since buffer size limit exceeded. Subsystem / Opcode = 0xX / 0xX Payload Length = 0xX	This error may indicate an incorrect configuration option for the iSCSI driver. It may also indicate a low memory condition.
Unrecoverable Error UE_LOW = 0xX UE_HIGH = 0xX Firmware Line Number = 0xX	This may be due to hardware errors or due to unhandled exceptions in the hardware or firmware.

Appendix A. Open-iSCSI Support

Open-iSCSI is a high-performance, transport independent, multi-platform implementation of Request for Comments (RFC) 3720.

The inbox Open-iSCSI driver is supported. However, for driver versions earlier than RHEL 6.2 and SLESS 11 SP2, due to limitations in some Open-iSCSI distribution tools, the inbox driver does not support ISNS or iSCSI boot. If you need to use ISNS or iSCSI boot, use the out-of-box iSCSI driver provided by Emulex.

For the inbox Open-iSCSI driver, use the Open-iSCSI tools to configure and manage Open-iSCSI. For the Emulex out-of-box iSCSI driver, use the Emulex OneCommand Manager application to configure and manage Open-iSCSI. The applications are not interchangeable between the drivers.

The iscsiadm utility is a command-line tool allowing discovery of iSCSI targets; logging into iSCSI targets; as well as, access and management of the open-iSCSI database on all Linux installations. This utility presents a set of operations that you can perform on iSCSI nodes, sessions, connections, and discovery records.

Discovering and Adding iSCSI Targets

When discovering and adding iSCSI targets, there are two sets of instructions depending on the operating system. One set is used for RHEL 5.6 and RHEL 5.7. The other set is for the other supported operating systems, which include: SLES 10 SP3 and later versions; SLES 11 SP1 and later versions; and RHEL 5.8 and later versions.

RHEL 5.6 and RHEL 5.7

For RHEL 5.6 and RHEL 5.7, you can add iSCSI targets using a discovery phase first (see the following section) or add them directly (see “Adding iSCSI Targets Directly (without Discovering)” on page 152).

Discovering and Adding iSCSI Targets

The following steps are used to discover and add a target using the Ethernet adapter for RHEL 5.6 and RHEL 5.7. In this example, based on the IPL file flashed, the adapter has two NIC functions and two iSCSI functions.

1. For the iSCSI function, set the IP address using the iSCSI Select utility.
 - The initiator IP is set to 192.168.65.99. This IP is set for the iSCSI function, but it does not get set for the Ethernet interface.
 - The target portal IP is 192.168.65.196.
2. Using the iSCSI Select, iSCSI Target Configuration option, see if the targets can be discovered, but do not add the targets yet.
3. The following command displays the specific host information.

```
iscsiadm -m host --print=4
```

Example output:

```
Host Number: 8
```

```

State: running
Transport: be2iscsi
Initiatorname: <empty>
IPaddress: <empty>
HWaddress: 00:00:c9:f2:73:8d
Netdev: <empty>

```

```

Host Number: 9
State: running
Transport: be2iscsi
Initiatorname: <empty>
IPaddress: <empty>
HWaddress: 00:00:c9:f2:73:91
Netdev: <empty>

```

4. Display the information for all Ethernet interfaces.

```
ifconfig -a
```

Example output:

```

eth0      Link encap:Ethernet  HWaddr 00:25:B3:DF:8D:AC
inet addr:10.192.199.36  Bcast:10.192.207.255  Mask:255.255.240.0
inet6 addr: fe80::225:b3ff:fedf:8dac/64  Scope:Link
UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
RX packets:306140 errors:0 dropped:0 overruns:0 frame:0
TX packets:2672 errors:0 dropped:0 overruns:0 carrier:0
collisions:0 txqueuelen:1000
RX bytes:23651804 (22.5 MiB)  TX bytes:476993 (465.8 KiB)
Interrupt:169 Memory:f4000000-f4012800

```

```

eth1      Link encap:Ethernet  HWaddr 00:25:B3:DF:8D:AE
BROADCAST MULTICAST  MTU:1500  Metric:1
RX packets:0 errors:0 dropped:0 overruns:0 frame:0
TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
collisions:0 txqueuelen:1000
RX bytes:0 (0.0 b)  TX bytes:0 (0.0 b)
Interrupt:154 Memory:f2000000-f2012800

```

```

eth2      Link encap:Ethernet  HWaddr 00:25:B3:DF:8D:B0
BROADCAST MULTICAST  MTU:1500  Metric:1
RX packets:0 errors:0 dropped:0 overruns:0 frame:0
TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
collisions:0 txqueuelen:1000
RX bytes:0 (0.0 b)  TX bytes:0 (0.0 b)
Interrupt:210 Memory:f8000000-f8012800

```

```

eth3      Link encap:Ethernet  HWaddr 00:25:B3:DF:8D:B2
BROADCAST MULTICAST  MTU:1500  Metric:1
RX packets:0 errors:0 dropped:0 overruns:0 frame:0
TX packets:0 errors:0 dropped:0 overruns:0 carrier:0

```

```

collisions:0 txqueuelen:1000
RX bytes:0 (0.0 b) TX bytes:0 (0.0 b)
Interrupt:218 Memory:f6000000-f6012800

eth4      Link encap:Ethernet HWaddr 00:00:C9:F2:73:8C
BROADCAST MULTICAST MTU:1500 Metric:1
RX packets:0 errors:0 dropped:0 overruns:0 frame:0
TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
collisions:0 txqueuelen:1000
RX bytes:0 (0.0 b) TX bytes:0 (0.0 b)

eth5      Link encap:Ethernet HWaddr 00:00:C9:F2:73:90
BROADCAST MULTICAST MTU:1500 Metric:1
RX packets:42 errors:0 dropped:0 overruns:0 frame:0
TX packets:75 errors:0 dropped:0 overruns:0 carrier:0
collisions:0 txqueuelen:1000
RX bytes:11738 (11.4 KiB) TX bytes:10590 (10.3 KiB)

```

5. Use the `ethtool` to find the interface on the Ethernet adapter for which the link is up.

```
ethtool eth5
```

Example output:

```

Settings for eth5:
Supported ports: [ FIBRE ]
Supported link modes:  10000baseT/Full
Supports auto-negotiation: No
Advertised link modes:  Not reported
Advertised auto-negotiation: No
Speed: 10000Mb/s
Duplex: Full
Port: FIBRE
PHYAD: 1
Transceiver: external
Auto-negotiation: off
Supports Wake-on: g
Wake-on: d
Link detected: yes

```

6. Set the IP address for the Ethernet interface.

```
ifconfig eth5 192.168.65.10 up
```

7. Display the ifaces that are present.

```
iscsiadm -m iface
```

Example output:

```

default tcp,<empty>,<empty>,<empty>,<empty>
iser iser,<empty>,<empty>,<empty>,<empty>
bnx2i.00:25:b3:df:8d:ad bnx2i,00:25:b3:df:8d:ad,<empty>,<empty>,<empty>
be2iscsi.00:00:c9:f2:73:91
be2iscsi,00:00:c9:f2:73:91,<empty>,<empty>,<empty> "This is the
interface for which IP was set using iSCSI Select"

```

```
be2iscsi.00:00:c9:f2:73:8d
be2iscsi,00:00:c9:f2:73:8d,<empty>,<empty>,<empty>
bnx2i.00:00:00:00:00:00 bnx2i,00:00:00:00:00:00,<empty>,<empty>,<empty>
```

8. Discover the targets using the Ethernet interface for which the IP was set.

```
iscsiadm -m discovery -t st -p 192.168.65.196:3260 -I
be2iscsi.00:00:c9:f2:73:91
```

Example output:

```
192.168.65.196:3260,1 iqn.tgt0
192.168.65.196:3260,1 iqn.tgt1
192.168.65.196:3260,1 iqn.tgt2
192.168.65.196:3260,1 iqn.tgt3
192.168.65.196:3260,1 iqn.tgt4
192.168.65.196:3260,1 iqn.tgt5
192.168.65.196:3260,1 iqn.tgt6
192.168.65.196:3260,1 iqn.tgt7
```

The “-I be2iscsi.00:00:c9:f2:73:91” option is passed during discovery so that the discovered portals are bound with the interface name, and the login occurs through those ifaces.

9. Log in to the discovered targets.

```
iscsiadm -m node -p 192.168.65.196:3260 -l
```

Example output:

```
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt7,
portal: 192.168.65.196,3260]
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt3,
portal: 192.168.65.196,3260]
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt1,
portal: 192.168.65.196,3260]
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt0,
portal: 192.168.65.196,3260]
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt6,
portal: 192.168.65.196,3260]
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt4,
portal: 192.168.65.196,3260]
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt2,
portal: 192.168.65.196,3260]
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt5,
portal: 192.168.65.196,3260]
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt7, portal:
192.168.65.196,3260] successful.
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt3, portal:
192.168.65.196,3260] successful.
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt1, portal:
192.168.65.196,3260] successful.
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt0, portal:
192.168.65.196,3260] successful.
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt6, portal:
192.168.65.196,3260] successful.
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt4, portal:
```

```
192.168.65.196,3260] successful.
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt2, portal:
192.168.65.196,3260] successful.
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt5, portal:
192.168.65.196,3260] successful.
```

10. Display the session information after a successful login.

Once the login is successful, the session information can be dumped with the following command.

```
iscsiadm -m session -P 3
```

11. Log out from the targets.

To log out from a single target, use the following command.

```
iscsiadm -m node -T <Target_Name> -u
```

To log out from all targets, use the following command.

```
iscsiadm -m node -u
```

Adding iSCSI Targets Directly (without Discovering)

If the target details are known, you can skip the discovery phase and add the targets to the iscsiadm node database directly. Then you can log in to those targets.

1. Add individual target details to the nodes database.

```
iscsiadm -m node -o new -T iqn.tgt1 -p 192.168.65.196:3260 -I
be2iscsi.00:00:c9:f2:73:91
```

Example output:

```
New iSCSI node
[be2iscsi:[hw=00:00:c9:f2:73:91,ip=,net_if=,iscsi_if=be2iscsi.00:00:c9:
f2:73:91] 192.168.65.196,3260,-1 iqn.tgt1] added
```

The “-I be2iscsi.00:00:c9:f2:73:91” option is passed so that the target is bound with the interface name, and the login occurs through those ifaces.

2. Log in to the specific target.

```
iscsiadm -m node -T iqn.tgt1 -l
```

Example output:

```
Logging in to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt1,
portal: 192.168.65.196,3260]
Login to [iface: be2iscsi.00:00:c9:f2:73:91, target: iqn.tgt1, portal:
192.168.65.196,3260] successful.
```

3. To display the session after a successful log in or to log out, see 10 and step 11 on page 152, respectively.

SLES 10 SP3, SLES 11 SP1, RHEL 5.8, and Later Versions

Discovering and Adding iSCSI Targets

The following steps are used to discover and add a target using the Ethernet adapter for SLES 10 SP3, SLES 11 SP1, RHEL 5.8, and later versions.

- For the iSCSI function, set the IP address using the iSCSI Select utility.
 - The initiator IP is set to 192.168.65.99. This IP is set for the iSCSI function, but it does not get set for the Ethernet interface.
 - The target portal IP is 192.168.65.196.
- Discover the targets using the Ethernet interface for which the IP was set.


```
iscsiadm -m discovery -t st -p 192.168.65.196:3260 -I
be2iscsi.00:00:c9:f2:73:91
```
- Log in to the discovered targets.


```
iscsiadm -m node -p 192.168.65.196:3260 -l
```
- To display the session after a successful log in or to log out, see 10 and step 11 on page 152, respectively.

iscsiadm Commands for Configuring the Target

After setting up the target and initiator machines, use the following procedure to configure the iSCSI target through Open-iSCSI.

Note: The following instructions assume that the initiator machine is a Linux machine with a OneConnect UCNA installed.

- Create a new interface (iface):


```
iscsiadm -m iface -o new -I <ifacename>
```

In this command, <ifacename> is the name you provide for the iface.
- Add a NIC media access control (MAC) address:


```
iscsiadm -m iface -I <ifacename> --op=update -n iface.hwaddress
-v <NIC mac_address>
```

In this command, <ifacename> is the name of the iface created in step 1 and <NIC mac_address> is the NIC MAC address you are adding for discovery.
- Add the Transport Name:


```
iscsiadm -m iface -I <ifacename> --op=update -n
iface.transport_name -v be2iscsi
```

In this command, <ifacename> is the name of the iface created in step 1.
- Restart the service:


```
service open-iscsi restart
```
- Perform target discovery using "SendTargets":


```
iscsiadm -m discovery -t st -p <ip:port> -I <iface> -P 1
```

In this command, <ip:port> is the IP address and port number and <iface> is the name of the iface created in step 1.

6. Add the iSCSI driver MAC address:

```
iscsiadm -m iface -I <ifacename> --op=update -n iface.hwaddress
-v <iSCSI MAC Address>
```

In this command, <ifacename> is the name of the iface created in step 1.

7. Set the initiator's IP address:

```
iscsiadm -m iface -I <ifacename> -o update -n iface.ipaddress -v
<IP>
```

In this command, <ifacename> is the name of the iface created in step 1 and <IP> is the IP address.

8. Log into the target:

```
iscsiadm -m node -T <targetname> -p <ip:port> -I <iface> -l
```

In this command, <targetname> is the target name you want to log into; <ip:port> is the IP address and port number; and <iface> is the name of the iface created in step 1.

9. Verify that the disk shows:

```
fdisk -l
```

10. Log out:

```
iscsiadm -m node -T <targetname> -p <ip:port> -I <iface> -u
```

In this command, <targetname> is the target name; <ip:port> is the IP address and port number; and <iface> is the name of the iface created in step 1.

Example

```
iscsiadm -m iface -o new -I iface0
iscsiadm -m iface -I iface0 --op=update -n iface.hwaddress -v 00:0F:1F:62:2B:BF
iscsiadm -m iface -I iface0 --op=update -n iface.transport_name -v be2iscsi
service open-iscsi restart
iscsiadm -m discovery -t st -p 20.0.0.107:3260 -I iface0 -P 1
iscsiadm -m iface -I iface0 --op=update -n iface.hwaddress -v 00:0F:1F:92:6B:BF
iscsiadm -m iface -I iface0 -o update -n iface.ipaddress -v 20.0.0.107
iscsiadm -m node -T iqn.tgt0 -p 20.0.0.107:3260 -I iface0 -l
fdisk -l
iscsiadm -m node -T iqn.tgt0 -p 20.0.0.107:3260 -I iface0 -u
```

Appendix B. Configuring iSCSI Through DHCP

IP Address Reservation

If you are using the Dynamic Host Configuration Protocol (DHCP) server to obtain an IP address for your iSCSI initiator, Emulex recommends that you set up a reservation. A reservation assigns a specific IP address based on the MAC address of your iSCSI function.

If you do not reserve an IP address through DHCP, you must set the lease length for the iSCSI initiator's IP address to unlimited. This allows the IP address lease not to expire.

DHCP Option 43 (Vendor-Specific Information)

This section describes the format for the data returned in DHCP option 43. The method and format for specifying the Vendor ID is outside the scope of this document and is not included here. The initiator offers this Vendor ID to the DHCP server to retrieve data in the format as described in "DHCP Option 43 Format" on page 156.

DHCP Option 43 Parameter Descriptions

Table B-1 describes the parameters used in the format (data string) for option 43.

Table B-1 DHCP Option 43 (Vendor-Specific Information)

Parameter	Description	Field Type
<TargetIP>	Replace with a valid IPv4 address in dotted decimal notation.	Mandatory
<TargetTCPPort>	Replace with a decimal number ranging from 1 to 65535 (inclusive). The default TCP port (3260) is assumed, if a value is not specified.	Optional
<LUN>	A hexadecimal representation of the Logical Unit Number of the boot device. Replace with an eight-byte number that should be specified as a hexadecimal number consisting of 16 digits, with an appropriate number of zeros padded to the left, if required. If a value is not provided, LUN 0 is assumed to be the boot LUN.	Optional
<TargetName>	Replace with a valid target iSCSI Qualified Name (IQN) name of up to 223 characters.	Mandatory
<InitiatorName>	Replace with a valid initiator IQN of up to 223 characters. If a value is not provided, the default initiator name (generated by the OneConnect UCNA based on the board's MAC address) is used.	Optional
<HeaderDigest>	Replace with either "E" or "D". <ul style="list-style-type: none"> "E" denotes that the header digest is enabled. "D" denotes that the header digest is disabled. 	Optional

Table B-1 DHCP Option 43 (Vendor-Specific Information)

Parameter	Description	Field Type
<DataDigest>	Replace with either “E” or “D”. <ul style="list-style-type: none"> “E” denotes that the data digest is enabled. “D” denotes that the data digest is disabled. If a value is not provided, it is assumed that the Data Digest is disabled by default.	Optional
<AuthenticationType>	If applicable, replace with “D”, “E”, or “M”. <ul style="list-style-type: none"> “D” denotes that authentication is disabled. “E” denotes that one-way CHAP is enabled (the username and secret to be used for one way CHAP must be specified by non-DHCP means). “M” denotes that mutual CHAP is enabled (user name and passwords required for mutual CHAP authentication must be specified by non-DHCP means). If a value is not specified, this field defaults to authentication disabled.	Optional

DHCP Option 43 Format

The following is the format of DHCP option 43 and its guidelines for creating the data string:

```
`iscsi:'<TargetIP>':'<TargetTCPPort>':'<LUN>':'<TargetName>':'<InitiatorName>':'<HeaderDigest>':'<DataDigest>':'<AuthenticationType>'
```

- Strings shown in quotes are part of the syntax and are mandatory.
- Fields enclosed in angular brackets (including the angular brackets) should be replaced with their corresponding values. Some of these fields are optional and may be skipped.
- If an optional field is skipped, a colon must be used as a placeholder to indicate the default value for that field.
- When specified, the value of each parameter should be enclosed in double quotes.
- All options are case sensitive.

Default Initiator Name and Data Digest Settings Example

The following is an example of default initiator name and data digest settings.

```
iscsi:"192.168.0.2":"3261":"0000000000000000E":"iqn.2009-4.com:1234567890":"E":"E"
```

In this example, the field values are:

- TargetIP: 192.168.0.2
- TargetTCPPort: 3261
- LUN: 0x0E
- TargetName: iqn.2009-04.com:1234567890

- InitiatorName: Not specified. Use the Initiator name already configured. Use the default name if none was configured.
- HeaderDigest: Enabled
- DataDigest: Not specified. Assume disabled.
- AuthenticationType: One-way CHAP is enabled.

Default TCP and Mutual CHAP Settings Example

The following is an example of default TCP port and mutual CHAP settings.

```
iscsi:"192.168.0.2"::"0000000000000000E":"iqn.2009-4.com:1234567890"  
::"E":"D":"M"
```

In this example, the field values are:

- TargetIP: 192.168.0.2
- TargetTCPPort: Use default from RFC 3720 (3260)
- LUN: 0x0E
- TargetName: iqn.2009-04.com:1234567890
- InitiatorName: Not specified. Use the Initiator name already configured. Use the default name if none was configured.
- HeaderDigest: Enabled
- DataDigest: Disabled
- AuthenticationType: Mutual CHAP is enabled.

Appendix C. OneConnect 10 GbE UCNA Port Speed Specifications

Port Speed Negotiation on Non-Mezzanine Cards

OneConnect 10 GbE UCNAs can support only one Ethernet port speed at a time. On non-mezzanine cards, its preference is 10 Gb/s. The type of module used (copper/optical) does not make a difference. If a 10 Gb/s module is plugged into one of the ports, the UCNA runs at a 10 Gb/s speed regardless of its other port's speed, even if I/Os are running on that port. This behavior is an adapter constraint; another adapter can be running on a different speed.

Table C-1 lists negotiated speed specifications per OneConnect 10 GbE UCNA port connection.

Table C-1 OneConnect 10 GbE UCNA Negotiated Speed Specifications

Port 0 Speed (Gb/s)	Port 1 Speed (Gb/s)	Port Link Status	OneConnect 10 GbE UCNA Speed (Gb/s)
10	10	Both ports are link up	10
10	1	Only Port 0 is link up	10
1	10	Only Port 1 is link up	10
1	1	Both ports are link up	1
1	-	Only Port 0 is link up	1
-	1	Only Port 1 is link up	1
10	-	Only Port 0 is link up	10
-	10	Only Port 1 is link up	10

Port Speed on UCNA Mezzanine Cards

For UCNA mezzanine cards, only one Ethernet port speed is supported at a time. Its speed is the first negotiated speed (either 1 Gb/s or 10 Gb/s), depending on the switch that is connected.

To change the speed on these cards:

1. Remove the switch from both of the ports.
2. Insert the switch for one port and wait for the link to come up.

The UCNA retains this speed until both links are down.

Appendix D. Updating Ethernet Firmware

The Emulex Ethernet driver supports updating the firmware image in the UCNA flash through the `request_firmware` interface in Linux. You can perform this update when the UCNA is online and passing network/storage traffic.

To update the Ethernet firmware image:

1. Copy the latest firmware image under the `/lib/firmware` directory:

```
# cp be3flash.ufi /lib/firmware
```

2. Start the update process.

- In Linux distributions that support the flash option in `ethtool` (for example, SLES 11 SPx and RHEL 6.x), use the following command:

```
# ethtool -f eth<X> be3flash.ufi 0
```

- In older Linux distributions (for example, SLES 10 SPx and RHEL 5.x), write the name of the flash image file in the `sysfs` node:

```
# echo 60 > /sys/class/firmware/timeout
```

```
# echo be3flash.ufi > /sys/class/net/eth<X>/flash_fw
```

3. Reboot the system to enable the new firmware image to take effect.

Appendix E. Ethtool -S Option Statistics

Table E-1 contains a list of ethtool -S option statistics and their descriptions. Table E-2 contains a list of transmit/receive statistics per receive queue basis.

Table E-1 Ethtool -S Option Statistics

Name	Description
rx_crc_errors	The number of packets dropped due to CRC errors.
rx_alignment_symbol_errors	The number of packets dropped due to L1 alignment errors. This counter is on a per-port basis.
rx_pause_frames	The number of Ethernet pause frames (flow control) received.
rx_control_frames	The number of control frames received.
rx_in_range_errors	Received packets dropped when the Ethernet length field is not equal to the actual Ethernet data length.
rx_out_range_errors	Received packets dropped when their length field is ≥ 1501 bytes and ≤ 1535 bytes.
rx_frame_too_long	Received packets dropped when they are longer than 9216 bytes.
rx_address_filtered	Received packets dropped when they don't pass the unicast or multicast address filtering.
rx_dropped_too_small	Received packets dropped when IP packet length field is less than the IP header length field.
rx_dropped_too_short	Received packets dropped when IP length field is greater than the actual packet length.
rx_dropped_header_too_small	Received packets dropped when the IP header length field is less than 5.
rx_dropped_tcp_length	Received packets dropped when the TCP header length field is less than 5 or the TCP header length + IP header length is more than IP packet length.
rx_dropped_runt	Dropped receive packets due to runt packets (for example, packets shorter than the Ethernet standard).
rxpp_fifo_overflow_drop, rx_input_fifo_overflow_drop	Number of received packets dropped when a FIFO for descriptors going into the packet demux block overflows. In normal operation, this FIFO must never overflow.
rx_ip_checksum_errs, rx_tcp_checksum_errs, rx_udp_checksum_errs	Packets dropped due to TCP/IP/UDP checksum errors.
rx_switched_unicast_packets, rx_switched_multicast_packets, rx_switched_broadcast_packets	The number of unicast, multicast, and broadcast packets switched internally.
tx_pauseframes	The number of Ethernet pause frames (flow control) transmitted per port.
tx_controlframes	The number of Ethernet control frames transmitted per port.

Table E-1 Ethtool -S Option Statistics (Continued)

Name	Description
rx_priority_pause_frames	The number of Ethernet priority pause frames (priority flow control) received per port.
pmem_fifo_overflow_drop	Received packets dropped when an internal FIFO going into main packet buffer tank (PMEM) overflows.
jabber_events	The number jabber packets received. Jabber packets are packets that are longer than the maximum size Ethernet frames and that have bad CRC.
rx_drops_no_pbuf	Packets dropped due to lack of available HW packet buffers used to temporarily hold the received packets.
rx_drops_no_erx_descr	Received packets dropped due to the input receive buffer descriptor FIFO overflowing.
rx_drops_no_tpre_descr	Packets dropped because the internal FIFO to the offloaded TCP receive processing block is full. This could happen only for offloaded iSCSI or FCoE traffic.
rx_drops_too_many_frags	Received packets dropped when they need more than 8 receive buffers. This counter will always be 0.
forwarded_packets	The number of packets generated by ASIC internally. These packets are not handed to the host. This counter is shared across ports and all functions (NIC/FCoE/iSCSI).
rx_drops_mtu	Received packets dropped when the frame length is more than 9018 bytes.
eth_red_drops	Received packets dropped due to ASIC's Random Early Drop policy.
on_die_temperature	Die temperature on the ASIC.
link_down_reason	The reason ASIC signaled the link status as down. The various values are: <ul style="list-style-type: none"> 0 - Link down due to reasons other than those listed here. 1 - Link down caused by Dynamic Control channel protocol. 3 - Link down triggered by Virtual NIC configuration (for example: zero bandwidth assigned to a VNIC). 4 - Link down caused by Ethernet Pause frame flooding. 5 - Link down due to physical thermal temperature going up.

Table E-2 Transmit/Receive Queue Statistics

Statistic	Description
rxq<x>:rx_bytes	The number bytes received by the driver.
rxq<x>:rx_pkts	The number of packets received by the driver.
rxq<x>:rx_compl	The number of receive completions signaled to the driver by ASIC.
rxq<x>:rx_mcast_pkts	The number of multicast packets received by the driver.
rxq<x>:rx_post_fai	The number of times the driver could not post received buffers to ASIC.

Table E-2 Transmit/Receive Queue Statistics (Continued)

Statistic	Description
rxq<x>:rx_drops_no_skbs	The number of times the driver could not allocate socket buffers.
rxq<x>:rx_drops_no_frags	Packets dropped due to insufficient buffers posted by the driver.
txq<x>:tx_compl	The number of transmit completions signaled by ASIC.
txq<x>:tx_bytes	The number of bytes transmitted by the driver.
txq<x>:tx_pkts	The number of packets transmitted by the driver.
txq<x>:tx_reqs	The number of transmit request generated by the driver.
txq<x>:tx_wrbs	The number of transmit work request buffers generated by the driver. WRB describes the nature of each transmit request to ASIC.
txq<x>:tx_stops	The number of times the driver requests the host to stop giving further transmit requests since the hardware transmit queue is filled up.